



The Science of Neural Networks

Ryan Gillard

Machine Learning on Google Cloud Platform

The Art of ML

Hyperparameter Tuning

A Pinch of Science

The Science of Neural Networks

Embeddings

Custom Estimator




Introduction to Neural Networks

Ryan Gillard


Feature crosses help
linear models work in
nonlinear problems



But there tends to
be a limit ...




Combine features as an alternative to feature crossing




Structure the model so that features are combined
Then the combinations may be combined

How to choose the combinations?
Get the model to learn them

A Linear Model can be
represented as nodes and edges




Add Complexity: Non-Linear?




Add Complexity: Non-Linear?

$$h_1 = w_1 * x_1 + w_4 * x_2 + w_7 * x_3$$




Add Complexity: Non-Linear?

$$h_2 = w_2 * x_1 + w_5 * x_2 + w_8 * x_3$$




Add Complexity: Non-Linear?

$$h_3 = w_3 * x_1 + w_6 * x_2 + w_9 * x_3$$



Add Complexity: Non-Linear?


$$output = w_{10} * h_1 + w_{11} * h_2 + w_{12} * h_3$$



Add Complexity: Non-Linear?

$$output = w_{10} * h_1 + w_{11} * h_2 + w_{12} * h_3$$

$$\begin{aligned} &= (w_{10} * w_1 + w_{11} * w_2 + w_{12} * w_3) * x_1 \\ &+ (w_{10} * w_4 + w_{11} * w_5 + w_{12} * w_6) * x_2 \\ &+ (w_{10} * w_7 + w_{11} * w_8 + w_{12} * w_9) * x_3 \end{aligned}$$




Add Complexity: Non-Linear?


$$output = w_{10} * h_1 + w_{11} * h_2 + w_{12} * h_3$$

$$\begin{aligned} &= (w_{10} * w_1 + w_{11} * w_2 + w_{12} * w_3) * x_1 \\ &+ (w_{10} * w_4 + w_{11} * w_5 + w_{12} * w_6) * x_2 \\ &+ (w_{10} * w_7 + w_{11} * w_8 + w_{12} * w_9) * x_3 \end{aligned}$$

$$= W_1 * x_1 + W_2 * x_2 + W_3 * x_3$$



More Complex: Non-Linear?



More Complex: Non-Linear?

$$H_1 = \begin{bmatrix} w_1 & w_2 & w_3 \\ w_4 & w_5 & w_6 \\ w_7 & w_8 & w_9 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}$$

3×1 3×3 3×1

More Complex: Non-Linear?

$$H_2 = \begin{bmatrix} w_{10} & w_{11} & w_{12} \\ w_{13} & w_{14} & w_{15} \\ w_{16} & w_{17} & w_{18} \end{bmatrix} \begin{bmatrix} w_1 & w_2 & w_3 \\ w_4 & w_5 & w_6 \\ w_7 & w_8 & w_9 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}$$

3×1 3×3 3×3 3×1

$$= \begin{bmatrix} W_1 & W_2 & W_3 \\ W_4 & W_5 & W_6 \\ W_7 & W_8 & W_9 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}$$

3×3 3×1


More Complex: Non-Linear?

$$\begin{aligned} \hat{y} &= \begin{bmatrix} w_{19} & w_{20} & w_{21} \end{bmatrix} \begin{bmatrix} w_{10} & w_{11} & w_{12} \\ w_{13} & w_{14} & w_{15} \\ w_{16} & w_{17} & w_{18} \end{bmatrix} \begin{bmatrix} w_1 & w_2 & w_3 \\ w_4 & w_5 & w_6 \\ w_7 & w_8 & w_9 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} \\ &= \begin{bmatrix} W'_1 & W'_2 & W'_3 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} \end{aligned}$$

Dimensions (row x column) for each matrix or vector:

- \hat{y} : 1x1
- $\begin{bmatrix} w_{19} & w_{20} & w_{21} \end{bmatrix}$: 1x3
- $\begin{bmatrix} w_{10} & w_{11} & w_{12} \\ w_{13} & w_{14} & w_{15} \\ w_{16} & w_{17} & w_{18} \end{bmatrix}$: 3x3
- $\begin{bmatrix} w_1 & w_2 & w_3 \\ w_4 & w_5 & w_6 \\ w_7 & w_8 & w_9 \end{bmatrix}$: 3x3
- $\begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}$: 3x1
- $\begin{bmatrix} W'_1 & W'_2 & W'_3 \end{bmatrix}$: 1x3
- $\begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}$: 3x1

Adding a Non-Linearity



Adding a Non-Linearity

$$\hat{y} = \underset{1 \times 1}{[w_{19} \quad w_{20} \quad w_{21}]} \underset{1 \times 3}{\begin{bmatrix} w_{10} & w_{11} & w_{12} \\ w_{13} & w_{14} & w_{15} \\ w_{16} & w_{17} & w_{18} \end{bmatrix}} \underset{3 \times 3}{f} \left(\underset{3 \times 3}{\begin{bmatrix} w_1 & w_2 & w_3 \\ w_4 & w_5 & w_6 \\ w_7 & w_8 & w_9 \end{bmatrix}} \underset{3 \times 1}{\begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}} \right)$$

$$= \underset{1 \times 3}{[w_{19} \quad w_{20} \quad w_{21}]} \underset{3 \times 3}{\begin{bmatrix} w_{10} & w_{11} & w_{12} \\ w_{13} & w_{14} & w_{15} \\ w_{16} & w_{17} & w_{18} \end{bmatrix}} \underset{3 \times 1}{\begin{bmatrix} \max(0, w_1 x_1 + w_2 x_2 + w_3 x_3) \\ \max(0, w_4 x_1 + w_5 x_2 + w_6 x_3) \\ \max(0, w_7 x_1 + w_8 x_2 + w_9 x_3) \end{bmatrix}}$$

Non-linearity Quiz

Why is it important adding non-linear activation functions to neural networks?


- A. Adds regularization
- B. Increases the number of dimensions
- C. Invokes early stopping
- D. Stops the layers from collapsing back into just a linear model

Non-linearity Quiz

Why is it important adding non-linear activation functions to neural networks?

- A. Adds regularization
- B. Increases the number of dimensions
- C. Invokes early stopping
- D. Stops the layers from collapsing back into just a linear model



Our favorite non-linearity is the
Rectified Linear Unit



$$f(x) = \max(0, x)$$

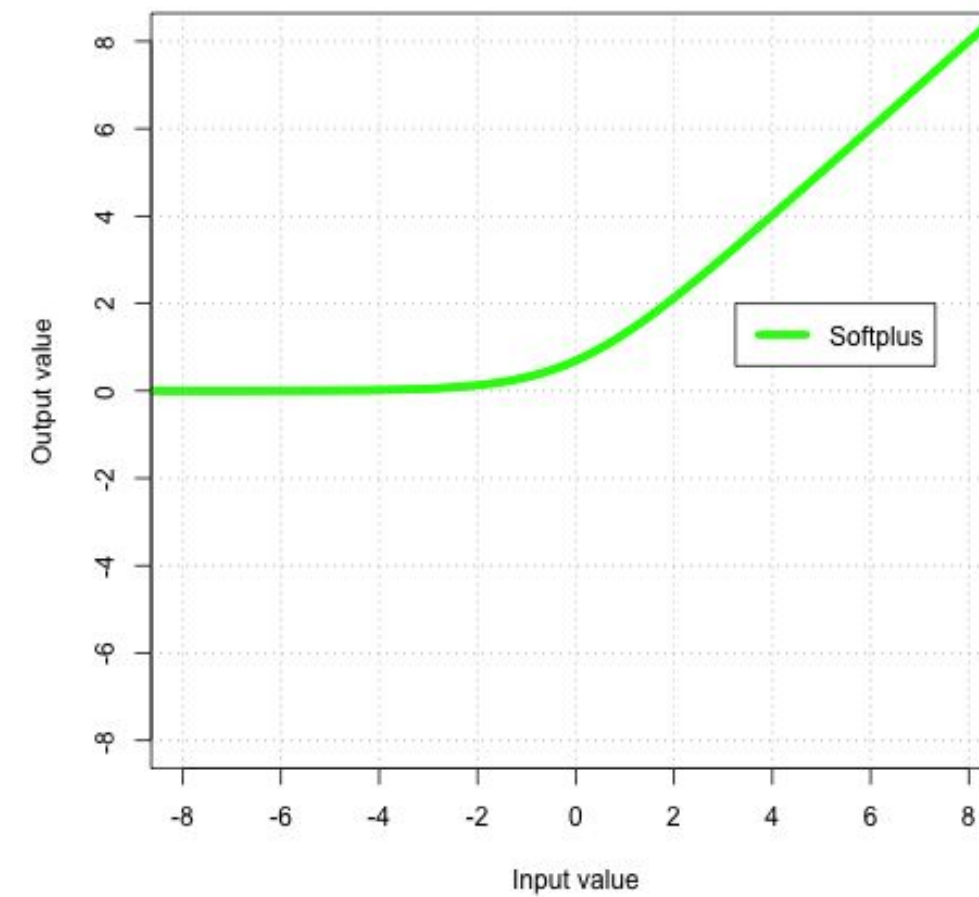

There are many different ReLU variants

$$\text{ReLU} = f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$$



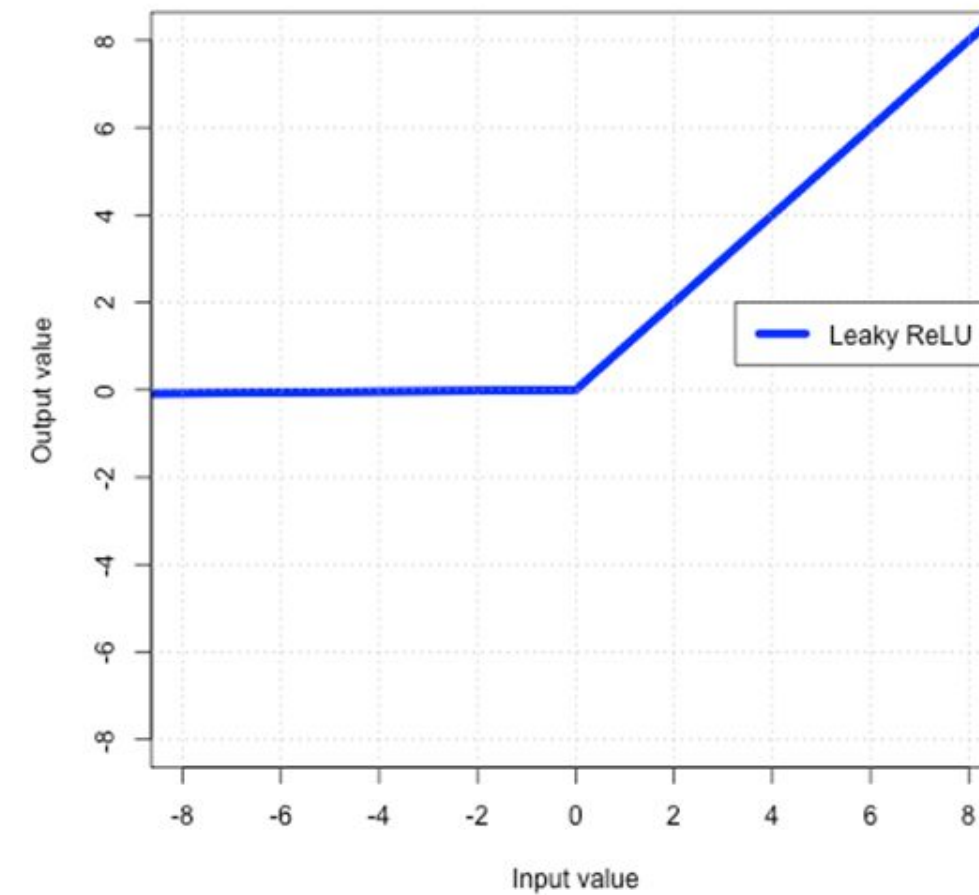

There are many different ReLU variants

$$\text{Softplus} = \ln(1 + e^x)$$





There are many different ReLU variants

$$\text{Leaky ReLU} = f(x) = \begin{cases} 0.01x & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$$



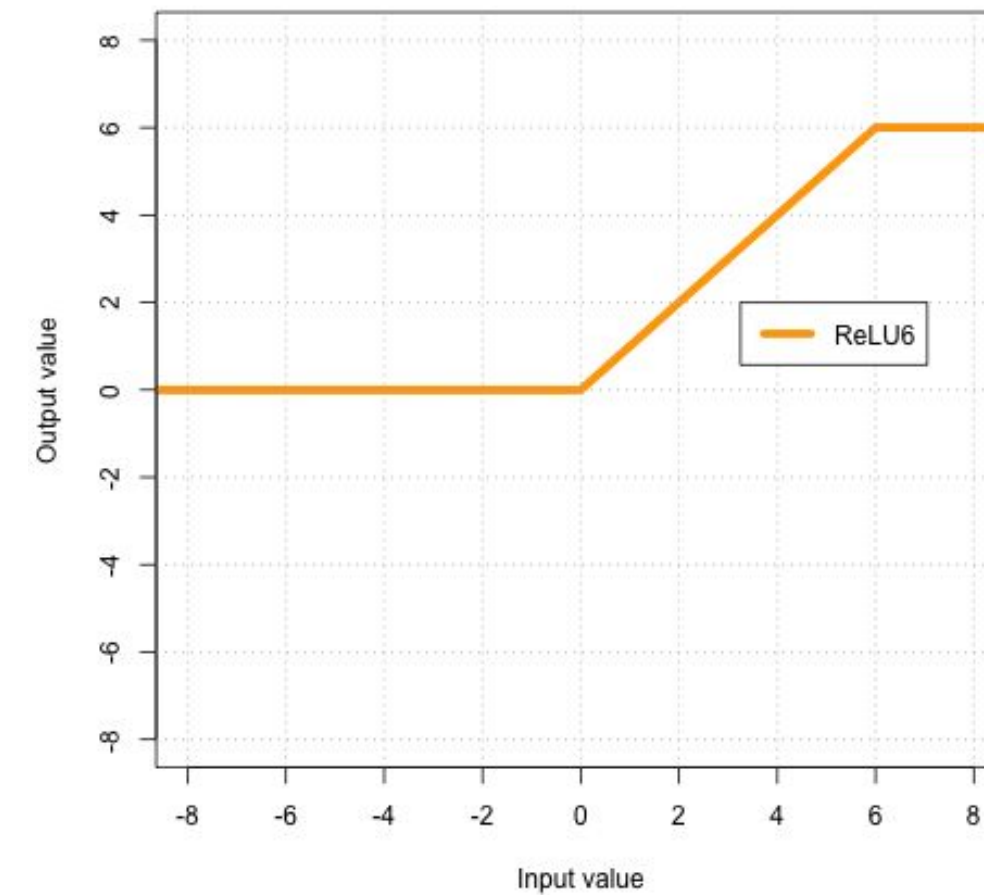
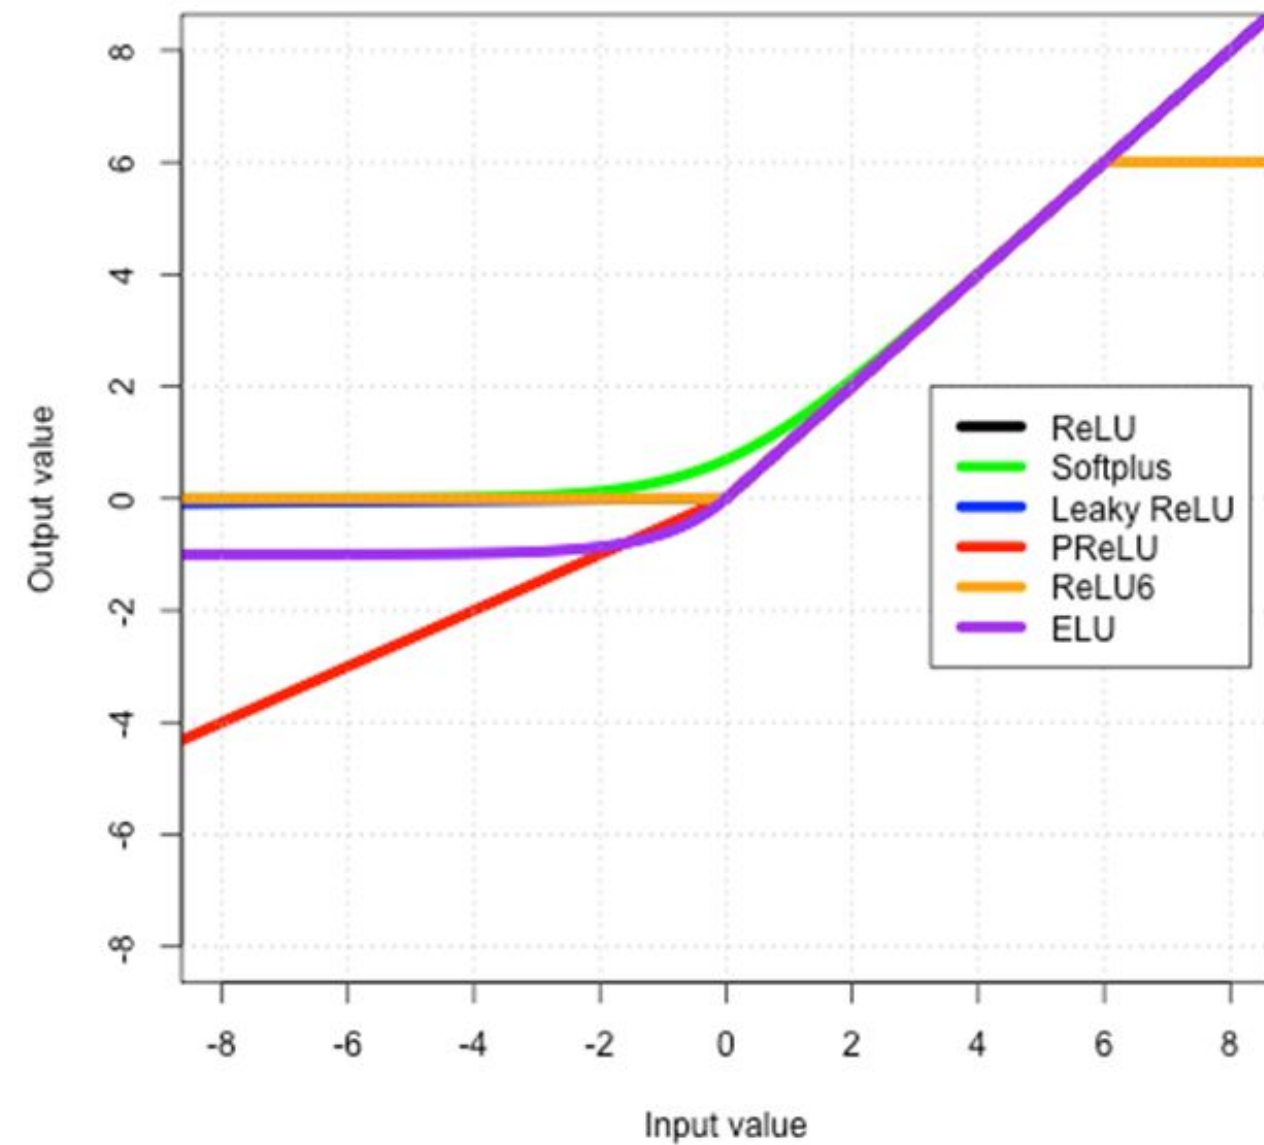
There are many different ReLU variants

$$PReLU = f(x) = \begin{cases} \alpha x & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$$



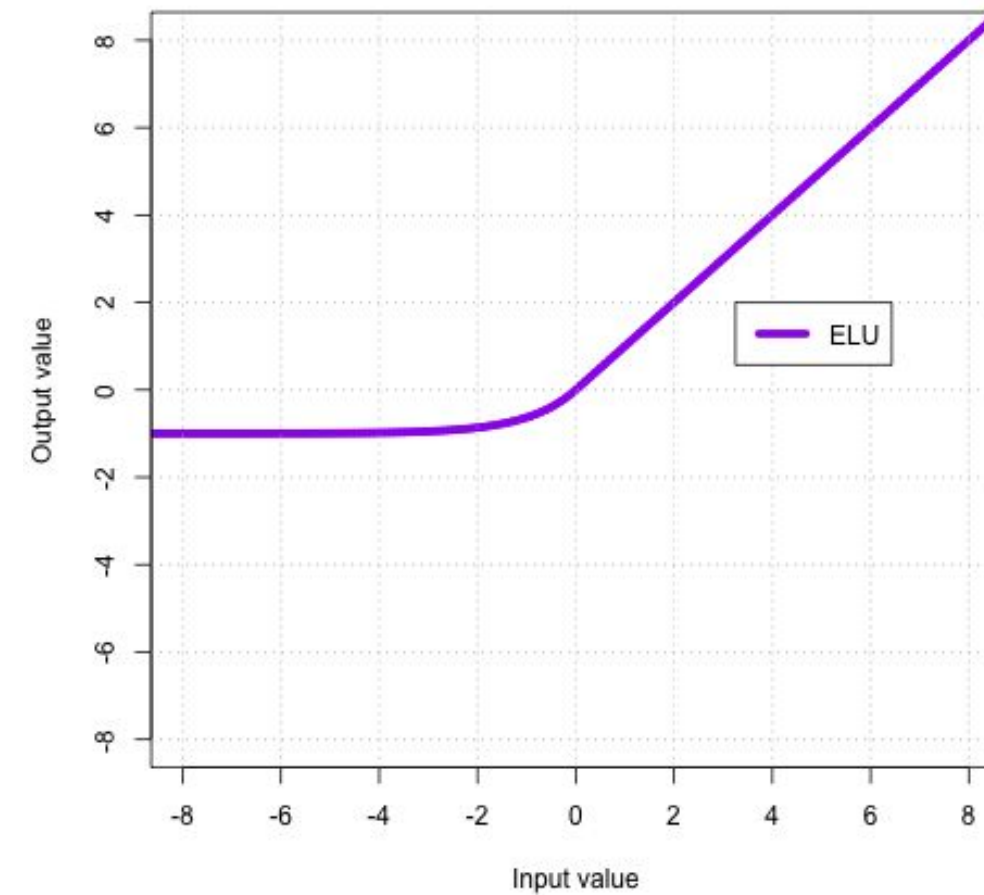

There are many different ReLU variants

$$\text{ReLU6} = \min(\max(0, x), 6)$$




There are many different ReLU variants

$$ELU = f(x) = \begin{cases} \alpha(e^x - 1) & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$$




Neural Nets Can Be Arbitrarily Complex



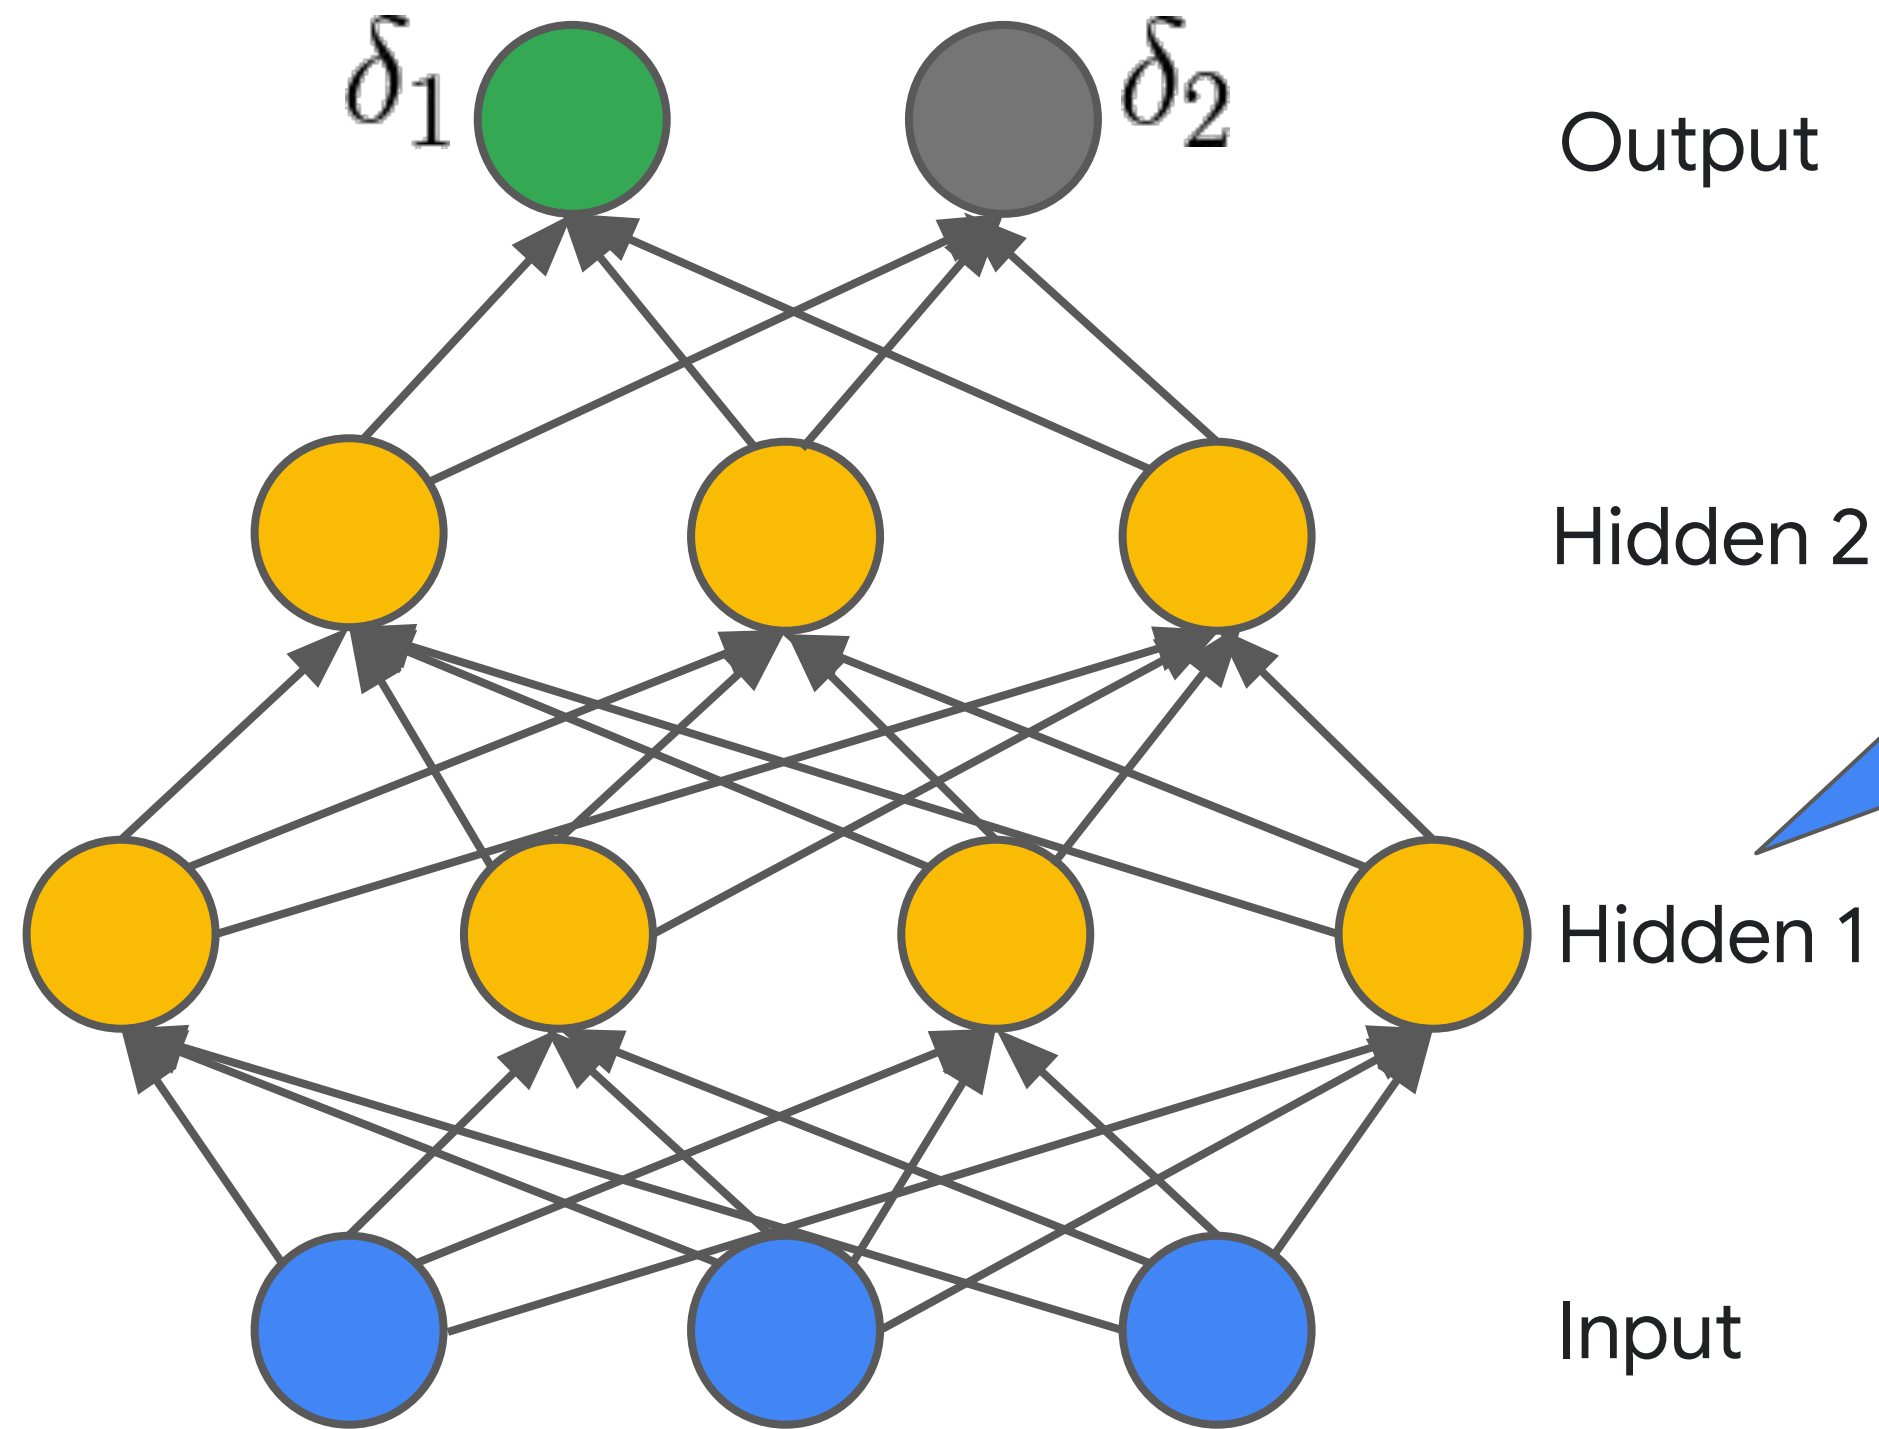
Training done via
BackProp algorithm:
gradient descent in very
non-convex space

Neural Nets Can Be Arbitrarily Complex



Training done via
BackProp algorithm:
gradient descent in very
non-convex space

Neural Nets Can Be Arbitrarily Complex



Training done via
BackProp algorithm:
gradient descent in very
non-convex space

Neural Network Complexity Quiz

Neural networks can be arbitrarily complex. To increase hidden dimensions, I can add _____. To increase function composition, I can add _____. If I have multiple labels per example, I can add _____.

- A. Layers, neurons, outputs
- B. Neurons, layers, outputs
- C. Layers, outputs, neurons
- D. Neurons, outputs, layers

Neural Network Complexity Quiz

Neural networks can be arbitrarily complex. To increase hidden dimensions, I can add _____. To increase function composition, I can add _____. If I have multiple labels per example, I can add _____.

- A. Layers, neurons, outputs
- ☒ B. Neurons, layers, outputs
- C. Layers, outputs, neurons
- D. Neurons, outputs, layers

Lab

Neural Networks playground

Lab: Neural Networks playground

Solve these problems in two ways:
one by feature engineering, the
other by adding layers:

<https://goo.gl/2eig4q>

<https://goo.gl/wXbGDW>

<https://goo.gl/i9r55D>

Lab: Neural Networks playground

Camtasia



Training neural networks

DNNRegressor usage is similar to LinearRegressor

```
myopt = tf.train.AdamOptimizer(learning_rate=0.01)

model = tf.estimator.DNNRegressor(model_dir=outdir,
                                   hidden_units=[100, 50, 20],
                                   feature_columns=INPUT_COLS,
                                   optimizer=myopt,
                                   dropout=0.1)

NSTEPS = (100 * len(traindf)) / BATCH_SIZE
model.train(input_fn=train_input_fn, steps=NSTEPS)
```

Use momentum-based optimizers e.g. Adagrad (the default) or Adam.

Specify number of hidden nodes.

Optionally, can also regularize using dropout

Three common failure modes for gradient descent

Gradients can vanish

Each additional layer
can successively
reduce signal vs. noise

Using ReLu instead of
sigmoid/tanh can help

Problem

Insight

Solution

Three common failure modes for gradient descent

Gradients can vanish

Each additional layer can successively reduce signal vs. noise

Using ReLu instead of sigmoid/tanh can help

Gradients can explode

Learning rates are important here

Batch normalization (useful knob) can help

Problem

Insight

Solution

Three common failure modes for gradient descent

Gradients can vanish

Each additional layer can successively reduce signal vs. noise

Using ReLu instead of sigmoid/tanh can help

Gradients can explode

Learning rates are important here


Batch normalization (useful knob) can help

ReLu layers can die Problem

Monitor fraction of zero weights in TensorBoard Insight


Lower your learning rates Solution

There are benefits if feature values are small numbers




Roughly zero-centered,
[-1, 1] range often
works well

There are benefits if feature values are small numbers




Roughly zero-centered,
 $[-1, 1]$ range often
works well




Small magnitudes help
gradient descent
converge and avoid
NaN trap


There are benefits if feature values are small numbers



Roughly zero-centered,
[-1, 1] range often
works well




Small magnitudes help
gradient descent
converge and avoid
NaN trap




Avoiding outlier
values helps with
generalization

We can use standard methods to make feature values scale to small numbers




Linear scaling

We can use standard methods to make feature values scale to small numbers




Linear scaling




Hard cap (clipping) to
max, min


We can use standard methods to make feature values scale to small numbers



Linear scaling



Hard cap (clipping) to
max, min



Log scaling

Gradient Descent Debugging Quiz

Which of these is good advice if my model is experiencing exploding gradients?

- A. Lower the learning rate
- B. Add weight regularization
- C. Add gradient clipping
- D. Add batch normalization
- E. See a doctor
- F. C,D
- G. A,C,D
- H. A,B,C,D

Gradient Descent Debugging Quiz

Which of these is good advice if my model is experiencing exploding gradients?


- A. Lower the learning rate
- B. Add weight regularization
- C. Add gradient clipping
- D. Add batch normalization
- E. See a doctor
- F. C,D
- G. A,C,D
- H. A,B,C,D

Dropout layers are a form of regularization


Dropout works by randomly “dropping out” unit activations in a network for a single gradient step

During training only!
In prediction all nodes are kept

Helps learn “multiple paths” --
think: ensemble models,
random forests



Dropout simulates ensemble learning




The more you drop out, the stronger the regularization

0.0 = no dropout
regularization

Intermediate values
more useful, a value of
dropout=0.2 is typical

1.0 = drop everything
out! learns nothing

0.0



1.0

Dropout Quiz

Dropout acts as another form of _____. It forces data to flow down _____ paths so that there is a more even spread. It also simulates _____ learning. Don't forget to scale the dropout activations by the inverse of the _____. We remove dropout during _____.

- A. Hyperparameter tuning, similar, deep, drop probability, training
- B. Hyperparameter tuning, multiple, deep, drop probability, inference
- C. Regularization, multiple, ensemble, keep probability, training
- D. Regularization, multiple, ensemble, drop probability, inference
- E. Regularization, multiple, ensemble, keep probability, inference
- F. Hyperparameter tuning, multiple, deep, keep probability, inference
- G. Regularization, similar, ensemble, keep probability, inference

Dropout Quiz

Dropout acts as another form of _____. It forces data to flow down _____ paths so that there is a more even spread. It also simulates _____ learning. Don't forget to scale the dropout activations by the inverse of the _____. We remove dropout during _____.

- A. Hyperparameter tuning, similar, deep, drop probability, training
- B. Hyperparameter tuning, multiple, deep, drop probability, inference
- C. Regularization, multiple, ensemble, keep probability, training
- D. Regularization, multiple, ensemble, drop probability, inference
- E. Regularization, multiple, ensemble, keep probability, inference
- F. Hyperparameter tuning, multiple, deep, keep probability, inference
- G. Regularization, similar, ensemble, keep probability, inference

Lab

Using Neural Networks to
build a ML model

Ryan Gillard

Lab: Using Neural Networks to build ML model

In this lab, you will use the DNNRegressor class in TensorFlow to predict median housing price

1 The data is based on 1990 census data from California.

2 This data is at the city block level, so these features reflect the total number of rooms in that block, or the total number of people who live on that block, respectively

Lab: Using Neural Networks to build ML model


Camtasia



Multi-class neural nets

Ryan Gillard

Logistic regression provides useful probabilities for binary-class problems



There are lots of multi-class problems



Pit

Stalls

Circle

Suite

How do we extend the logits idea to multi-class classifiers?

There are lots of multi-class problems

Model 1	
<u>Positive Class</u>	<u>Negative Class</u>
Pit	Other

Model 2	
<u>Positive Class</u>	<u>Negative Class</u>
Stalls	Other

Model 3	
<u>Positive Class</u>	<u>Negative Class</u>
Circle	Other

Model 4	
<u>Positive Class</u>	<u>Negative Class</u>
Suite	Other

There are lots of multi-class problems

Model 1	
<u>Positive Class</u>	<u>Negative Class</u>
Pit	Stalls

Model 2	
<u>Positive Class</u>	<u>Negative Class</u>
Pit	Circle


Model 3	
<u>Positive Class</u>	<u>Negative Class</u>
Pit	Suite

Model 4	
<u>Positive Class</u>	<u>Negative Class</u>
Stalls	Circle

Model 5	
<u>Positive Class</u>	<u>Negative Class</u>
Stalls	Suite


Model 6	
<u>Positive Class</u>	<u>Negative Class</u>
Circle	Suite

Idea: Use separate output nodes for each possible class



Add additional constraint, that total of outputs = 1.0

$$p(y = j|\mathbf{x}) = \frac{\exp(\mathbf{w}_j^T \mathbf{x} + b_j)}{\sum_{k \in K} \exp(\mathbf{w}_k^T \mathbf{x} + b_k)}$$



Use one softmax loss for all possible classes

```
logits = tf.matmul(...) # logits for each output node
                        -> shape = [batch_size, num_classes]
labels = ...           # one-hot encoding in [0, num_classes)
                        -> shape = [batch_size, num_classes]

loss = tf.reduce_mean(
    tf.nn.softmax_cross_entropy_with_logits_v2(
        logits, labels) -> shape = [batch_size]
)
```

Use one softmax loss for all possible classes

```
logits = tf.matmul(...) # logits for each output node
                        -> shape = [batch_size, num_classes]
labels = ...           # index in [0, num_classes)
                        -> shape = [batch_size]

loss = tf.reduce_mean(
    tf.nn.sparse_softmax_cross_entropy_with_logits(
        logits, labels) -> shape = [batch_size]
)
```

Use softmax only when classes are mutually exclusive

“Multi-Class, Single-Label Classification”


An example may be a member of only one class.

Are there multi-class settings where examples may belong to more than one class?

```
tf.nn.sigmoid_cross_entropy_with_logits(  
    logits, labels) -> shape = [batch_size, num_classes]
```

If you have hundreds or thousands of classes, loss computation can become a significant bottleneck

Need to evaluate every output node for every example



Approximate versions of softmax exist



Candidate Sampling calculates for all the positive labels, but only for a random sample of negatives: `tf.nn.sampled_softmax_loss`

Noise-contrastive approximates the denominator of softmax by modeling the distribution of outputs: `tf.nn.nce_loss`

Softmax Quiz

For our classification output, if we have both mutually exclusive labels and probabilities, we should use _____. If the labels are mutually exclusive, but the probabilities aren't, we should use _____. If our labels aren't mutually exclusive, we should use _____.

- I. `tf.nn.sigmoid_cross_entropy_with_logits`
- II. `tf.nn.sparse_softmax_cross_entropy_with_logits`
- III. `tf.nn.softmax_cross_entropy_with_logits_v2`

- A. III, II, I
- B. I, II, III
- C. III, I, II
- D. II, III, I

Softmax Quiz

For our classification output, if we have both mutually exclusive labels and probabilities, we should use _____. If the labels are mutually exclusive, but the probabilities aren't, we should use _____. If our labels aren't mutually exclusive, we should use _____.

- I. `tf.nn.sigmoid_cross_entropy_with_logits`
- II. `tf.nn.sparse_softmax_cross_entropy_with_logits`
- III. `tf.nn.softmax_cross_entropy_with_logits_v2`

- A. III, II, I
- B. I, II, III
- C. III, I, II
- D. II, III, I

cloud.google.com