



Embeddings

Lak Lakshmanan



Machine Learning on Google Cloud Platform

The Art of ML

Hyperparameter Tuning

A Pinch of Science

The Science of Neural Networks

Embeddings

Custom Estimator



Learn how to...

Use embeddings to:

Manage sparse data



Learn how to...

Use embeddings to:

- Manage sparse data

- Reduce dimensionality



Learn how to...

Use embeddings to:

- Manage sparse data

- Reduce dimensionality

- Increase model generalization



Learn how to...

Use embeddings to:

- Manage sparse data

- Reduce dimensionality

- Increase model generalization

- Cluster observations



Learn how to...

Use embeddings to:

- Manage sparse data
- Reduce dimensionality
- Increase model generalization
- Cluster observations

Create reusable embeddings



Learn how to...

Use embeddings to:

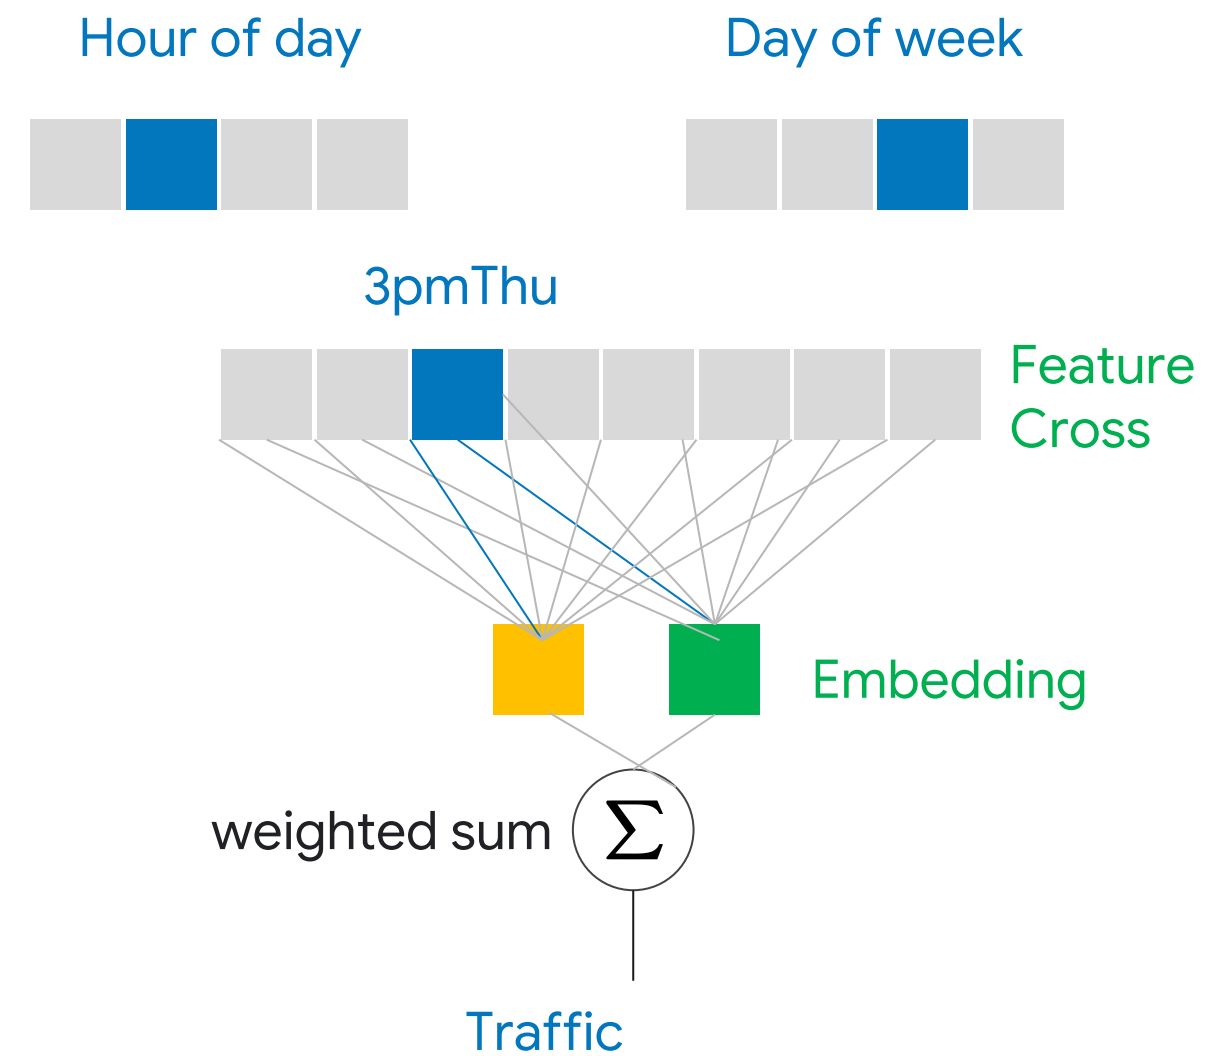
- Manage sparse data
- Reduce dimensionality
- Increase model generalization
- Cluster observations

Create reusable embeddings

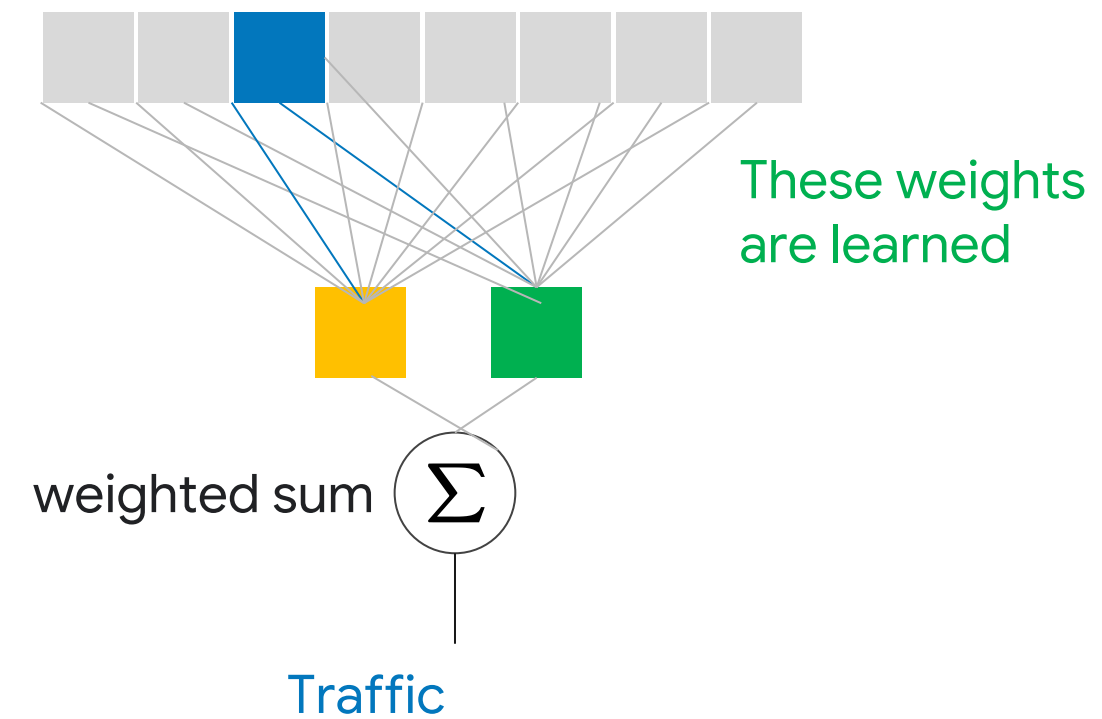
Explore embeddings in
TensorBoard




Creating an embedding column from a feature cross



The weights in the embedding column are learned from data



The model learns how to embed the feature cross in lower-dimensional space



8am Tue	0.8	0.7
9am Wed	0.7	0.9
11 am Tue	0.1	0.6
2 pm Wed	0.1	0.7
2 am Tue	0	0.1
2 am Wed	0	0.1

similar

different



Embedding a feature cross in TensorFlow

```
import tf.feature_column as fc

day_hr = fc.crossed_column(
    [dayofweek, hourofday],
    24x7 )

day_hr_em = fc.embedding_column(
    day_hr,
    2 )
```



Transfer Learning of embeddings from similar ML models

```
import tf.feature_column as fc
```

```
day_hr = fc.crossed_column(  
    [dayofweek, hourofday],  
    24x7 )
```

```
day_hr_em = fc.embedding_column(  
    day_hr,  
    2,  
    ckpt_to_load_from='london/*ckpt-1000*',  
    tensor_name_in_ckpt='dayhr_embed',  
    trainable=False  
)
```



Transfer Learning of embeddings from similar ML models

First layer: the feature cross

Second layer: a mystery box labeled latent factor

Third layer: the embedding

Fourth layer: one side: image of traffic

Second side: image of people watching TV





Representing feature columns as sparse vectors

These are all different ways to
create a categorical column

If you know the keys beforehand:

```
tf.feature_column.categorical_column_with_vocabulary_list('employeeId',  
    vocabulary_list = ['8345', '72345', '87654', '98723', '23451']),
```

If your data is already indexed; i.e., has integers in [0-N):

```
tf.feature_column.categorical_column_with_identity('employeeId',  
    num_buckets = 5)
```

If you don't have a vocabulary of all possible values:

```
tf.feature_column.categorical_column_with_hash_bucket('employeeId',  
    hash_bucket_size = 500)
```



How do you recommend movies to customers?



500,000 movies



One approach is to organize movies by similarity (1D)

Average age of viewers



Shrek



Incredibles



The Triplets
of Belleville



Harry Potter




Star Wars



Bleu



The Dark
Knight Rises



Memento



Using a second dimension gives us more freedom in organizing movies by similarity



Shrek



Incredibles



Harry Potter



Star Wars



The Dark
Knight Rises

The Triplets
of Belleville



Bleu




Memento

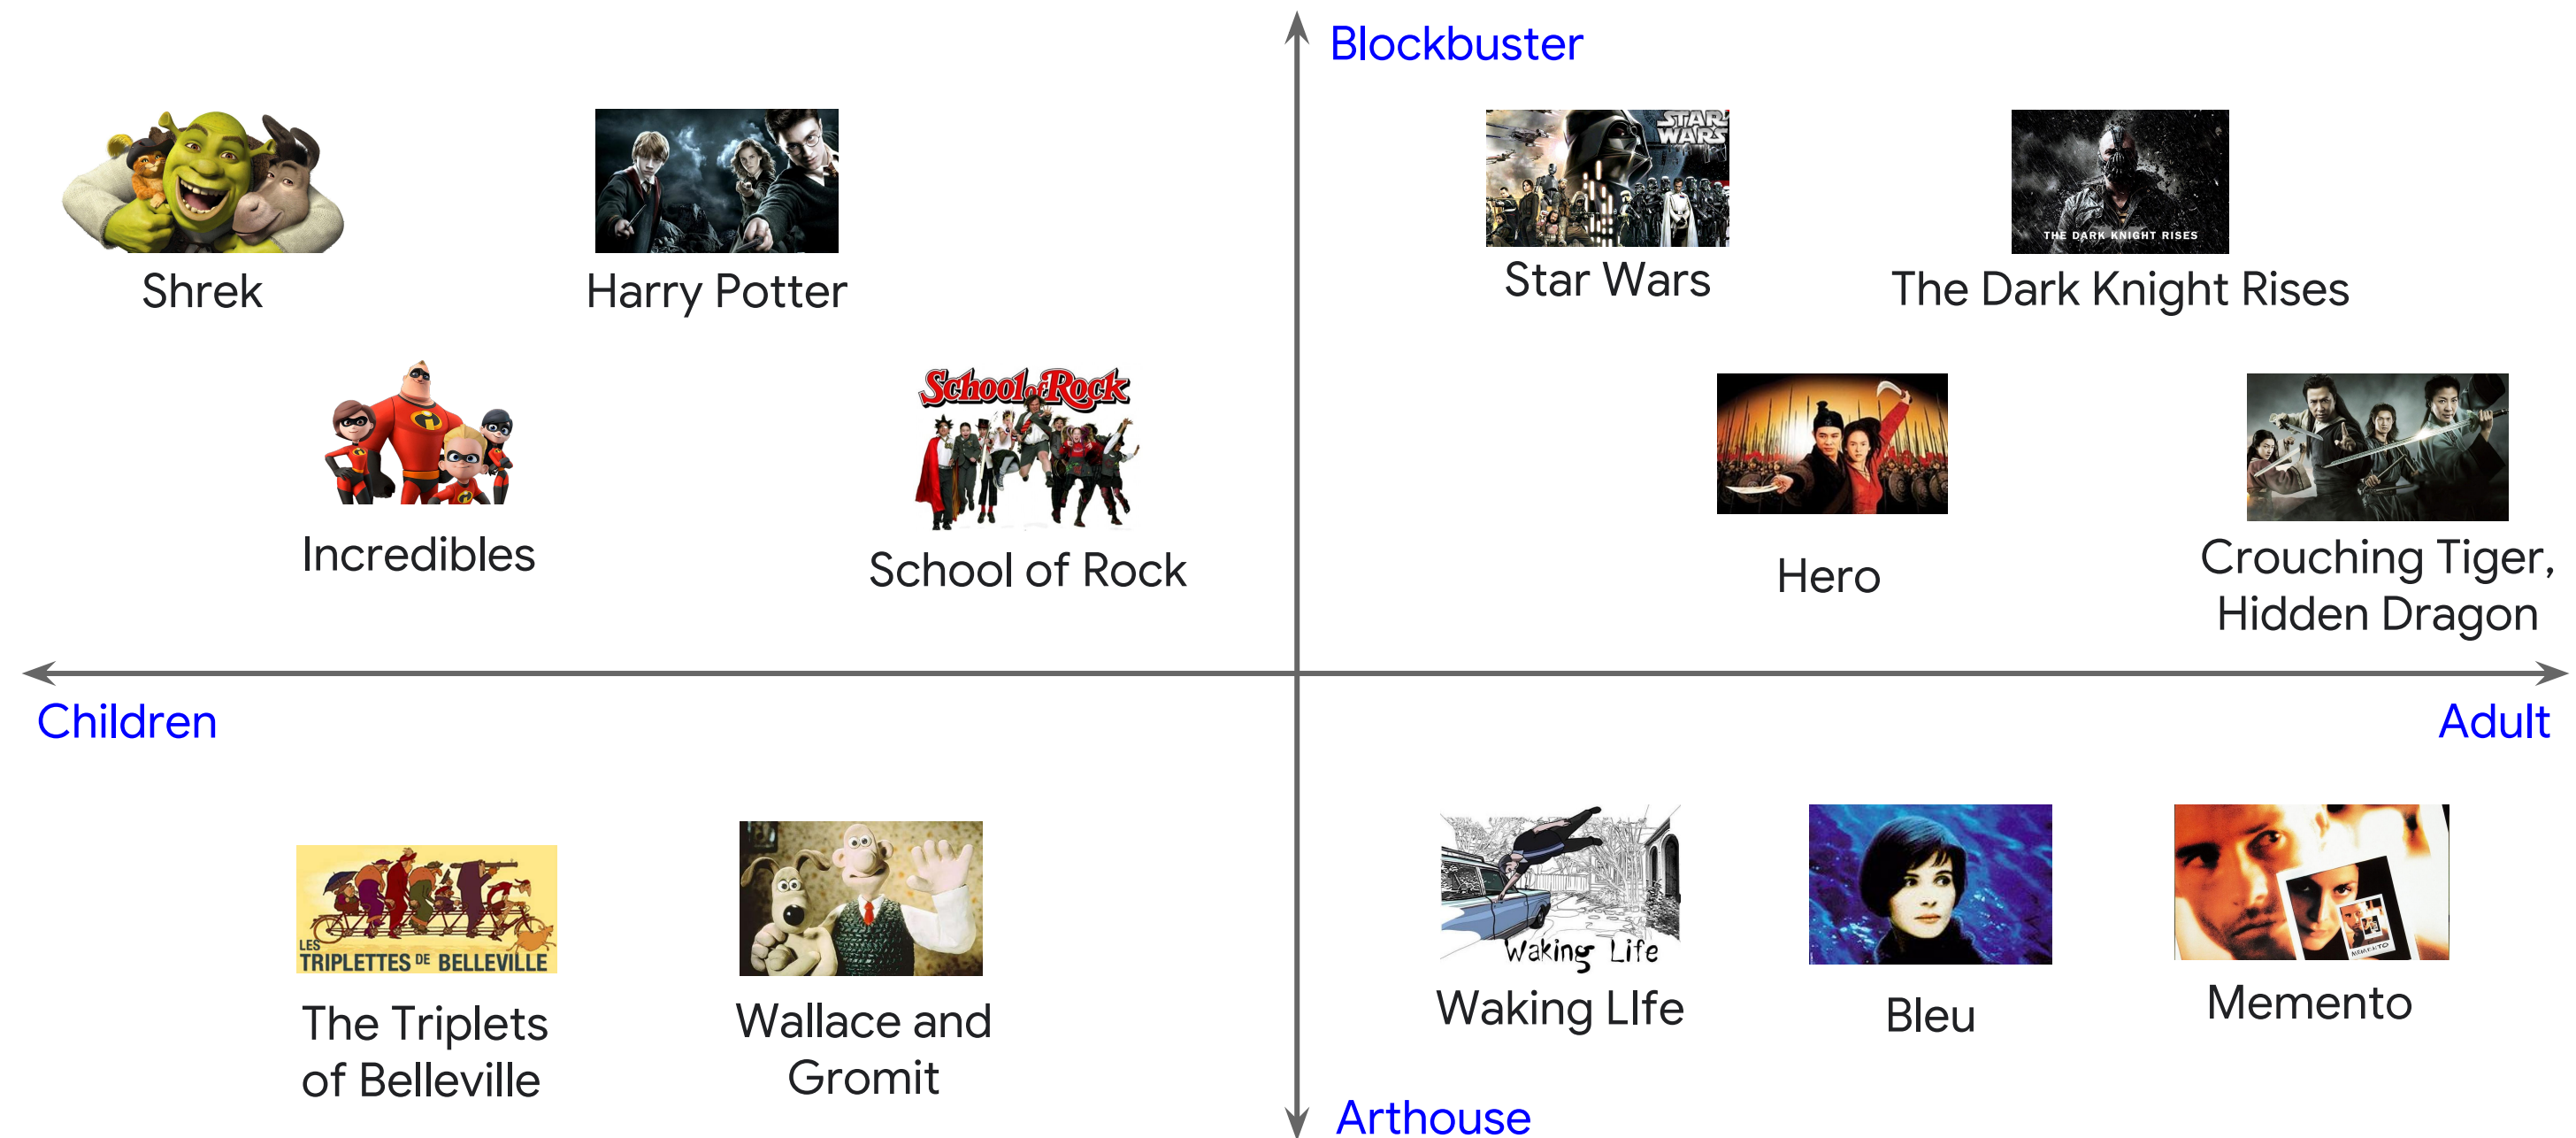
Gross ticket sales



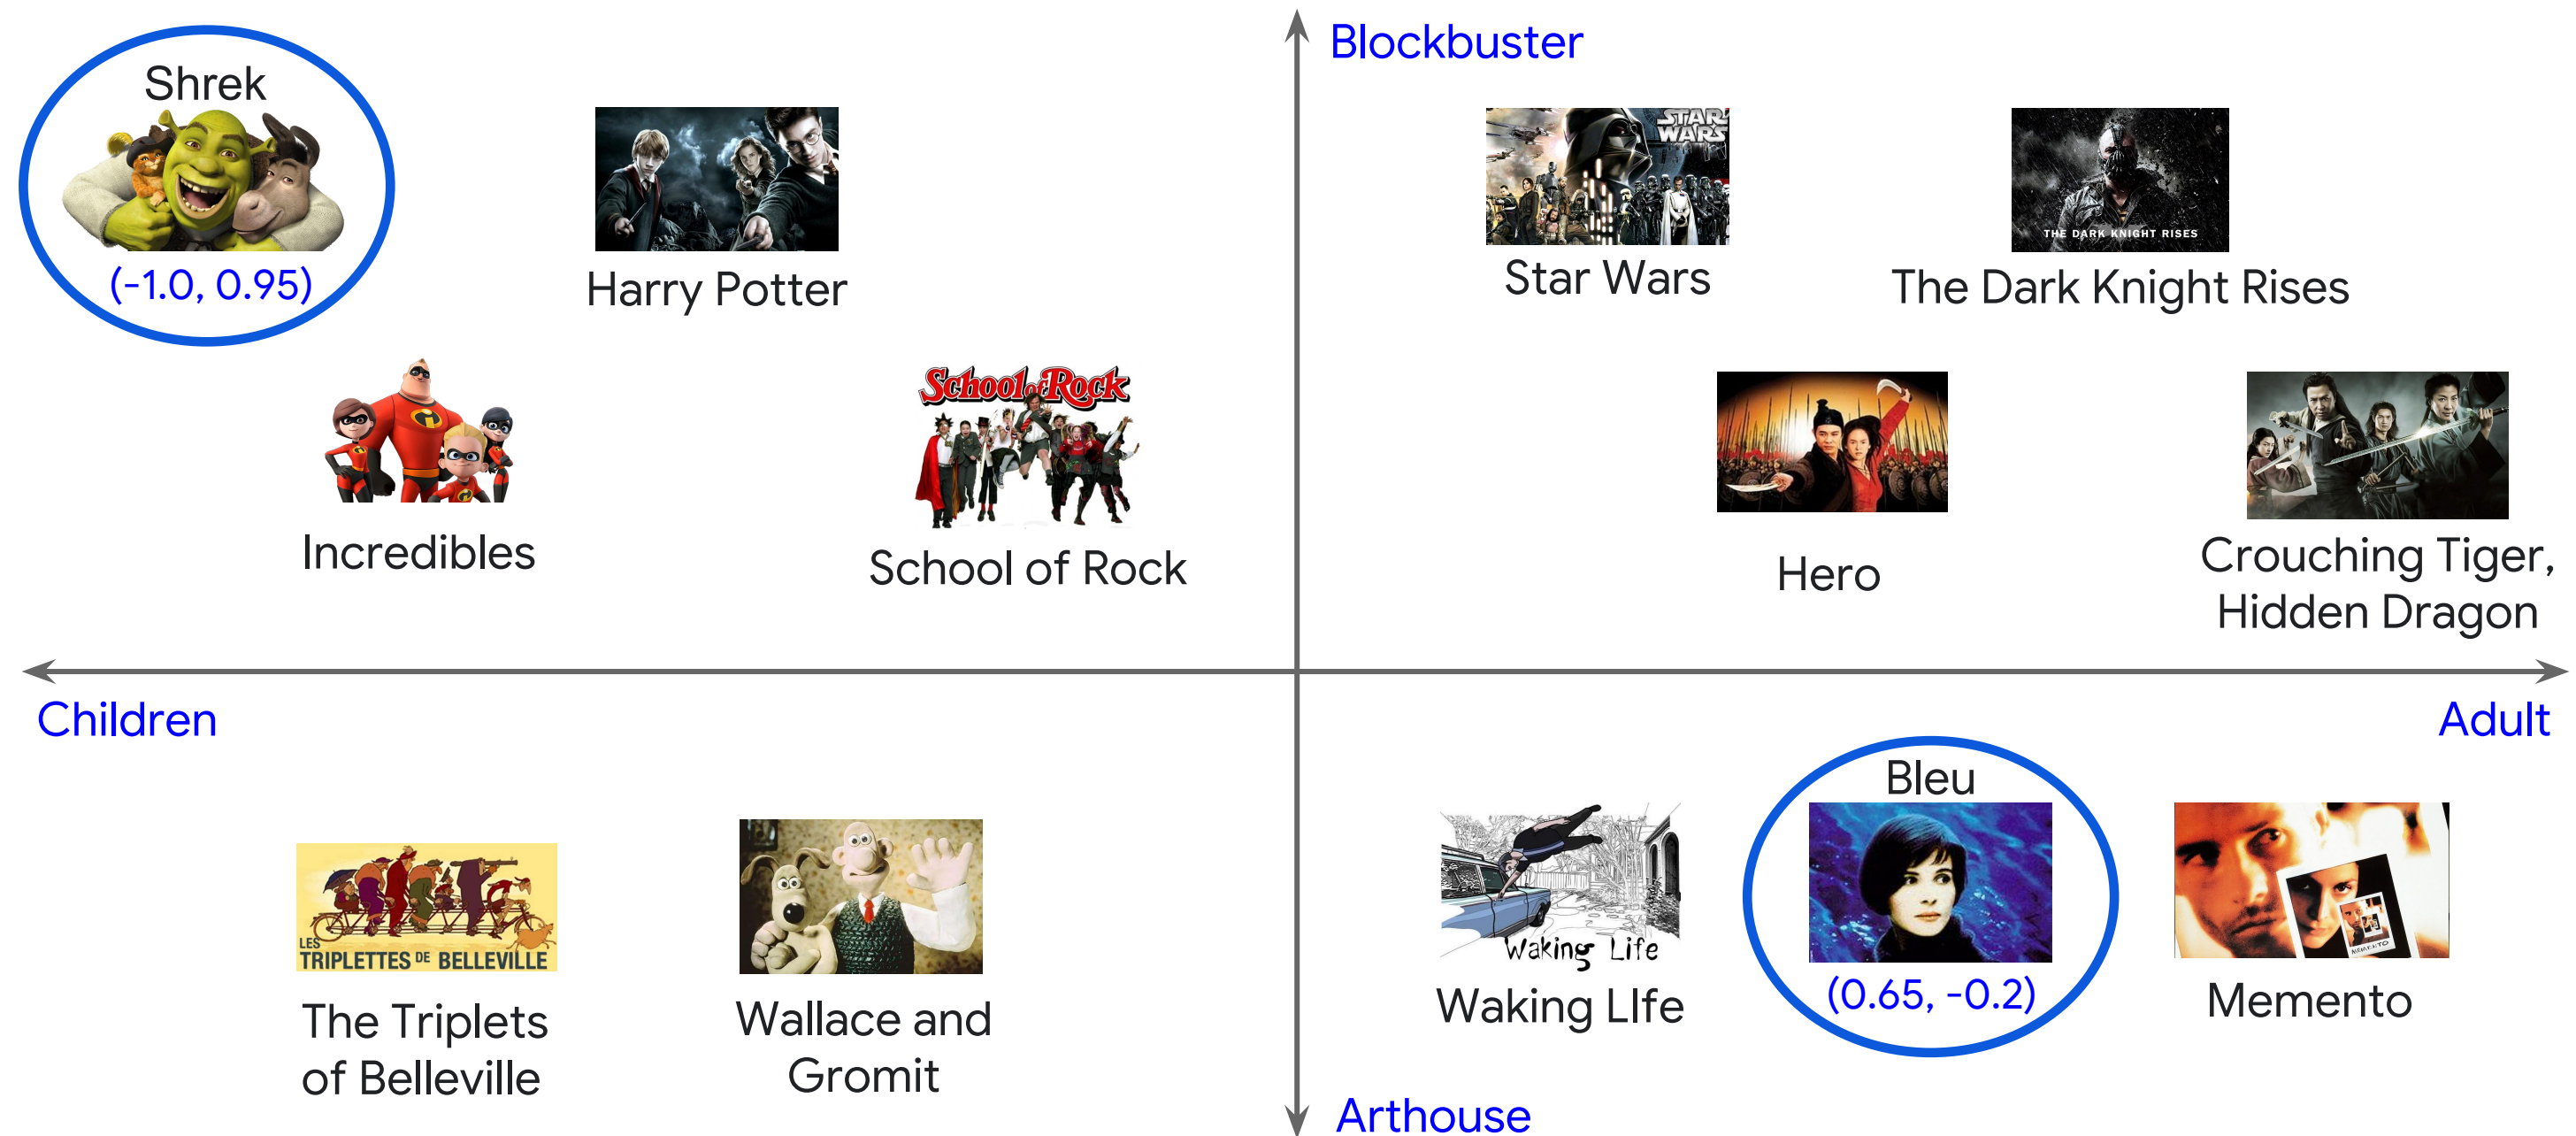
A d -dimensional embedding assumes that user interest in movies can be approximated by d aspects




We could give the axes names, but it is not essential ...



The coordinates are called the 2D embedding for the movie



It's easier to train a model with d inputs than
a model with N inputs





Embeddings can be learned from data

Example	movieId
0	
1	
2	
3	
4	



#	Shrek	Incredible	Triplets ...	Harry Potter	Star Wars
0	1	0	0	0	0
1	0	1	0	0	0
2	0	0	1	0	0
3	0	0	0	1	0
4	1	0	0	0	0



Represent input as sparse vector



4-dimensional embedding
implies we represent movies
with 4 latent factors

Dense representations are inefficient in space and compute

Dense
representation

(0, 1, 0, 1, 0, ..., 0, 1)




				...			
✓					✓	✓	✓
✓	✓		✓			✓	
		✓			✓		✓
	✓		✓				✓



So, use a sparse representation to hold the example

Build a dictionary mapping each feature to an integer from 0, ..., # movies - 1

Efficiently represent the sparse vector as just the movies the user watched:



Represented as: (1, 3, 999999)

0	1	2	3	...	999999
				...	
✓					✓
✓	✓		✓		✓
		✓			✓
	✓		✓		
					✓





Code to create an embedded feature column in TensorFlow

Example	movieId
0	
1	
2	
3	
4	



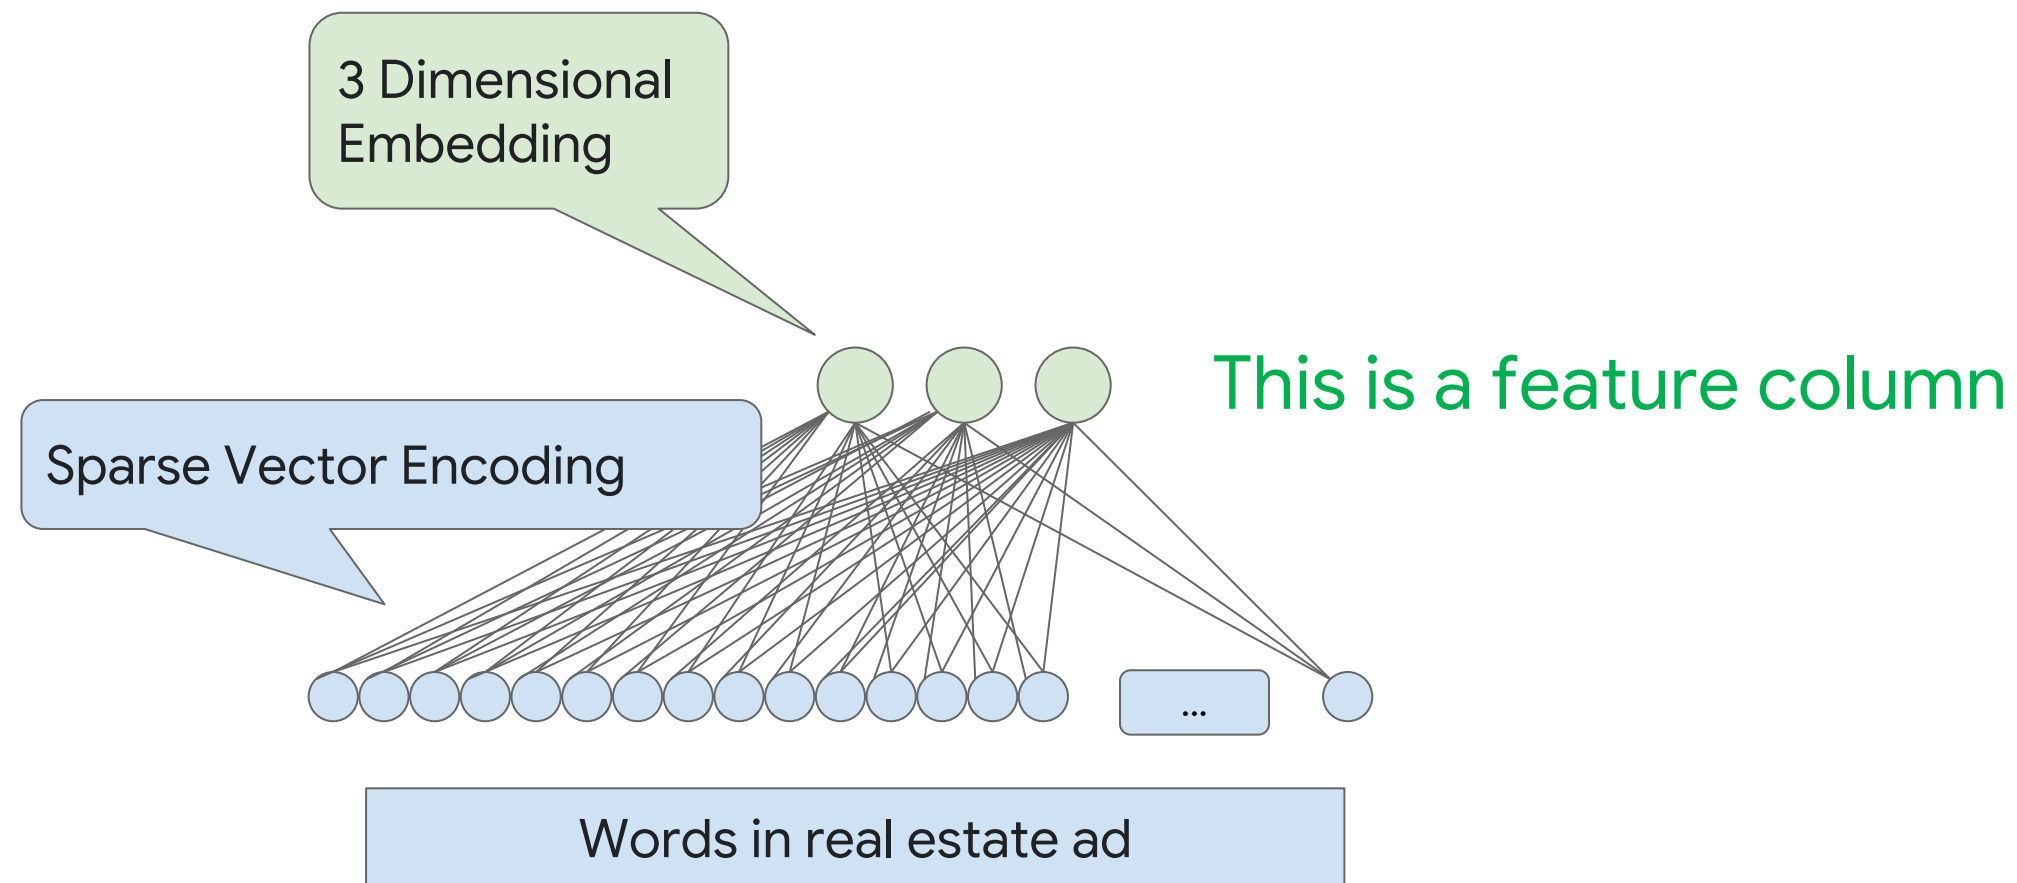
#	Shrek	Incredible	Triplets ...	Harry Potter	Star Wars
0	1	0	0	0	0
1	0	1	0	0	0
2	0	0	1	0	0
3	0	0	0	1	0
4	1	0	0	0	0



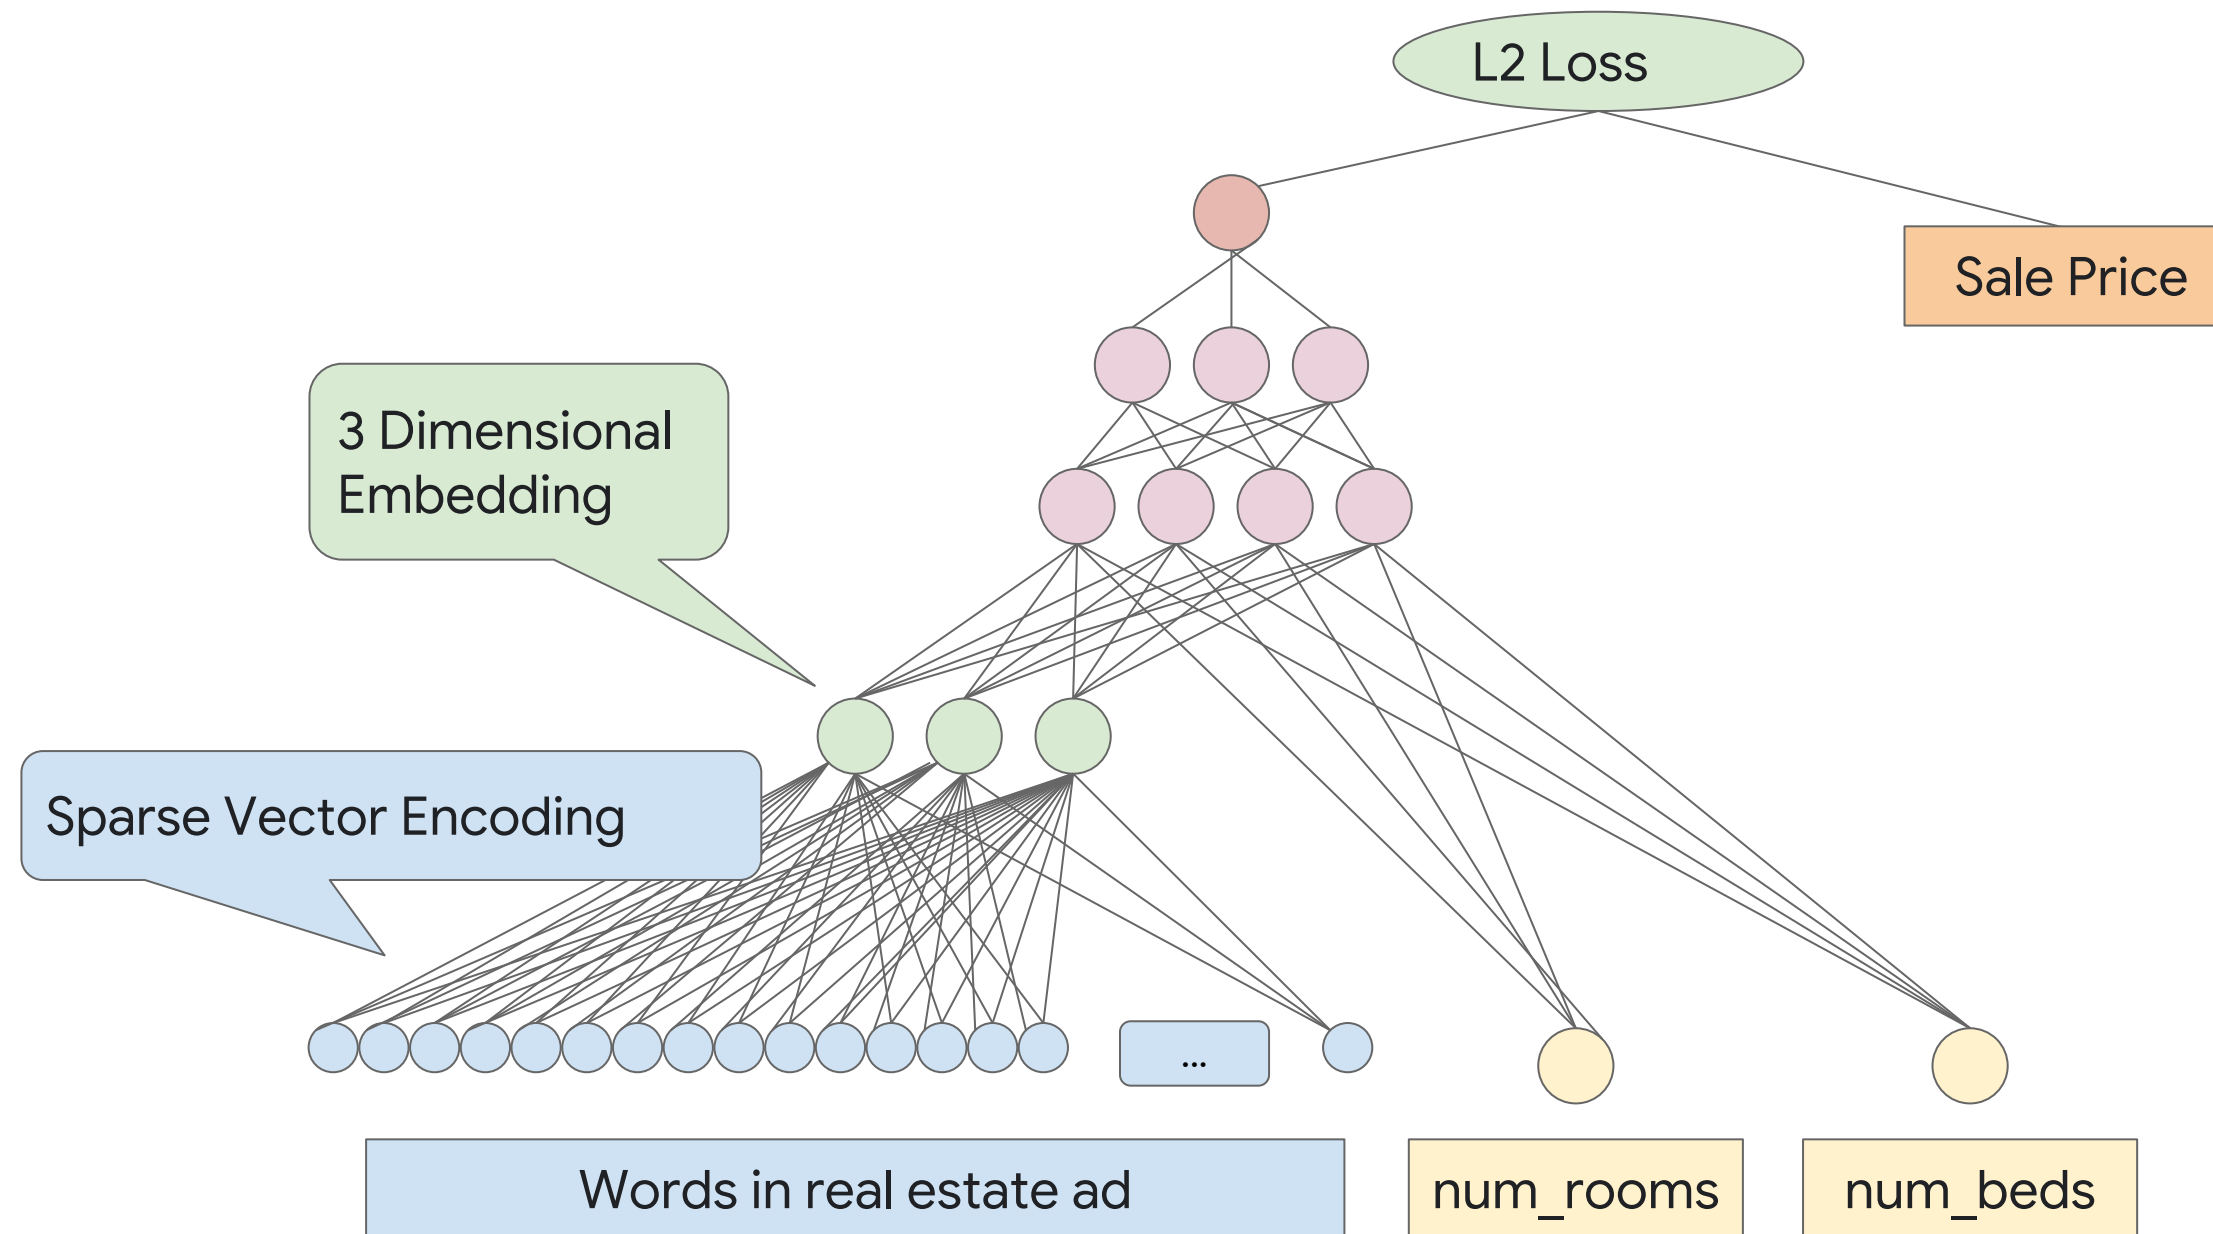
```
sparse_movie = layers.sparse_column_with_keys('movieId',  
                                              keys=[...])  
embedded_movie = layers.embedding_column(sparse_movie, 100)
```

Embeddings are feature columns that function like layers

```
sparse_word = fc.categorical_column_with_vocabulary_list('word', vocabulary_list=englishWords)
embedded_word = fc.embedding_column(sparse_word, 3)
```




The weights in the embedding layer are learned through backprop just as with other weights




Embeddings can be thought of as latent features


Deep Network



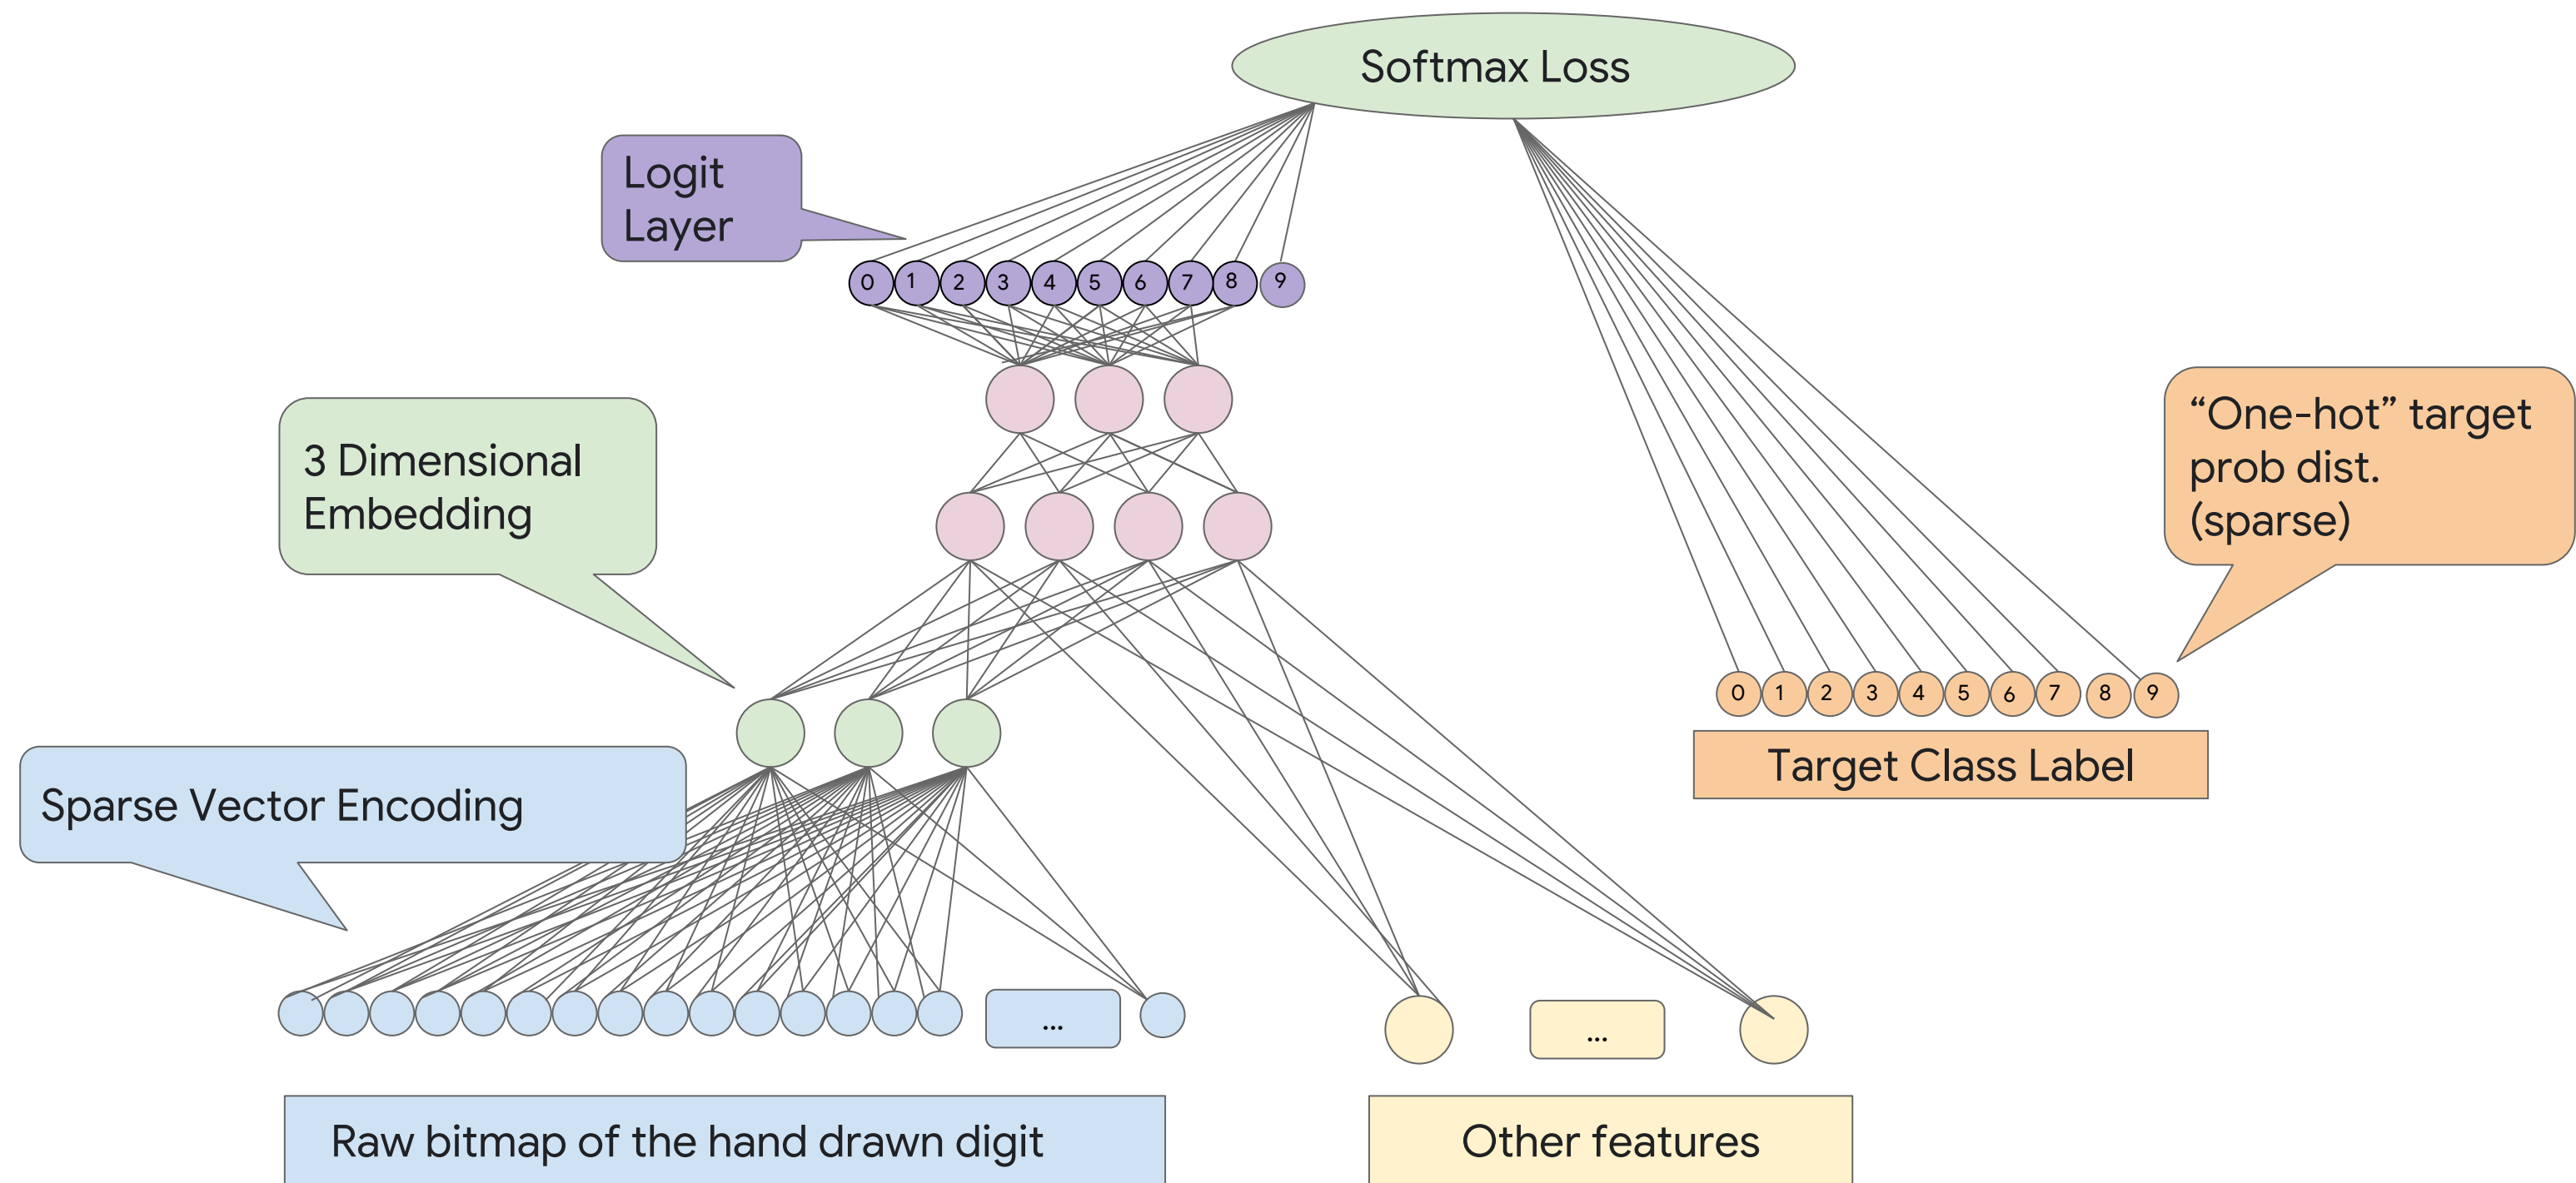
Geometric view of a single movie embedding



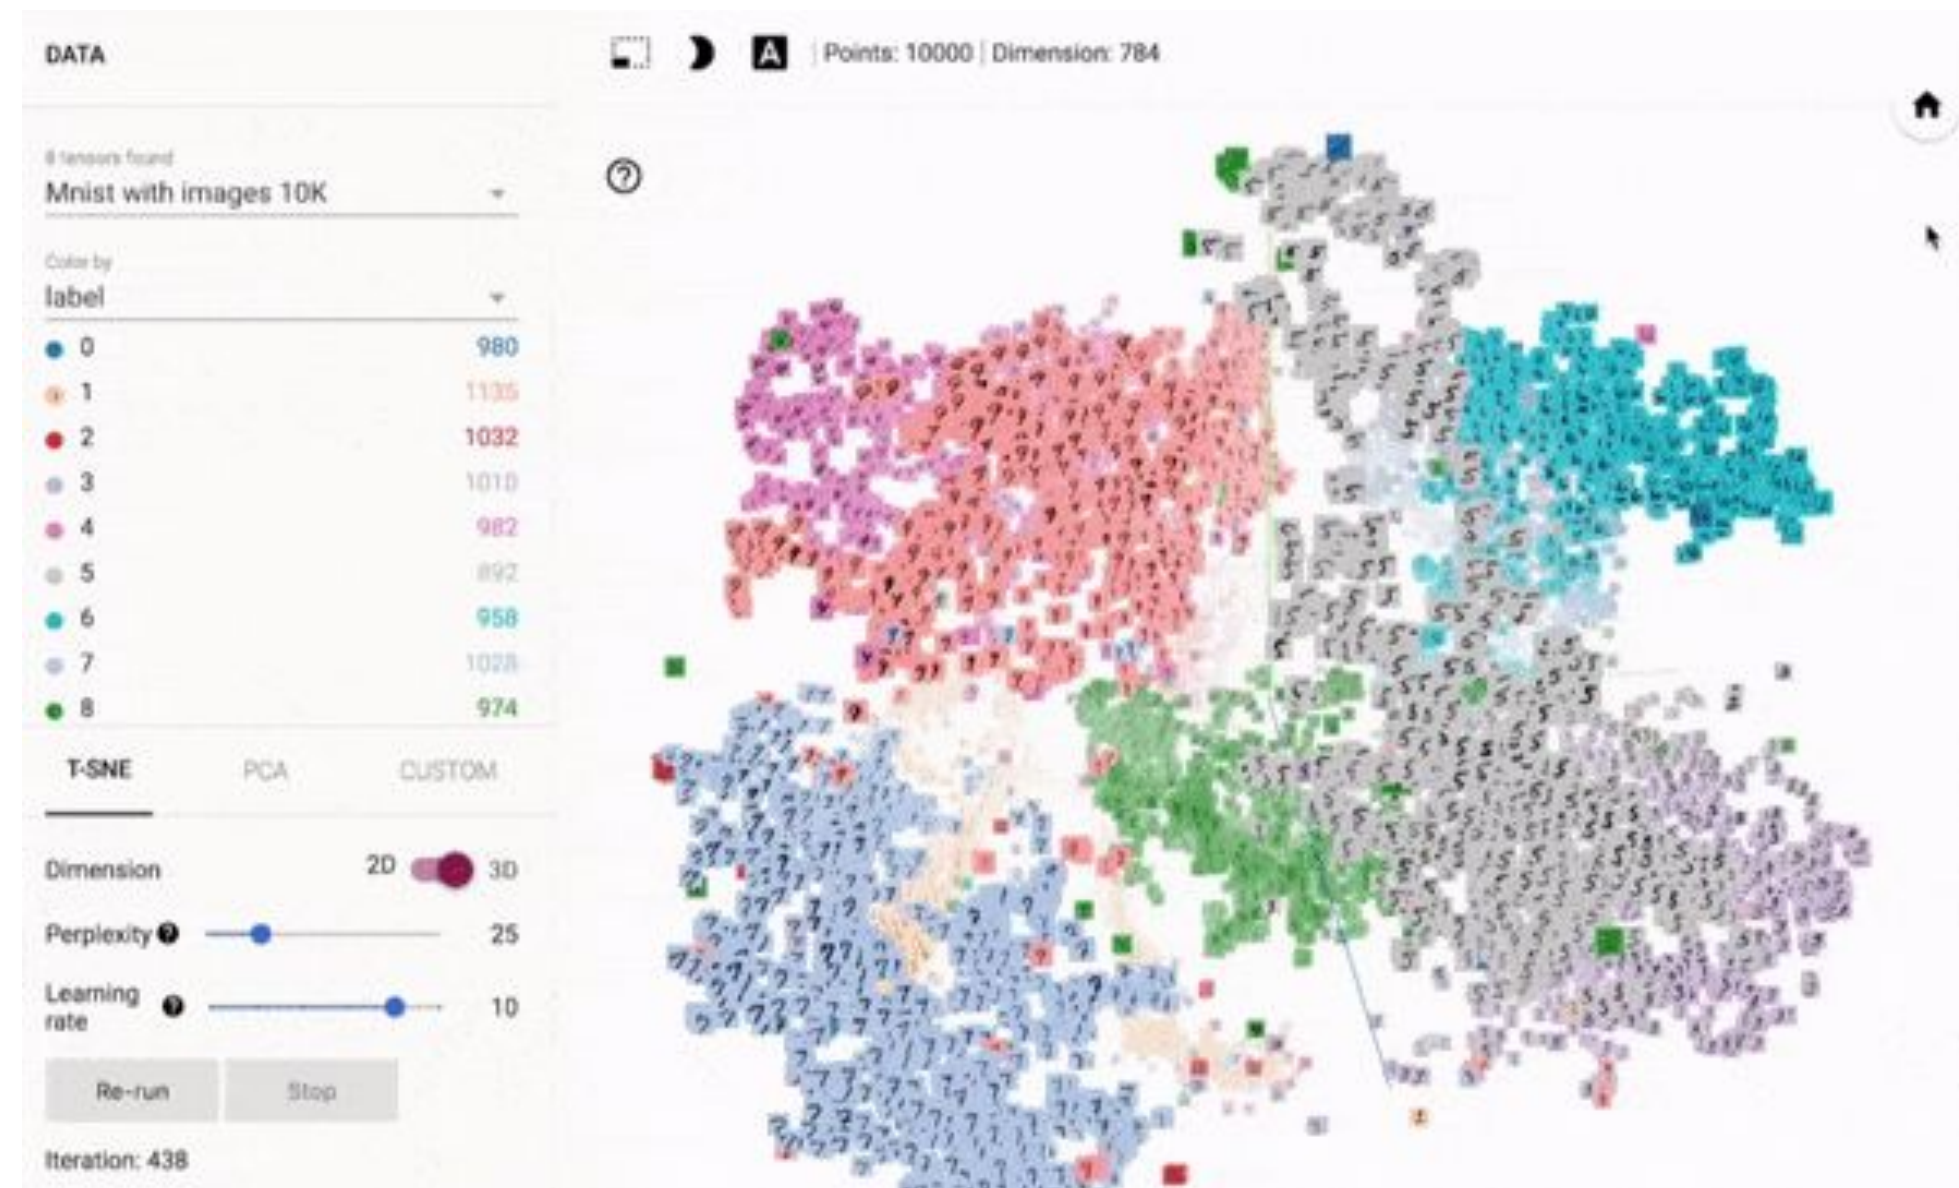
Embeddings provide dimensionality reduction




Embeddings provide dimensionality reduction



The result of embedding is such that similar items are close to each other



You can take advantage of this similarity property of embeddings



A good starting point for number of embedding dimensions

Higher dimensions ->
more accuracy



Higher dimensions ->
overfitting, slow
training

$$\text{dimensions} \approx \sqrt[4]{\text{possible values}}$$

Empirical tradeoff



cloud.google.com

