



---

Custom Estimator

Lak Lakshmanan

# Machine Learning on Google Cloud Platform

---

The Art of ML

Hyperparameter Tuning

A Pinch of Science

The Science of Neural Networks

Embeddings

**Custom Estimator**

# Learn how to...

---

Go beyond canned estimators

# Learn how to...

---

Go beyond canned estimators

Write a custom estimator

# Learn how to...

---

Go beyond canned estimators

Write a custom estimator

Gain control over model  
functions

# Learn how to...

---

Go beyond canned estimators

Write a custom estimator

Gain control over model  
functions

Incorporate Keras models into  
Estimator

# Estimator provides a lot of benefits

Quick model

Checkpointing

Out-of-memory datasets

Train / eval / monitor

Distributed training

Hyper-parameter tuning on ML-Engine

Production: serving predictions from a trained model

# Suppose that you want to use a model structure from a research paper ...

```
def decode_line(row):
    cols = tf.decode_csv(row, record_defaults=[[0], ['house'], [0]])
    features = {'sq_footage': cols[0], 'type': cols[1]}
    label = cols[2] # price
    return features, label

def input_fn():
    return tf.data.Dataset.list_files("train.csv-*") \
        .flat_map(tf.data.TextLineDataset) \
        .map(decode_line).shuffle(1000) \
        .repeat(15).batch(128).
        .make_one_shot_iterator().get_next()

model = tf.estimator.LinearRegressor(featcols, './outdir')
model.train(input_fn)
```

	property type	
sq_footage		PRICE in K\$
1001,	house,	501
2001,	house,	1001
3001,	house,	1501
1001,	apt,	701
2001,	apt,	1301
3001,	apt,	1901
1101,	house,	526
2101,	house,	1026



```
run_config =  
tf.estimator.RunConfig(model_dir=output_dir, ...)  
  
estimator =  
tf.estimator.LinearRegressor(featscols,  
config=run_config)  
  
train_spec =  
tf.estimator.TrainSpec(input_fn=train_input_fn,  
max_steps=1000)  
  
export_latest =  
tf.estimator.LatestExporter(serving_input_receiver_fn=serving_input_fn)  
  
eval_spec =  
tf.estimator.EvalSpec(input_fn=eval_input_fn,  
exporters=export_latest)  
  
tf.estimator.train_and_evaluate(estimator,  
train_spec, eval_spec)
```


```
run_config =  
tf.estimator.RunConfig(model_dir=output_dir, ...)  
  
estimator =  
tf.estimator.LinearRegressor(featscols,  
config=run_config)  
  
train_spec =  
tf.estimator.TrainSpec(input_fn=train_input_fn,  
max_steps=1000)  
  
export_latest =  
tf.estimator.LatestExporter(serving_input_receiver_fn=serving_input_fn)  
  
eval_spec =  
tf.estimator.EvalSpec(input_fn=eval_input_fn,  
exporters=export_latest)  
  
tf.estimator.train_and_evaluate(estimator,  
train_spec, eval_spec)
```

```
run_config =  
tf.estimator.RunConfig(model_dir=output_dir, ...)  
  
estimator =  
tf.estimator.LinearRegressor(featscols,  
config=run_config)  
  
train_spec =  
tf.estimator.TrainSpec(input_fn=train_input_fn,  
max_steps=1000)  
  
export_latest =  
tf.estimator.LatestExporter(serving_input_receiver_fn=serving_input_fn)  
  
eval_spec =  
tf.estimator.EvalSpec(input_fn=eval_input_fn,  
exporters=export_latest)  
  
tf.estimator.train_and_evaluate(estimator,  
train_spec, eval_spec)
```

```
run_config =  
tf.estimator.RunConfig(model_dir=output_dir, ...)  
  
estimator =  
tf.estimator.LinearRegressor(featscols,  
config=run_config)  
  
train_spec =  
tf.estimator.TrainSpec(input_fn=train_input_fn,  
max_steps=1000)  
  
export_latest =  
tf.estimator.LatestExporter(serving_input_receiver_fn=serving_input_fn)  
  
eval_spec =  
tf.estimator.EvalSpec(input_fn=eval_input_fn,  
exporters=export_latest)  
  
tf.estimator.train_and_evaluate(estimator,  
train_spec, eval_spec)
```

```
run_config =  
tf.estimator.RunConfig(model_dir=output_dir, ...)  
  
estimator =  
tf.estimator.LinearRegressor(featcols,  
config=run_config)  
  
train_spec =  
tf.estimator.TrainSpec(input_fn=train_input_fn,  
max_steps=1000)  
  
export_latest =  
tf.estimator.LatestExporter(serving_input_receiver_fn=serving_input_fn)  
  
eval_spec =  
tf.estimator.EvalSpec(input_fn=eval_input_fn,  
exporters=export_latest)  
  
tf.estimator.train_and_evaluate(estimator,  
train_spec, eval_spec)
```

# Canned Estimators are sometimes insufficient



# Suppose that you want to use a model structure from a research paper ...

## Implement the model using low-level TensorFlow ops

```
def model_from_research_paper(timeseries):  
    x = tf.split(timeseries, N_INPUTS, 1)  
    lstm_cell = rnn.BasicLSTMCell(LSTM_SIZE, forget_bias=1.0)  
    outputs, _ = rnn.static_rnn(lstm_cell, x, dtype=tf.float32)  
    outputs = outputs[-1]  
    weight = tf.Variable(tf.random_normal([LSTM_SIZE, N_OUTPUTS]))  
    bias = tf.Variable(tf.random_normal([N_OUTPUTS]))  
    predictions = tf.matmul(outputs, weight) + bias  
    return predictions
```

How do we wrap this custom model into Estimator framework?

# Create train\_and\_evaluate function with the base-class Estimator

```
def train_and_evaluate(output_dir, ...):  
    estimator = tf.estimators.Estimator(model_fn = myfunc,  
    model_dir = output_dir),  
    train_spec = get_train()  
    exporter = ...  
    eval_spec = get_valid()  
    tf.estimator.train_and_evaluate(estimator, train_spec, eval_spec)
```

myfunc (above) is a EstimatorSpec



# The 6 things in a EstimatorSpec

```
def myfunc(features, targets, mode):  
    # Code up the model  
    predictions = model_from_research_paper(features[INCOL])  
  
    # Set up loss function, training/eval ops  
    ... #(next slide)  
  
    # Create dictionary of output tensors  
    predictions_dict = {"predicted": predictions}  
  
    # Create export outputs  
    export_outputs = {"regression_export_outputs":  
        tf.estimator.export.RegressionOutput(value = predictions)}  
  
    # Return EstimatorSpec  
    return tf.estimator.EstimatorSpec(  
        mode = mode,  
        predictions = predictions_dict,  
        loss = loss,  
        train_op = train_op,  
        eval_metric_ops = eval_metric_ops,  
        export_outputs = export_outputs)
```

1. Mode is pass-through
2. Any tensors you want to return
3. Loss metric
4. Training op
5. Eval ops
6. Export outputs

# The ops are set up in the appropriate mode

```
if mode == tf.estimator.ModeKeys.TRAIN or  
    mode == tf.estimator.ModeKeys.EVAL:  
    loss = tf.losses.mean_squared_error(targets, predictions)  
    train_op = tf.contrib.layers.optimize_loss(  
        loss=loss,  
        global_step=tf.contrib.framework.get_global_step(),  
        learning_rate=0.01,  
        optimizer="SGD")  
    eval_metric_ops = {  
        "rmse": tf.metrics.root_mean_squared_error(targets,  
predictions)  
    }  
else:  
    loss = None  
    train_op = None  
    eval_metric_ops = None
```

# Create train\_and\_evaluate function with the base-class Estimator

```
def train_and_evaluate(output_dir, ...):  
    estimator = tf.estimators.Estimator(model_fn = myfunc,  
    model_dir = output_dir),  
    train_spec = get_train()  
    exporter = ...  
    eval_spec = get_valid()  
    tf.estimator.train_and_evaluate(estimator, train_spec, eval_spec)
```

myfunc (above) is an EstimatorSpec

# Lab

---

## Implementing a Custom Estimator

Lak Lakshmanan

# Lab: Implementing a Custom Estimator

You will build a custom estimator to predict time series values

# Create train\_and\_evaluate function with the base-class Estimator

```
def train_and_evaluate(output_dir, ...):  
    estimator = tf.estimators.Estimator(model_fn = myfunc,  
    model_dir = output_dir),  
    train_spec = get_train()  
    exporter = ...  
    eval_spec = get_valid()  
    tf.estimator.train_and_evaluate(estimator, train_spec, eval_spec)
```

myfunc (above) is an EstimatorSpec; how does it work with Keras?

Keras is a high-level deep neural networks library  
that supports multiple backends



theano



# Keras is easy to use for fast prototyping

```
model = Sequential()
model.add(Embedding(max_features, output_dim=256))
model.add(LSTM(128))
model.add(Dropout(0.5))
model.add(Dense(1, activation='sigmoid'))

model.compile(loss='binary_crossentropy',
              optimizer='rmsprop',
              metrics=['accuracy'])

model.fit(x_train, y_train, batch_size=16, epochs=10)
score = model.evaluate(x_test, y_test, batch_size=16)
```



# From a compiled Keras model, you can get an Estimator

```
from tensorflow import keras

model = Sequential()
model.add(Embedding(max_features, output_dim=256))
model.add(LSTM(128))
model.add(Dropout(0.5))
model.add(Dense(1, activation='sigmoid'))

model.compile(loss='binary_crossentropy',
              optimizer='rmsprop',
              metrics=['accuracy'])

estimator = keras.estimator.model_to_estimator(keras_model=model)

model.fit(x_train, y_train, batch_size=16, epochs=10)
score = model.evaluate(x_test, y_test, batch_size=16)
```

# From a compiled Keras model, you can get an Estimator

```
def train_and_evaluate(output_dir):  
    estimator = make_keras_estimator(output_dir)  
    train_spec = tf.estimator.TrainSpec(train_fn,  
                                        max_steps = 1000)  
    exporter = LatestExporter('exporter', serving_input_fn)  
    eval_spec = tf.estimator.EvalSpec(eval_fn,  
                                     steps = None,  
                                     exporters = exporter)  
    tf.estimator.train_and_evaluate(estimator, train_spec,  
    eval_spec)
```

# The connection between the input features and Keras is through a naming convention

```
model = keras.models.Sequential()
model.add(keras.layers.Dense(..., name='XYZ'))

def train_input_fn():
    ...
    features = {
        'XYZ_input': some_tensor,
    }
    return features, labels
```

cloud.google.com