



Estimator API

Introduction to TensorFlow

Martin Gerner

Learn how to...

Learn how to...

Create production-ready
machine learning models the
easy way

Learn how to...

Create production-ready
machine learning models the
easy way

Train on large datasets that do
not fit in memory

Learn how to...

Create production-ready machine learning models the easy way

Train on large datasets that do not fit in memory

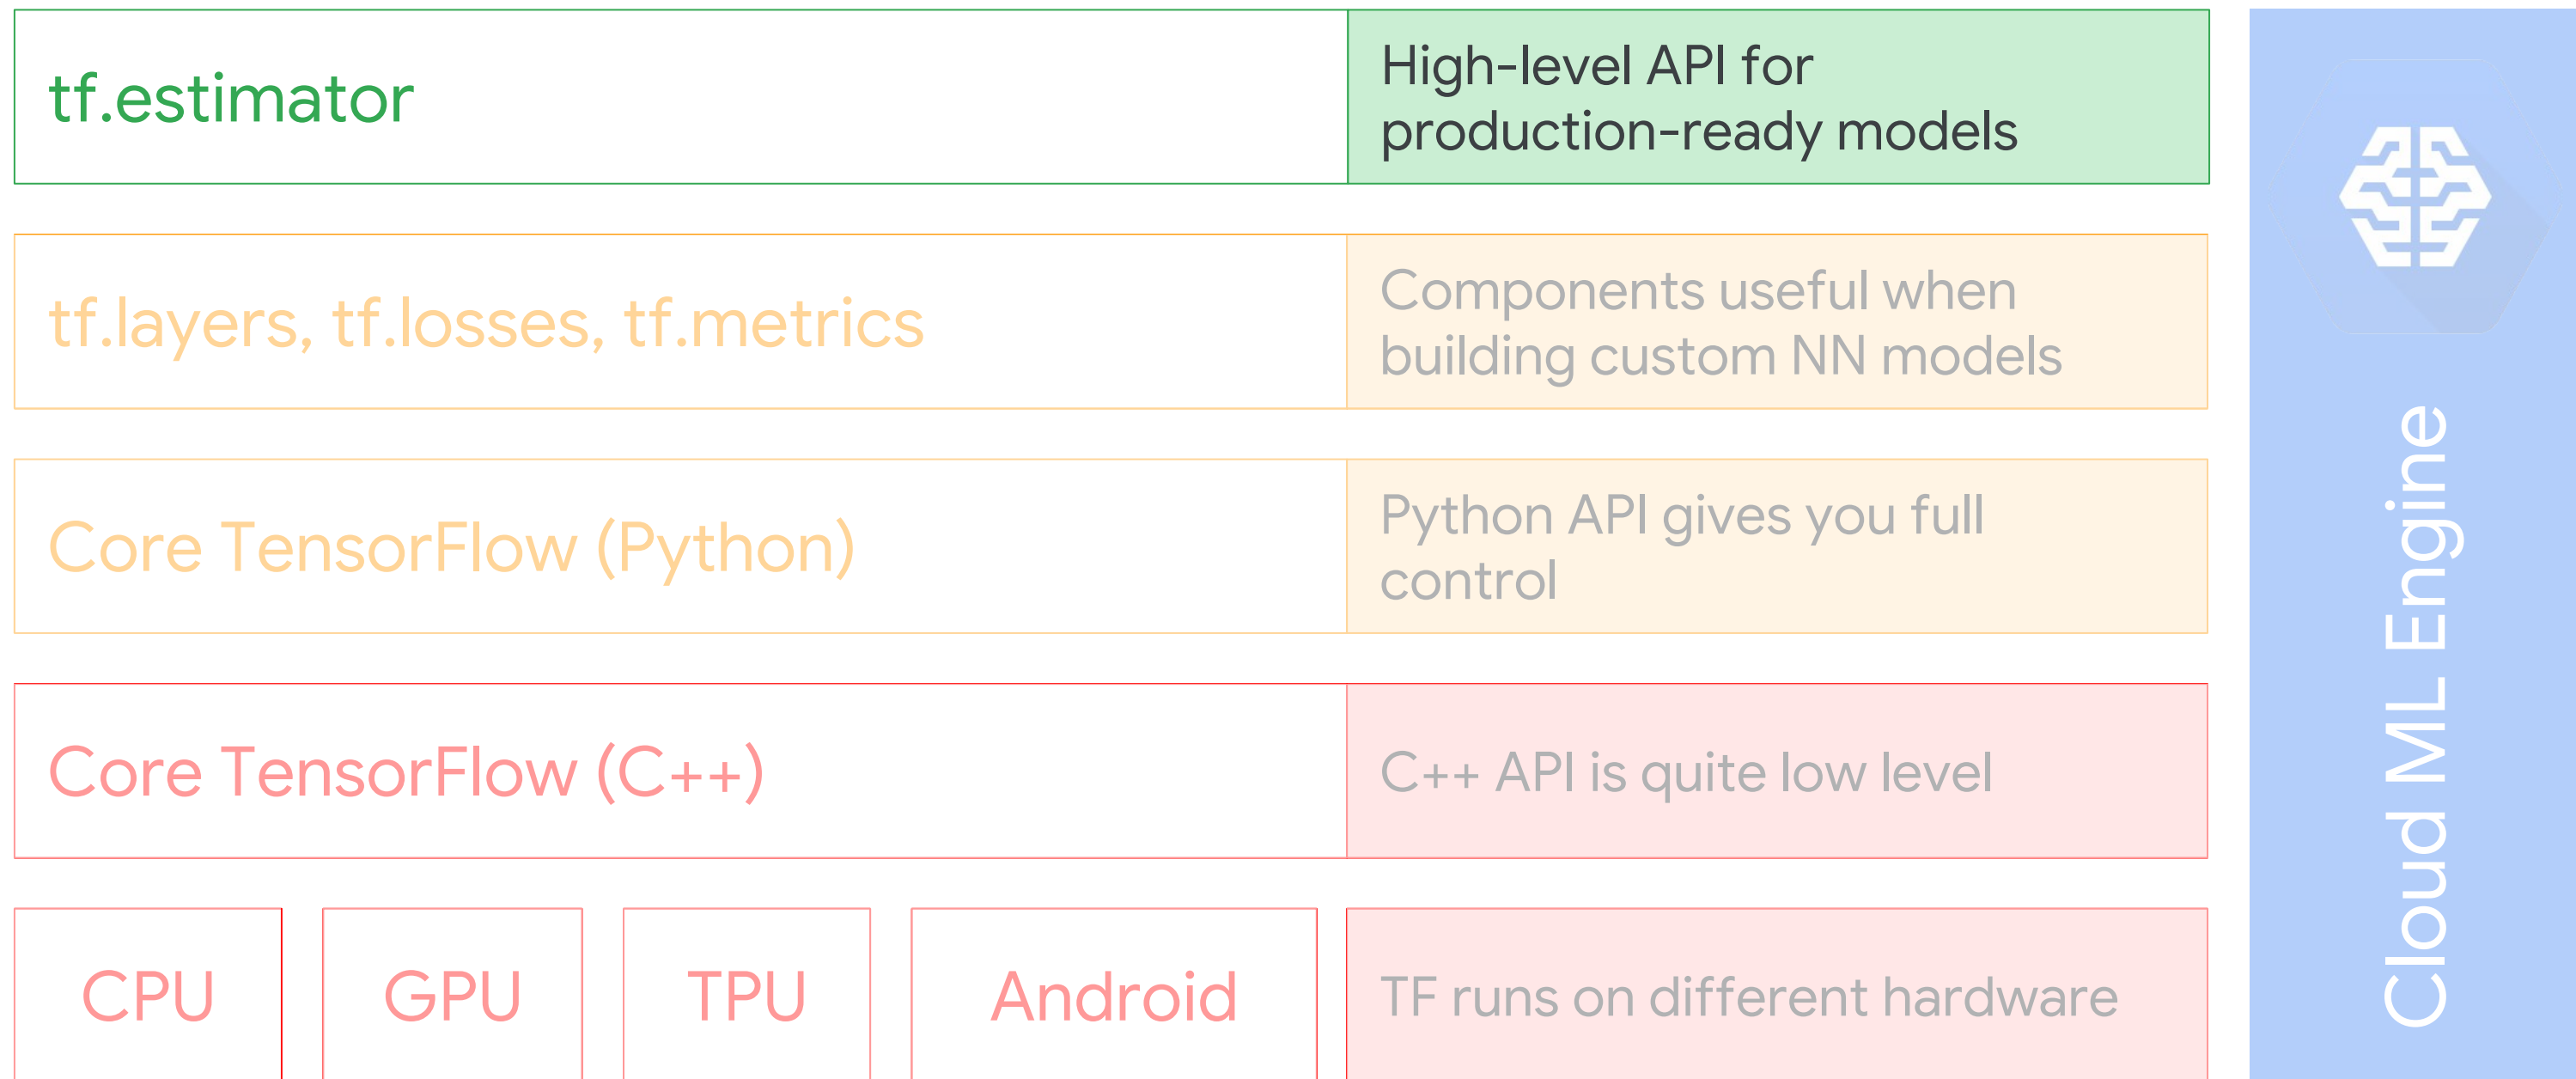
Monitor your training metrics in Tensorboard



Estimator API

Martin Gerner

Estimators wrap up a large amount of boilerplate code, on top of the model itself



From small to big to prod with the Estimator API

From small to big to prod with the Estimator API

Quick model

From small to big to prod with the Estimator API

Quick model

Checkpointing

From small to big to prod with the Estimator API

Quick model

Checkpointing

Out-of-memory datasets

From small to big to prod with the Estimator API

Quick model

Checkpointing

Out-of-memory datasets

Train / eval / monitor

From small to big to prod with the Estimator API

Quick model

Checkpointing

Out-of-memory datasets

Train / eval / monitor

Distributed training

From small to big to prod with the Estimator API

Quick model

Checkpointing

Out-of-memory datasets

Train / eval / monitor

Distributed training

Hyper-parameter tuning on ML-Engine

Production: serving predictions from a trained model

Pre-made estimators that can all be used in the same
way

tf.estimator.Estimator.

Pre-made estimators that can all be used in the same
way

tf.estimator.Estimator.

Pre-made estimators that can all be used in the same
way

tf.estimator.Estimator.

LinearRegressor

DNNRegressor

DNNLinearCombinedRegressor

...

LinearClassifier

DNNClassifier

DNNLinearCombinedClassifier

...

Pre-made regressors

Pre-made classifiers

your custom Estimator

Pre-made estimators that can all be used in the same way

tf.estimator.Estimator.

LinearRegressor

DNNRegressor

DNNLinearCombinedRegressor

...

LinearClassifier

DNNClassifier

DNNLinearCombinedClassifier

...

Pre-made regressors

Pre-made classifiers

your custom Estimator

Predict property value from
historical data



Predict property value from
historical data

“Features”

???

???

Predict property value from historical data

“Features”

Square footage
House / apartment



ML Model



Output

Price: \$400,000

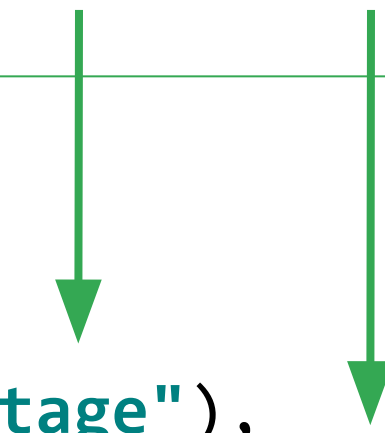
Feature columns tell the model what inputs to expect

“Features”

```
import tensorflow as tf

featcols = [
    tf.feature_column.numeric_column("sq_footage"),
    tf.feature_column.categorical_column_with_vocabulary_list("type",
                                                             ["house", "apt"])
]

model = tf.estimator.LinearRegressor(featcols)
```



Feature columns tell the model what inputs to expect

```
import tensorflow as tf

featcols = [
    tf.feature_column.numeric_column("sq_footage"),
    tf.feature_column.categorical_column_with_vocabulary_list("type",
                                                             ["house", "apt"])
]

model = tf.estimator.LinearRegressor(featcols)
```

Feature columns tell the model what inputs to expect

```
import tensorflow as tf

featcols = [
    tf.feature_column.numeric_column("sq_footage"),
    tf.feature_column.categorical_column_with_vocabulary_list("type",
                                                             ["house", "apt"])
]

model = tf.estimator.LinearRegressor(featcols)
```


Feature columns tell the model what inputs to expect

```
import tensorflow as tf

featcols = [
    tf.feature_column.numeric_column("sq_footage"),
    tf.feature_column.categorical_column_with_vocabulary_list("type",
                                                              ["house", "apt"])
]

model = tf.estimator.LinearRegressor(featcols)
```

Feature columns tell the model what inputs to expect

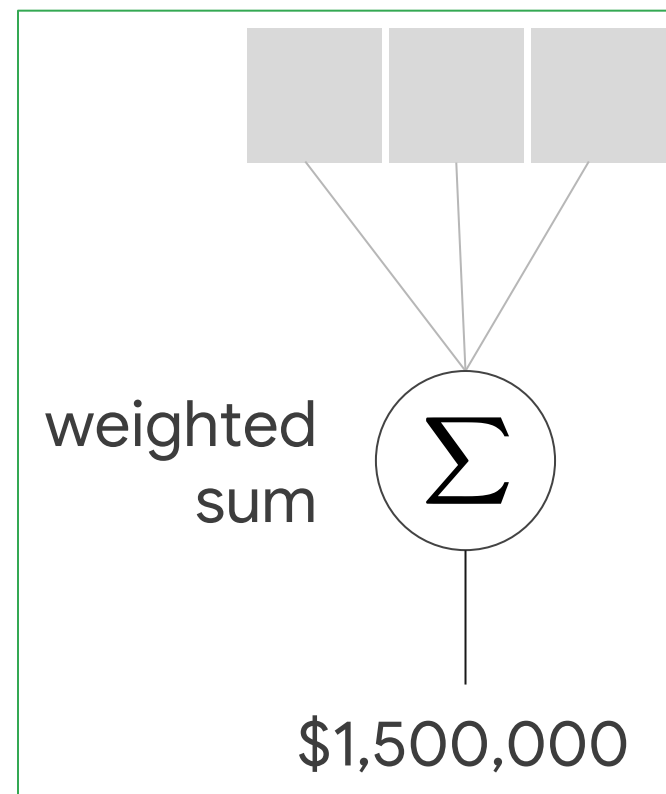
```
import tensorflow as tf

featcols = [
    tf.feature_column.numeric_column("sq_footage"),
    tf.feature_column.categorical_column_with_vocabulary_list("type",
                                                             ["house", "apt"])
]

model = tf.estimator.LinearRegressor(featcols)
```

↑
“Model”: predicts PRICE

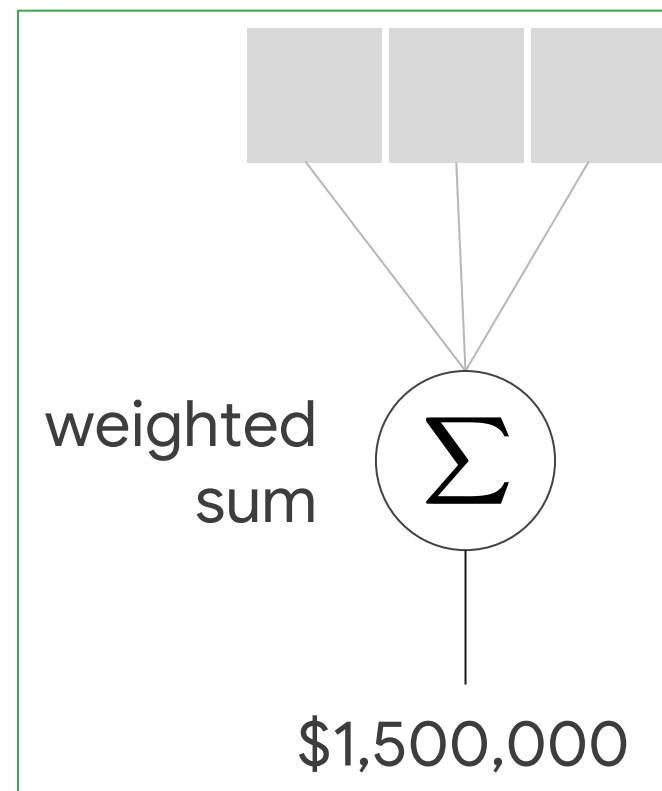
Under the hood: feature columns take care of packing the inputs into the input vector of the model



`tf.estimator.LinearRegressor`

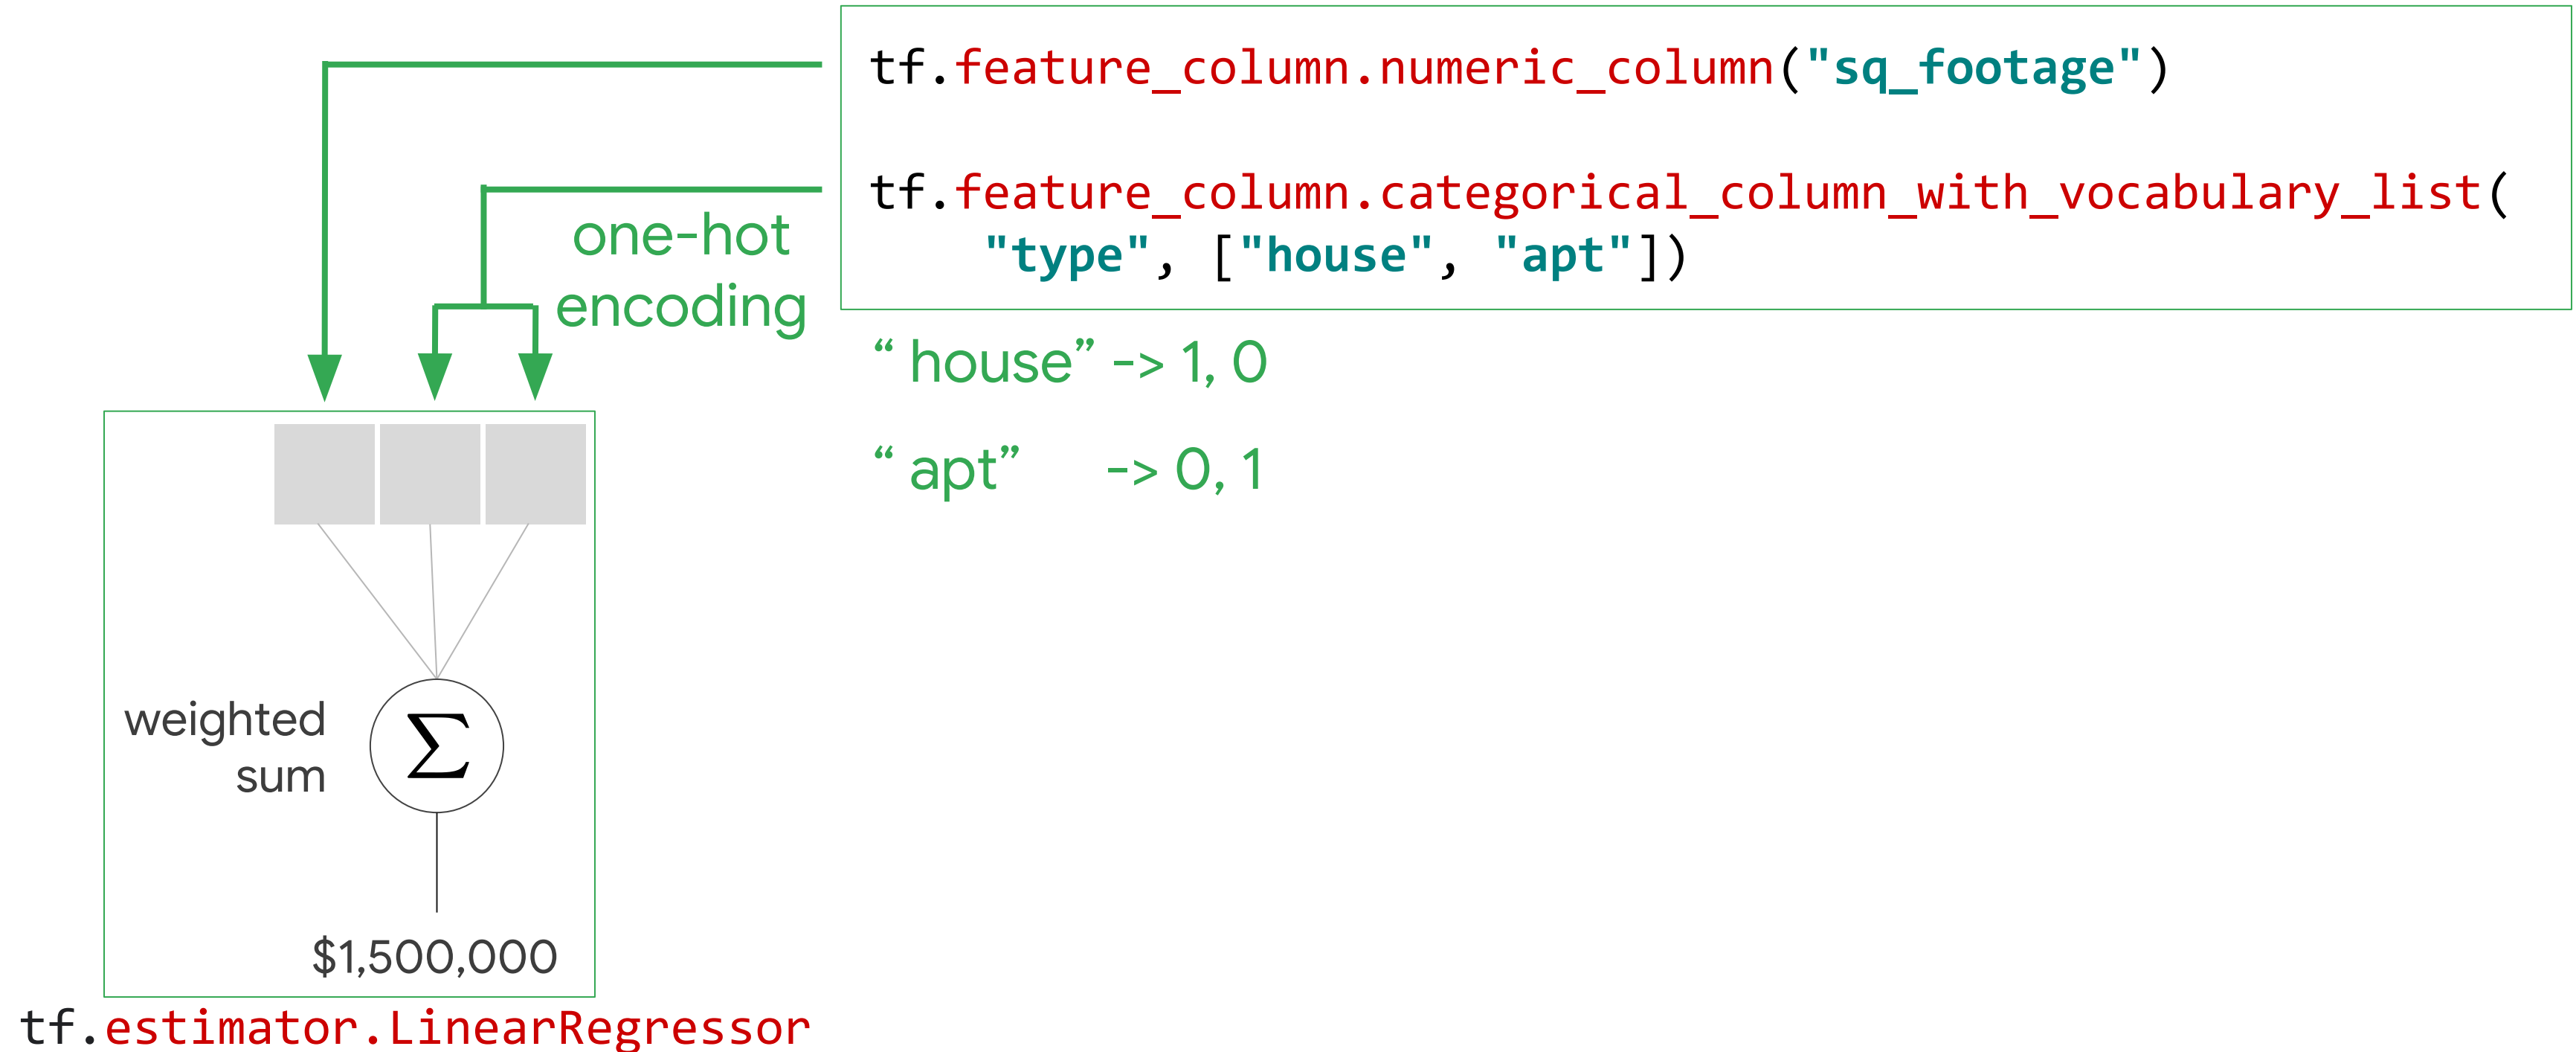
Under the hood: feature columns take care of packing the inputs into the input vector of the model

```
tf.feature_column.numeric_column("sq_footage")
```

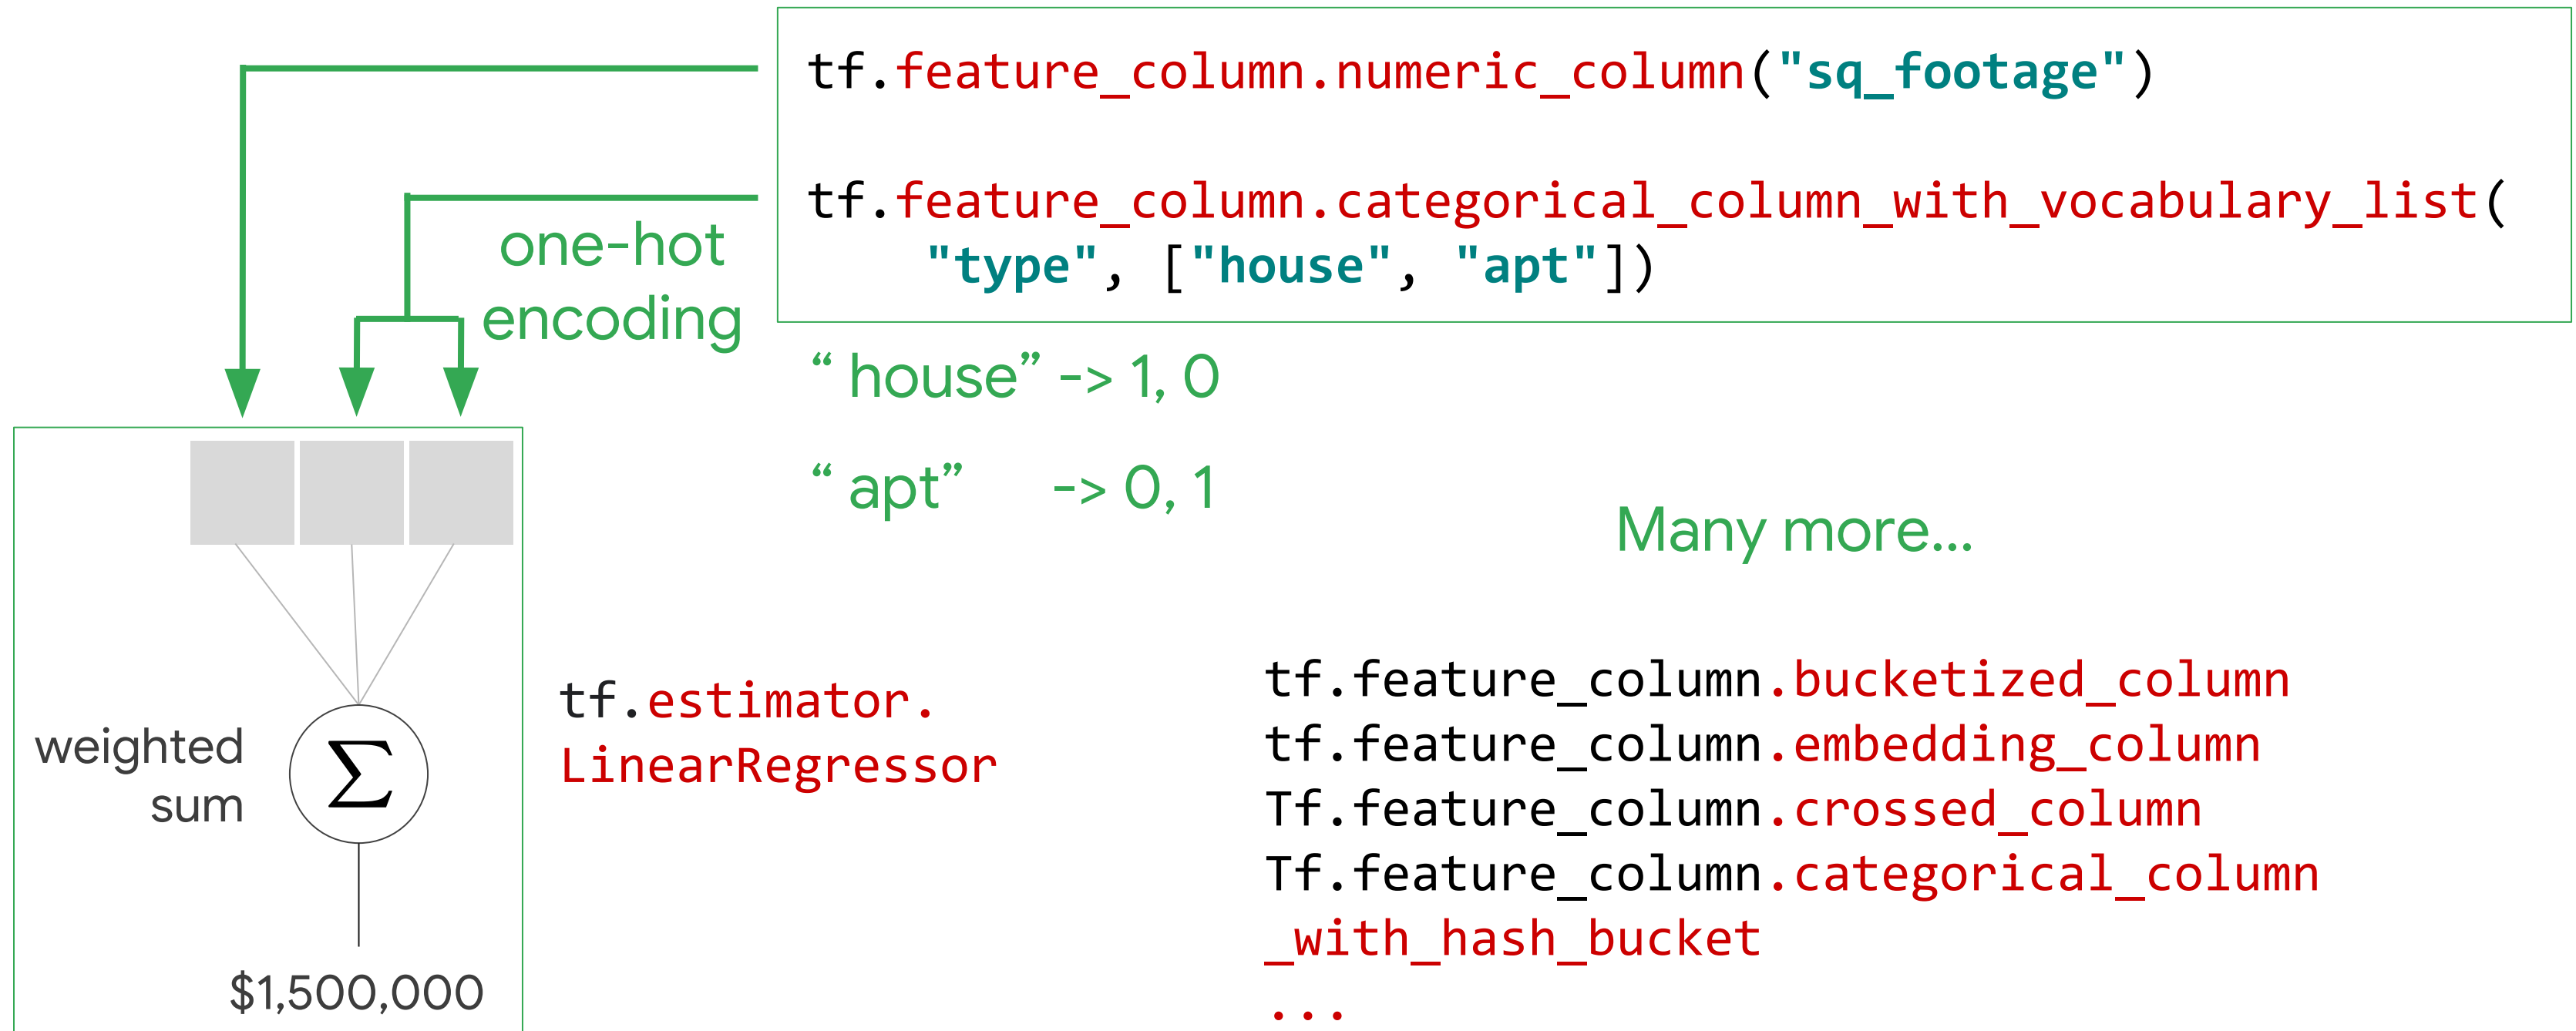


```
tf.estimator.LinearRegressor
```

Under the hood: feature columns take care of packing the inputs into the input vector of the model



Under the hood: feature columns take care of packing the inputs into the input vector of the model



Training: feed in training input data and train for 100 epochs

```
def train_input_fn():  
    features = {"sq_footage": [ 1000,    2000,    3000,    1000,    2000,    3000],  
               "type":       ["house", "house", "house", "apt",  "apt",  "apt"]}  
               # prices in thousands  
    labels = [ 500,    1000,    1500,    700,    1300,    1900]  
    return features, labels
```

```
model.train(train_input_fn, steps=100)
```

Training: feed in training input data and train for 100 epochs

```
def train_input_fn():  
    features = {"sq_footage": [ 1000,    2000,    3000,    1000,    2000,    3000],  
               "type":       ["house", "house", "house", "apt",  "apt",  "apt"]}   
               # prices in thousands  
    labels = [ 500,    1000,    1500,    700,    1300,    1900]  
    return features, labels
```

```
model.train(train_input_fn, steps=100)
```


Predictions: once trained, the model can be used for prediction

```
def predict_input_fn():  
    features = {"sq_footage": [1500, 1800],  
               "type": ["house", "apt"]}  
    return features  
  
predictions = model.predict(predict_input_fn)
```

Predictions: once trained, the model can be used
for prediction

```
def predict_input_fn():  
    features = {"sq_footage": [1500, 1800],  
               "type": ["house", "apt"]}  
    return features  
  
predictions = model.predict(predict_input_fn)
```

Predictions: once trained, the model can be used for prediction

```
def predict_input_fn():  
    features = {"sq_footage": [1500, 1800],  
               "type":       ["house", "apt"]}  
    return features
```

```
predictions = model.predict(predict_input_fn)
```

↑
generator

```
print(next(predictions))  
print(next(predictions))
```

→

```
{'predictions': array([855.93], dtype=float32)}  
{'predictions': array([859.07], dtype=float32)}
```

Pick an Estimator, train, predict

```
import tensorflow as tf

featcols = [
    tf.feature_column.numeric_column("sq_footage"),
    tf.feature_column.categorical_column_with_vocabulary_list("type", ["house", "apt"])
]
```

Pick an Estimator, train, predict

```
import tensorflow as tf

featcols = [
    tf.feature_column.numeric_column("sq_footage"),
    tf.feature_column.categorical_column_with_vocabulary_list("type", ["house", "apt"])
]

model = tf.estimator.LinearRegressor(featcols)
```

Pick an Estimator, train, predict

```
import tensorflow as tf

featcols = [
    tf.feature_column.numeric_column("sq_footage"),
    tf.feature_column.categorical_column_with_vocabulary_list("type", ["house", "apt"])
]

model = tf.estimator.LinearRegressor(featcols)

model.train(train_input_fn, steps=100)
```

Pick an Estimator, train, predict

```
import tensorflow as tf

featcols = [
    tf.feature_column.numeric_column("sq_footage"),
    tf.feature_column.categorical_column_with_vocabulary_list("type", ["house", "apt"])
]

model = tf.estimator.LinearRegressor(featcols)

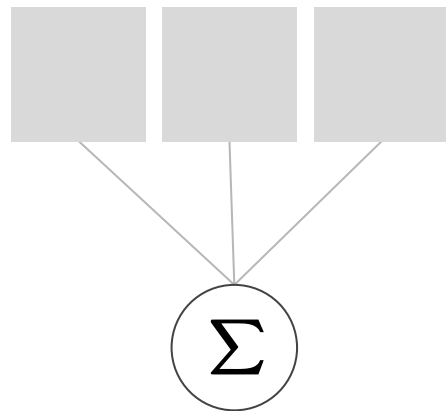
model.train(train_input_fn, steps=100)

predictions = model.predict(predict_input_fn)
```



To use a different pre-made estimator,
just change the class name and supply
appropriate parameters

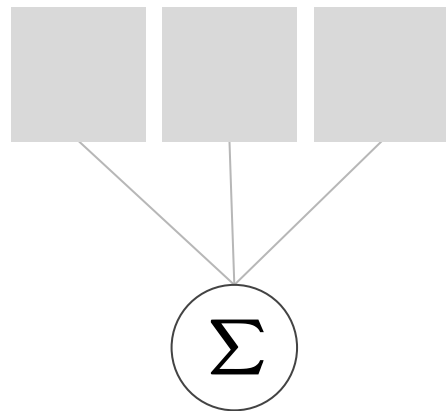
```
model = tf.estimator.DNNRegressor(featcols,  
                                   hidden_units=[3, 2])
```



tf.estimator.
LinearRegressor

To use a different pre-made estimator,
just change the class name and supply
appropriate parameters

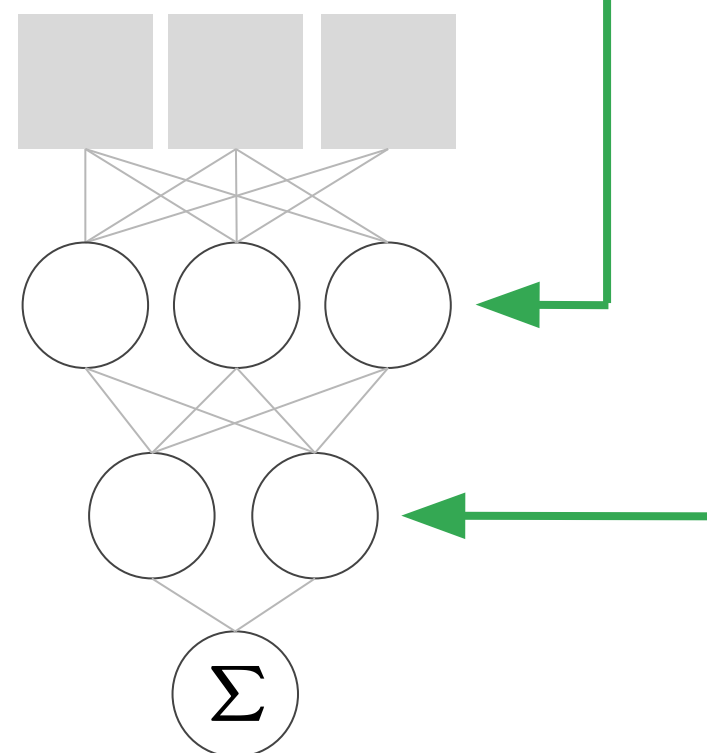
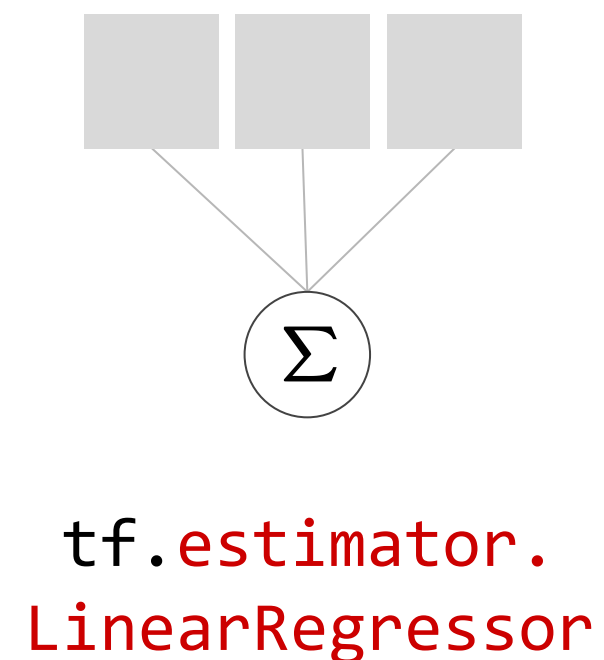
```
model = tf.estimator.DNNRegressor(featscols,  
                                   hidden_units=[3, 2])
```



tf.estimator.
LinearRegressor

To use a different pre-made estimator,
just change the class name and supply
appropriate parameters

```
model = tf.estimator.DNNRegressor(featcols,  
                                   hidden_units=[3, 2])
```



For example, here are some of the things you can change about the DNN Regressor

```
tf.estimator.DNNRegressor(feature_columns=...,  
                           hidden_units=[10, 5],
```

For example, here are some of the things you can change about the DNN Regressor

```
tf.estimator.DNNRegressor(feature_columns=...,  
                           hidden_units=[10, 5],  
                           activation_fn=tf.nn.relu,
```

For example, here are some of the things you can change about the DNN Regressor

```
tf.estimator.DNNRegressor(feature_columns=...,  
                           hidden_units=[10, 5],  
                           activation_fn=tf.nn.relu,  
                           dropout=0.2,  
                           optimizer="Adam")
```

Model checkpoints

1 Continue training

Model checkpoints

1 Continue training

2 Resume on failure

Model checkpoints

1

Continue training

2

Resume on failure

3

Predict from trained
model

Estimators automatically checkpoint training

Where to put the checkpoints



```
model = tf.estimator.LinearRegressor(featcols, './model_trained')  
  
model.train(train_input_fn, steps=100)
```

```
%ls model_trained
```

```
checkpoint
```

```
graph.pbtxt
```

```
model.ckpt-100.data-00000-of-00001
```

```
model.ckpt-100.index
```

```
model.ckpt-100.meta
```

```
model.ckpt-1.data-00000-of-00001
```

```
model.ckpt-1.index
```

```
model.ckpt-1.meta
```

We can now restore and predict with the model

```
trained_model = tf.estimator.LinearRegressor(featcols, './model_trained')  
predictions = trained_model.predict(pred_input_fn)
```

```
INFO:tensorflow:Restoring parameters from  
model_trained/model.ckpt-100
```

```
{'predictions': array([855.93], dtype=float32)}  
{'predictions': array([859.07], dtype=float32)}
```

Training also resumes from the last checkpoint

```
model = tf.estimator.LinearRegressor(featcols, './model_trained')  
model.train(train_input_fn, steps=100)
```

Training continues



Training also resumes from the last checkpoint

```
model = tf.estimator.LinearRegressor(featcols, './model_trained')  
model.train(train_input_fn, steps=100)
```

Training continues



```
INFO:tensorflow:Restoring parameters from  
model_trained/model.ckpt-100
```

In memory data: usually numpy arrays or
Pandas dataframes - you can use them directly

```
def numpy_train_input_fn(sqft, prop_type, price): #np arrays
    return tf.estimator.inputs.numpy_input_fn(
        x = {"sq_footage": sqft, "type": prop_type},
        y = price,
        batch_size=128,
        num_epochs=10,
        shuffle=True,
        queue_capacity=1000
    )
```

```
def pandas_train_input_fn(df): # a Pandas dataframe
    return tf.estimator.inputs.pandas_input_fn(
        x = df, # "sq_footage", "type" selected
            # automatically because of feature
            # columns definition
        y = df['price'],
        batch_size=128,
        num_epochs=10,
        shuffle=True,
        queue_capacity=1000
    )
```

In memory data: usually numpy arrays or
Pandas dataframes - you can use them directly

```
def numpy_train_input_fn(sqft, prop_type, price): #np arrays
    return tf.estimator.inputs.numpy_input_fn(
        x = {"sq_footage": sqft, "type": prop_type},
        y = price,
        batch_size=128,
        num_epochs=10,
        shuffle=True,
        queue_capacity=1000
    )
```

```
def pandas_train_input_fn(df): # a Pandas dataframe
    return tf.estimator.inputs.pandas_input_fn(
        x = df, # "sq_footage", "type" selected
            # automatically because of feature
            # columns definition
        y = df['price'],
        batch_size=128,
        num_epochs=10,
        shuffle=True,
        queue_capacity=1000
    )
```

In memory data: usually numpy arrays or
Pandas dataframes - you can use them directly

```
def numpy_train_input_fn(sqft, prop_type, price): #np arrays
    return tf.estimator.inputs.numpy_input_fn(
        x = {"sq_footage": sqft, "type": prop_type},
        y = price,
        batch_size=128,
        num_epochs=10,
        shuffle=True,
        queue_capacity=1000
    )
```

```
def pandas_train_input_fn(df): # a Pandas dataframe
    return tf.estimator.inputs.pandas_input_fn(
        x = df, # "sq_footage", "type" selected
                # automatically because of feature
                # columns definition
        y = df['price'],
        batch_size=128,
        num_epochs=10,
        shuffle=True,
        queue_capacity=1000
    )
```

In memory data: usually numpy arrays or
Pandas dataframes - you can use them directly

```
def numpy_train_input_fn(sqft, prop_type, price): #np arrays
    return tf.estimator.inputs.numpy_input_fn(
        x = {"sq_footage": sqft, "type": prop_type},
        y = price,
        batch_size=128,
        num_epochs=10,
        shuffle=True,
        queue_capacity=1000
    )
```

```
def pandas_train_input_fn(df): # a Pandas dataframe
    return tf.estimator.inputs.pandas_input_fn(
        x = df, # "sq_footage", "type" selected
            # automatically because of feature
            # columns definition
        y = df['price'],
        batch_size=128,
        num_epochs=10,
        shuffle=True,
        queue_capacity=1000
    )
```


In memory data: usually numpy arrays or
Pandas dataframes - you can use them directly

```
def numpy_train_input_fn(sqft, prop_type, price): #np arrays
    return tf.estimator.inputs.numpy_input_fn(
        x = {"sq_footage": sqft, "type": prop_type},
        y = price,
        batch_size=128,
        num_epochs=10,
        shuffle=True,
        queue_capacity=1000
    )
```

```
def pandas_train_input_fn(df): # a Pandas dataframe
    return tf.estimator.inputs.pandas_input_fn(
        x = df, # "sq_footage", "type" selected
            # automatically because of feature
            # columns definition
        y = df['price'],
        batch_size=128,
        num_epochs=10,
        shuffle=True,
        queue_capacity=1000
    )
```

Training happens until input is exhausted or number of steps is reached

```
def pandas_train_input_fn(df): # a Pandas dataframe
    return tf.estimator.inputs.pandas_input_fn(
        x = df,
        y = df['price'],
        batch_size=128,
        num_epochs=10,
        shuffle=True
    )
```

```
model.train(pandas_train_input_fn(df))
```

Trains until input exhausted (10 epochs)
starting from checkpoint



Training happens until input is exhausted or number of steps is reached

```
def pandas_train_input_fn(df): # a Pandas dataframe
    return tf.estimator.inputs.pandas_input_fn(
        x = df,
        y = df['price'],
        batch_size=128,
        num_epochs=10,
        shuffle=True
    )
```

```
model.train(pandas_train_input_fn(df))
```

```
model.train(pandas_train_input_fn(df), steps=1000)
```

Trains until input exhausted (10 epochs)
starting from checkpoint



```
graph LR
    A["Trains until input exhausted (10 epochs)  
starting from checkpoint"] --> B["model.train(pandas_train_input_fn(df))"]
```

1000 additional steps
from checkpoint



```
graph LR
    C["1000 additional steps  
from checkpoint"] --> D["model.train(pandas_train_input_fn(df), steps=1000)"]
```

Training happens until input is exhausted or number of steps is reached

```
def pandas_train_input_fn(df): # a Pandas dataframe
    return tf.estimator.inputs.pandas_input_fn(
        x = df,
        y = df['price'],
        batch_size=128,
        num_epochs=10,
        shuffle=True
    )
```

```
model.train(pandas_train_input_fn(df))
```

Trains until input exhausted (10 epochs)
starting from checkpoint

```
model.train(pandas_train_input_fn(df), steps=1000)
```

1000 additional steps
from checkpoint

```
model.train(pandas_train_input_fn(df), max_steps=1000)
```

1000 steps - might be
nothing if checkpoint
already there

To add a new feature, add it to the list of feature columns and make sure it is present in data frame

```
featcols = [  
    tf.feature_column.numeric_column("sq_footage"),  
    tf.feature_column.categorical_column_with_vocabulary_list("type",  
                                                             ["house", "apt"])  
]  
  
model = tf.estimator.LinearRegressor(featcols)  
  
def train_input_fn(df): # a Pandas dataframe  
    return tf.estimator.inputs.pandas_input_fn(  
        x = df,  
        y = df['price'],  
        batch_size=128, num_epochs=10, shuffle=True  
    )  
  
model.train(train_input_fn(df))
```

To add a new feature, add it to the list of feature columns and make sure it is present in data frame

```
featcols = [  
    tf.feature_column.numeric_column("sq_footage"),  
    tf.feature_column.categorical_column_with_vocabulary_list("type",  
                                                             ["house", "apt"]),  
    tf.feature_column.numeric_column("nbeds"),  
]  
  
model = tf.estimator.LinearRegressor(featcols)  
  
def train_input_fn(df): # a Pandas dataframe  
    return tf.estimator.inputs.pandas_input_fn(  
        x = df,  
        y = df['price'],  
        batch_size=128, num_epochs=10, shuffle=True  
    )  
  
model.train(train_input_fn(df))
```

Lab

Implementing a Machine
Learning model in TensorFlow
using Estimator API



Train on large datasets with
Dataset API

Lak Lakshmanan









Real World ML Models



Problem	Solution
Out of memory data	?
Distribution	?
Need to evaluate during training	?
Deployments that scale	?

Real World ML Models

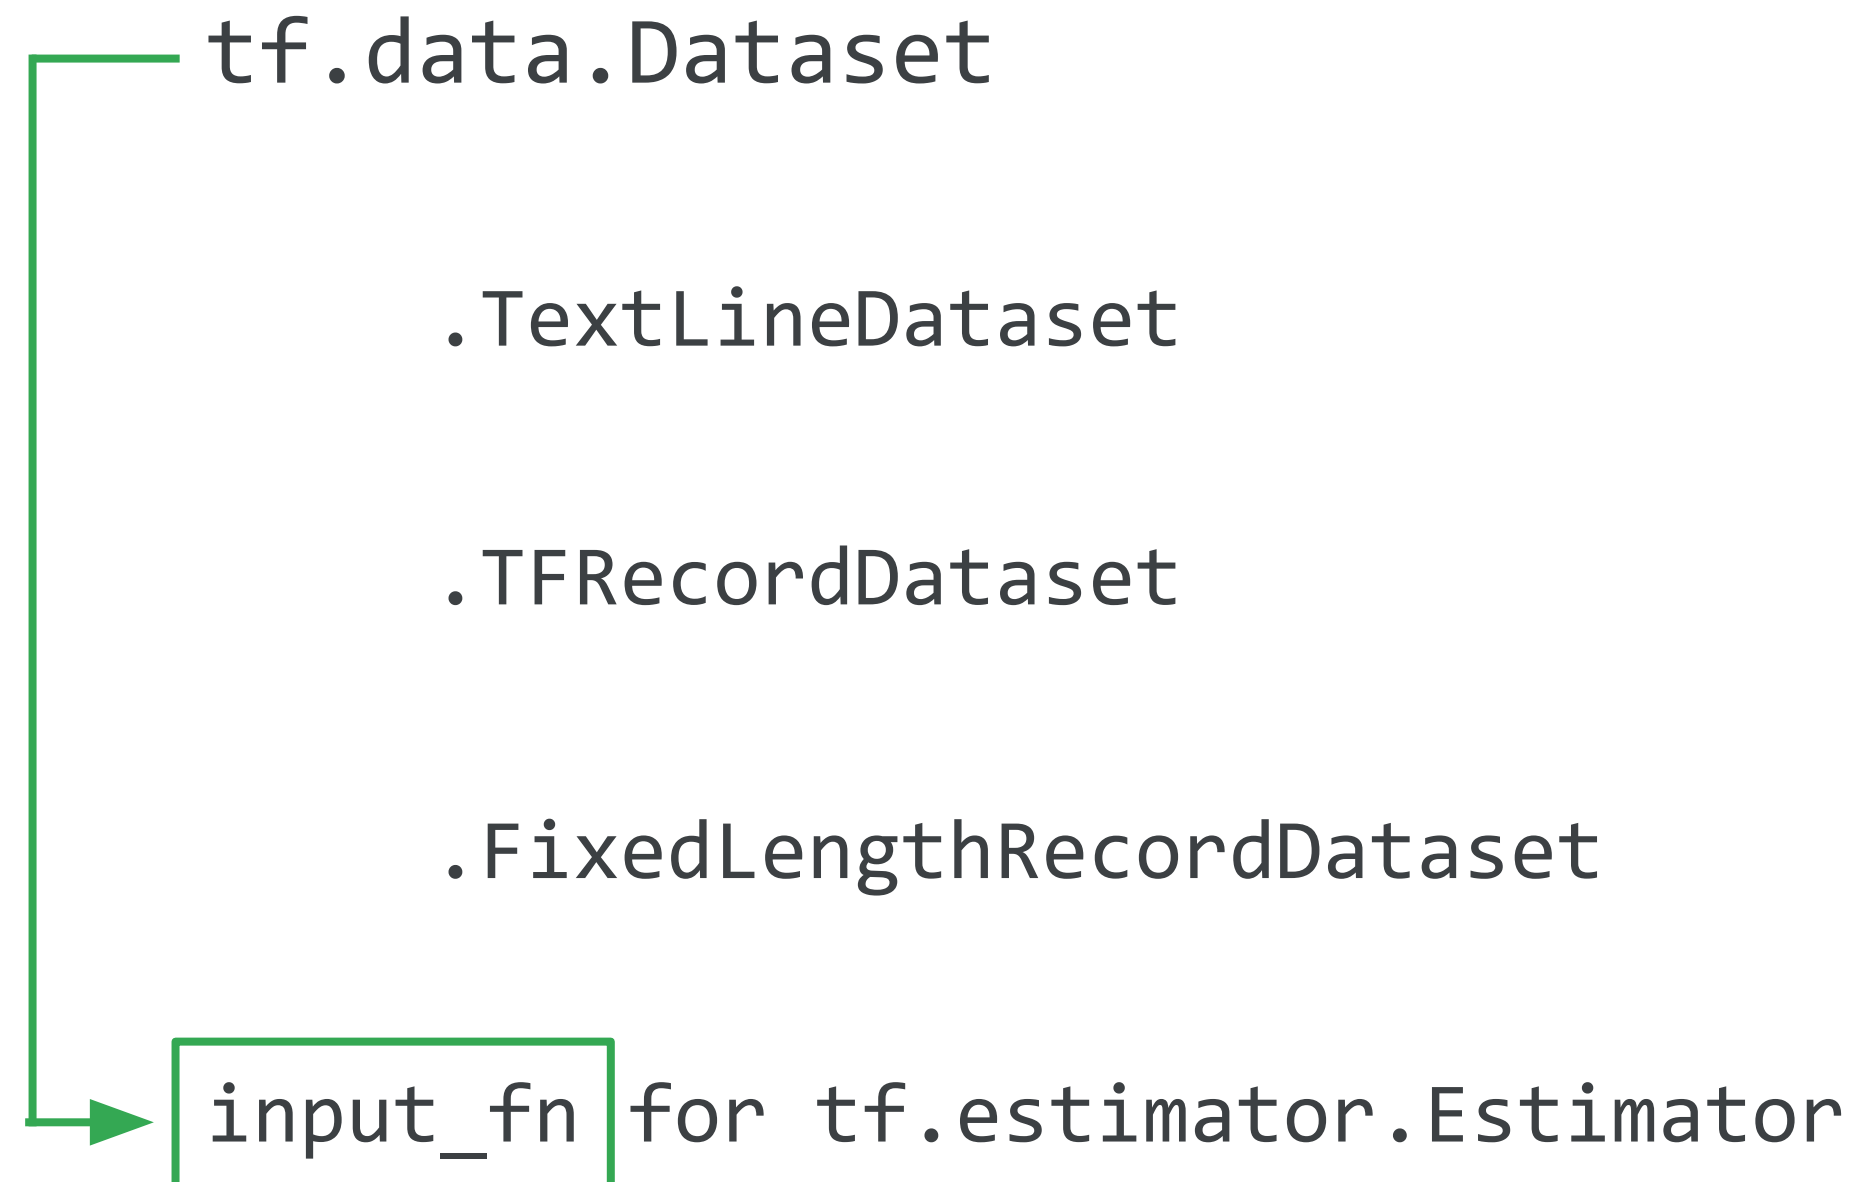
Problem	Solution
Out of memory data	Use the Dataset API
Distribution	?
Need to evaluate during training	?
Deployments that scale	?

Out-of-memory datasets tend to be sharded into multiple files

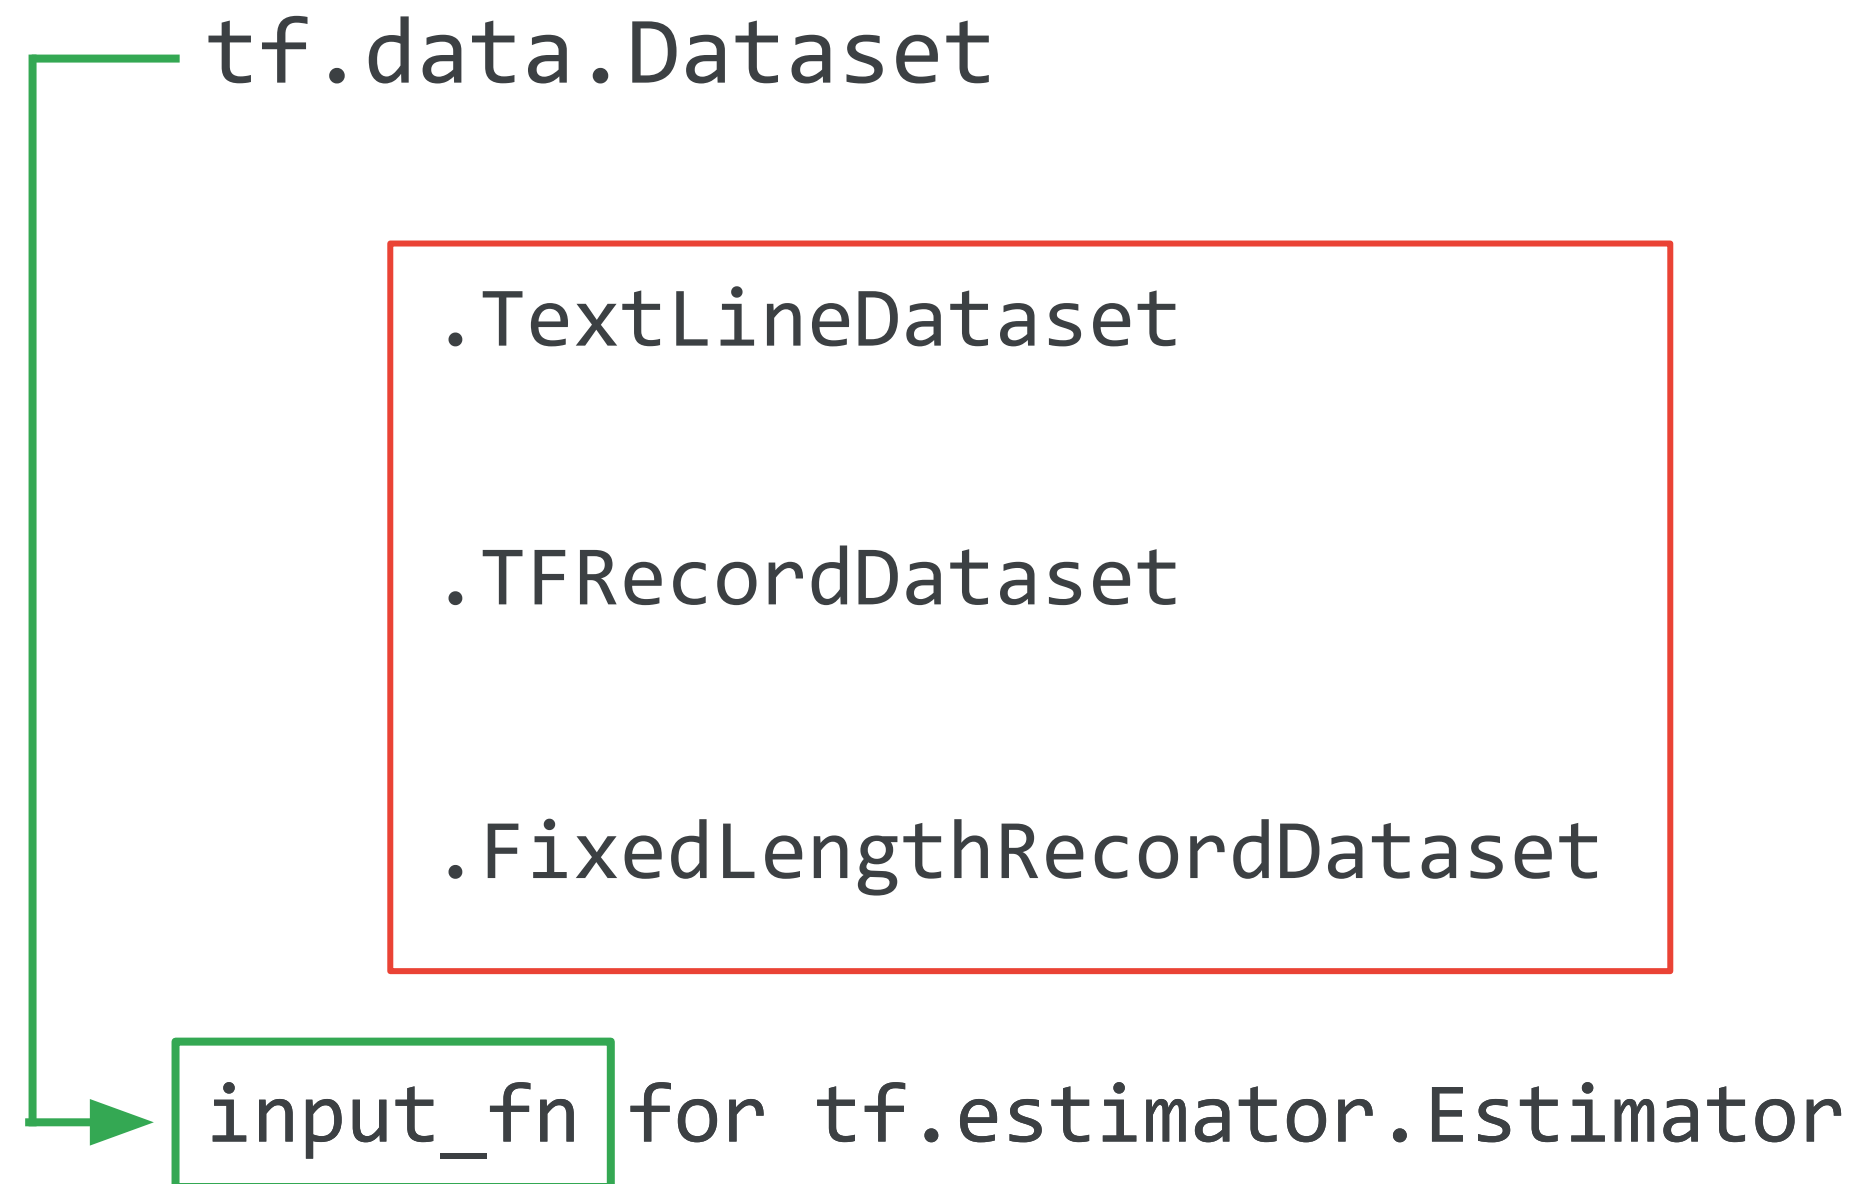
<input type="checkbox"/>		train.csv-00000-of-00011	9.23 MB
<input type="checkbox"/>		train.csv-00001-of-00011	16.82 MB
<input type="checkbox"/>		train.csv-00002-of-00011	44.18 MB
<input type="checkbox"/>		train.csv-00003-of-00011	14.63 MB
<input type="checkbox"/>		train.csv-00004-of-00011	45.58 MB
<input type="checkbox"/>		train.csv-00005-of-00011	11.29 MB
<input type="checkbox"/>		train.csv-00006-of-00011	10.24 MB
<input type="checkbox"/>		train.csv-00007-of-00011	49.75 MB

<input type="checkbox"/>		valid.csv-00000-of-00001	2.31 MB
<input type="checkbox"/>		valid.csv-00000-of-00009	19.47 MB
<input type="checkbox"/>		valid.csv-00001-of-00009	11.6 MB
<input type="checkbox"/>		valid.csv-00002-of-00009	9.5 MB
<input type="checkbox"/>		valid.csv-00003-of-00009	18.29 MB

Datasets can be created from different file formats.
They generate input functions for Estimators.



Datasets can be created from different file formats.
They generate input functions for Estimators.



Read one CSV file using TextLineDataset

```
def decode_line(row):  
    cols = tf.decode_csv(row, record_defaults=[[0], ['house'], [0]])  
    features = {'sq_footage': cols[0], 'type': cols[1]}  
    label = cols[2] # price  
    return features, label
```

```
dataset = tf.data.TextLineDataset("train_1.csv") \  
    .map(decode_line)
```

```
dataset = dataset.shuffle(1000) \  
    .repeat(15) \  
    .batch(128)
```

```
def input_fn():  
    features, label = dataset.make_one_shot_iterator().get_next()  
    return features, label
```

```
model.train(input_fn)
```

property type
sq_footage PRICE in K\$

1001,	house,	501
2001,	house,	1001
3001,	house,	1501
1001,	apt,	701
2001,	apt,	1301
3001,	apt,	1901
1101,	house,	526
2101,	house,	1026

Datasets handle shuffling, epochs, batching, ...

```
dataset = dataset.shuffle(1000) \
    .repeat(15) \
    .batch(128)
```

Shuffle buffer size

Nb of epochs

property type
sq_footage PRICE in K\$

1001	house	501
2001	house	1001
3001	house	1501
1001	apt	701
2001	apt	1301
3001	apt	1901
1101	house	526
2101	house	1026

They support arbitrary transformations with map()

```
def decode_line(txt_line):  
    cols = tf.decode_csv(txt_line, record_defaults=[[0], ['house'], [0]])  
    features = {'sq_footage': cols[0], 'type': cols[1]}  
    label = cols[2] # price  
    return features, label
```

```
dataset = tf.data.TextLineDataset("train_1.csv") \  
    .map(decode_line)
```

Dataset of
text lines

sq_footage property type PRICE in K\$

1001,	house,	501
2001,	house,	1001
3001,	house,	1501
1001,	apt,	701
2001,	apt,	1301
3001,	apt,	1901
1101,	house,	526
2101,	house,	1026

They support arbitrary transformations with `map()`

```
def decode_line(txt_line):  
    cols = tf.decode_csv(txt_line, record_defaults=[[0], ['house'], [0]])  
    features = {'sq_footage': cols[0], 'type': cols[1]}  
    label = cols[2] # price  
    return features, label
```

```
dataset = tf.data.TextLineDataset("train_1.csv") \  
    .map(decode_line)
```



Dataset of pairs
(features, label)

sq_footage property type PRICE in K\$

1001,	house,	501
2001,	house,	1001
3001,	house,	1501
1001,	apt,	701
2001,	apt,	1301
3001,	apt,	1901
1101,	house,	526
2101,	house,	1026

Datasets help create `input_fn`'s for Estimators

```
dataset = ...
```

```
def input_fn():  
    features, label = dataset.make_one_shot_iterator().get_next()  
    return features, label
```

```
model.train(input_fn)
```

property type
sq_footage PRICE in K\$

1001,	house,	501
2001,	house,	1001
3001,	house,	1501
1001,	apt,	701
2001,	apt,	1301
3001,	apt,	1901
1101,	house,	526
2101,	house,	1026

Datasets help create `input_fn`'s for Estimators

```
dataset = ...
```

```
def input_fn():  
    features, label = dataset.make_one_shot_iterator().get_next()  
    return features, label
```

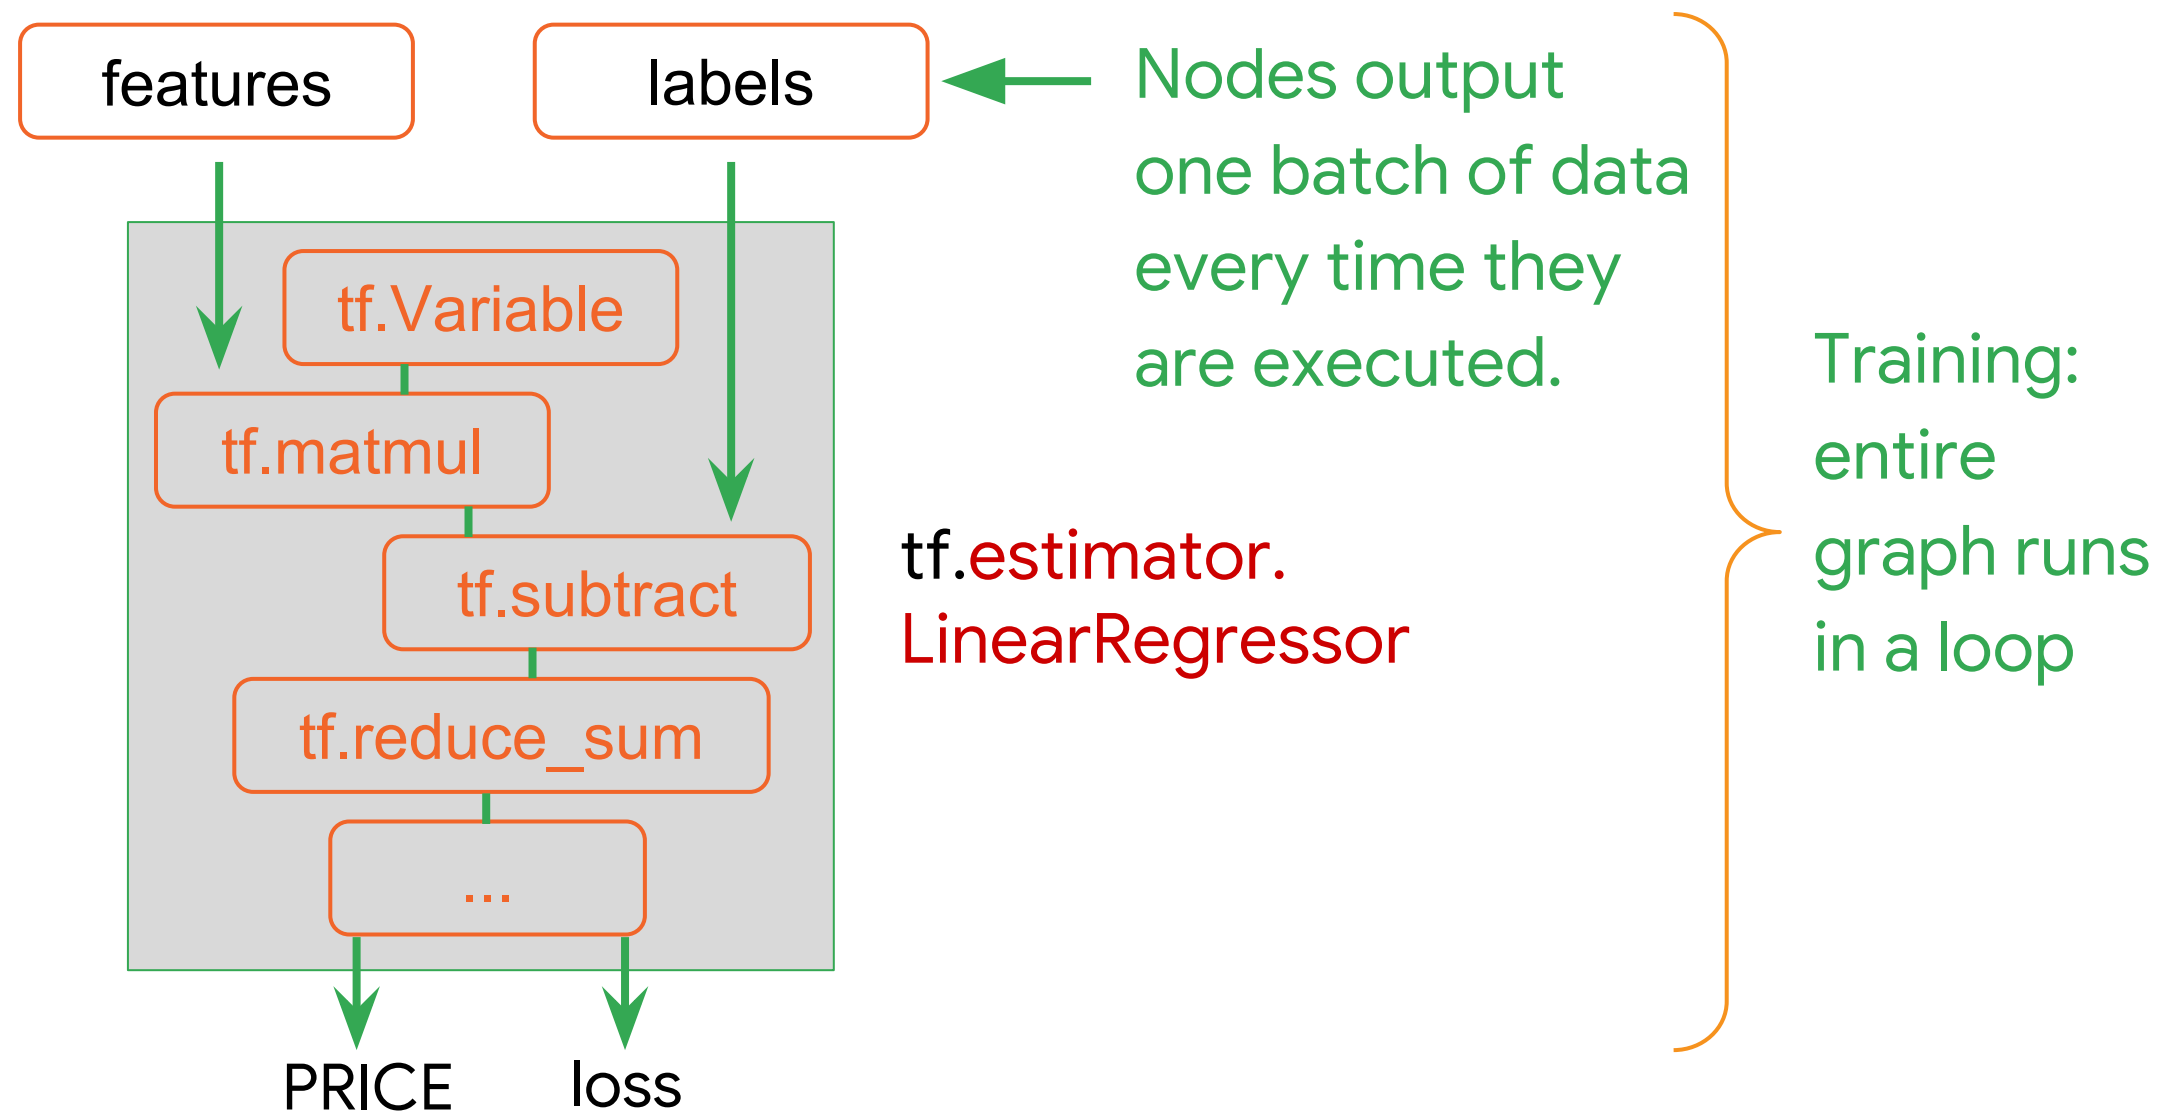
```
model.train(input_fn)
```

sq_footage property type PRICE in K\$

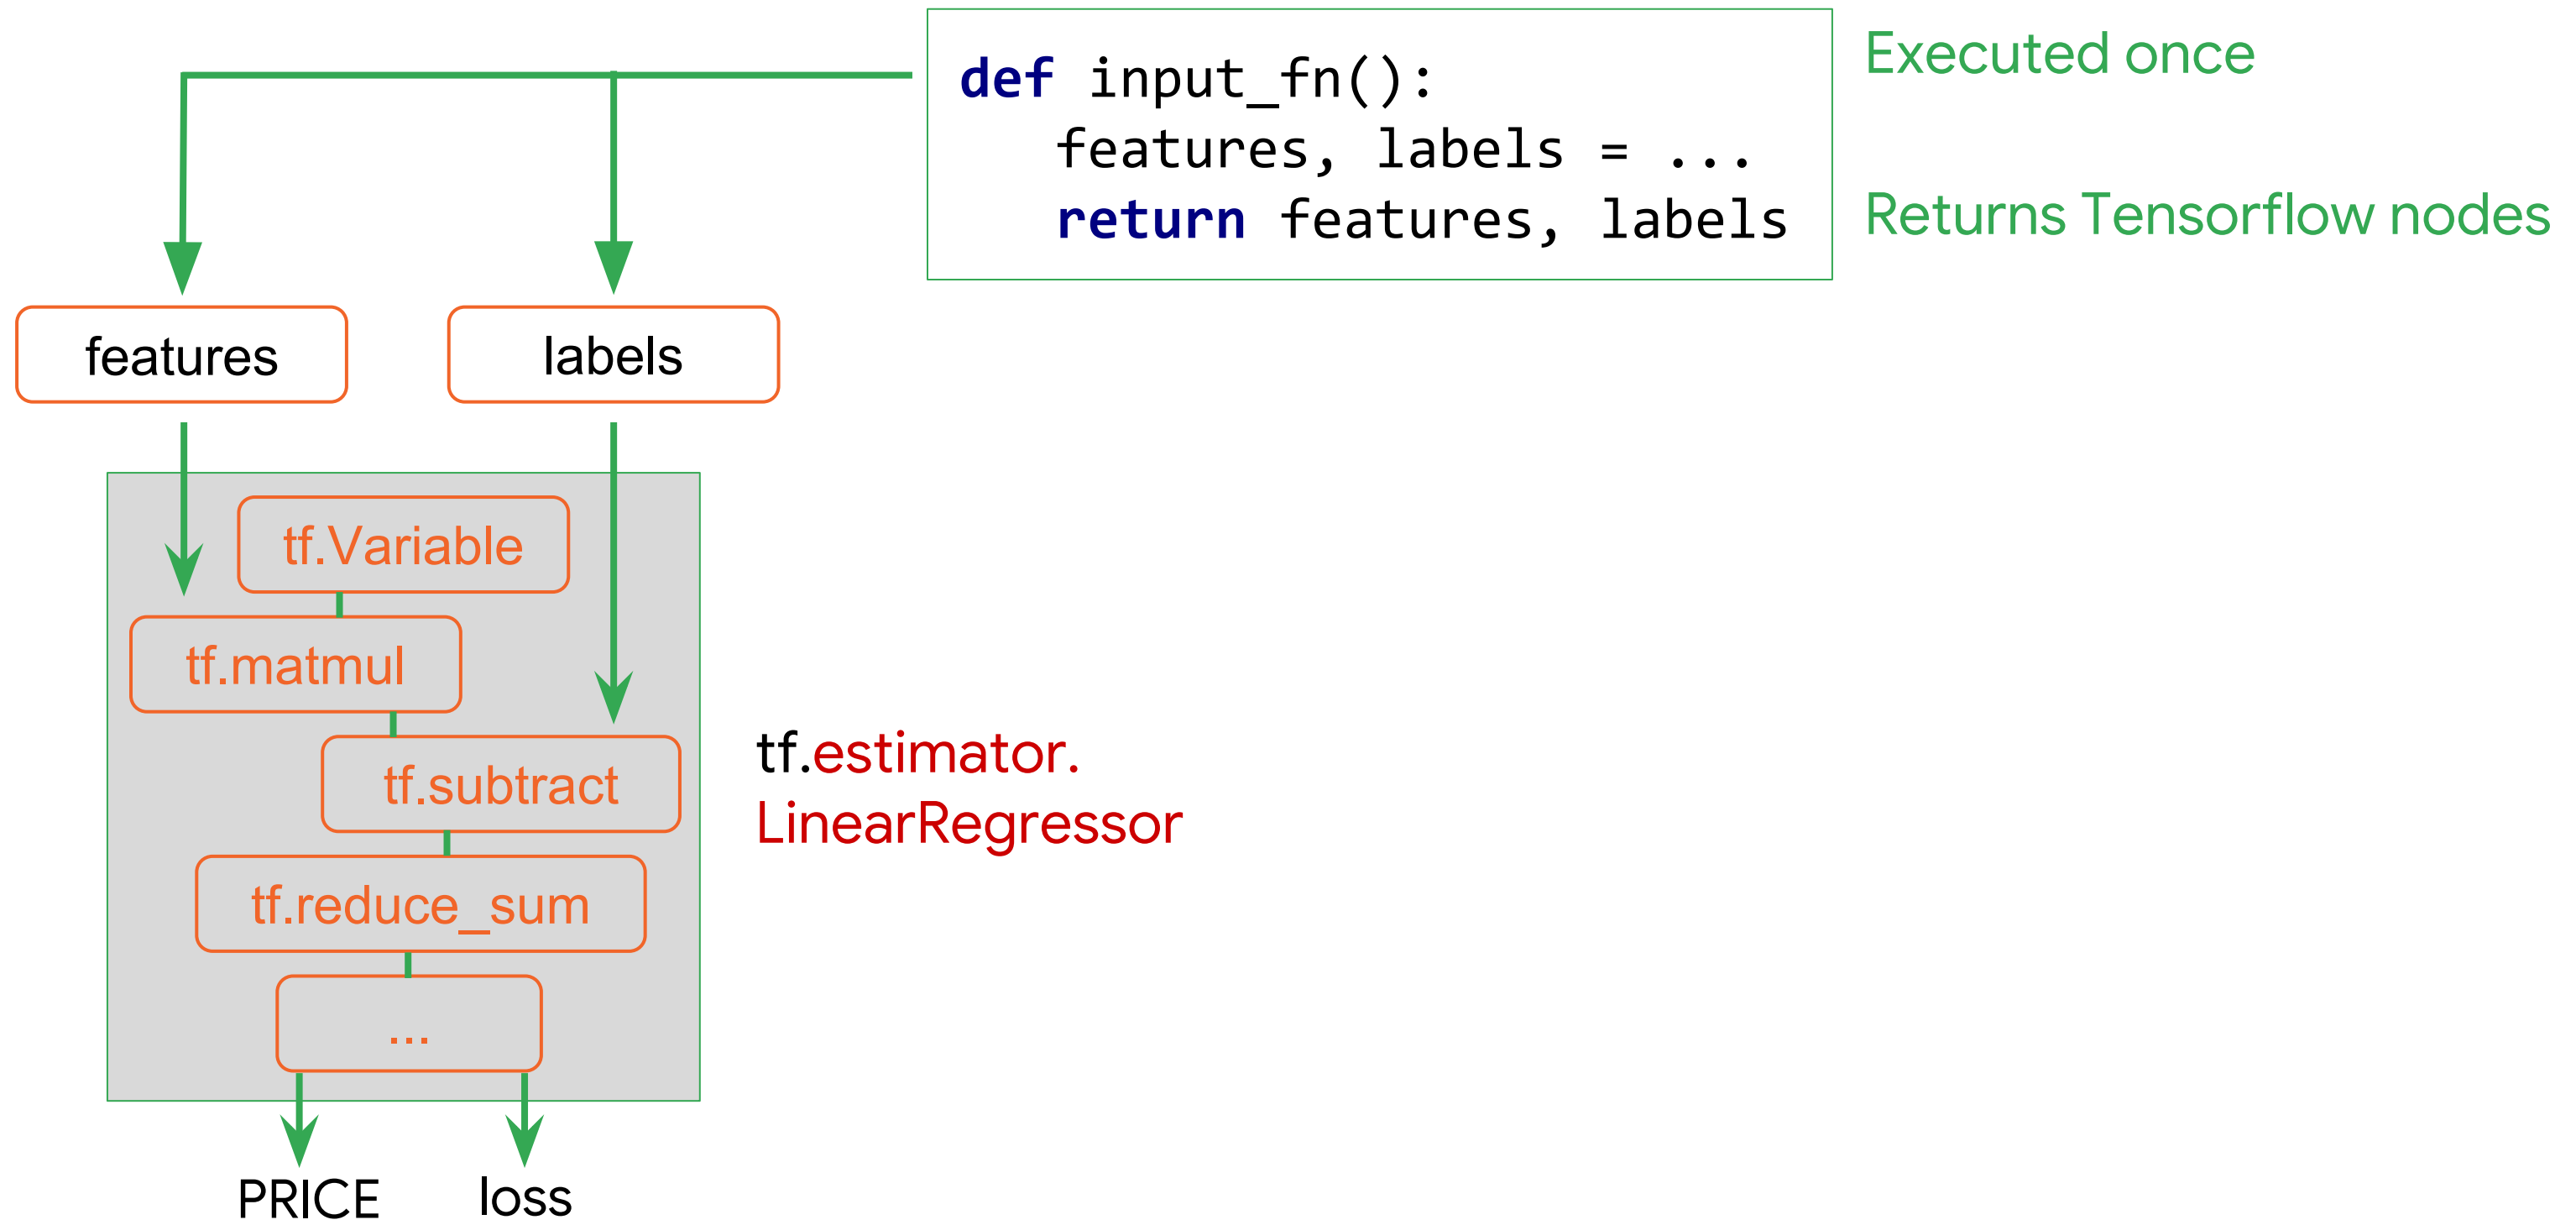
1001,	house,	501
2001,	house,	1001
3001,	house,	1501
1001,	apt,	701
2001,	apt,	1301
3001,	apt,	1901
1101,	house,	526
2101,	house,	1026

All the tf. commands
that you write in Python
do not actually
process any data, they
just build graphs

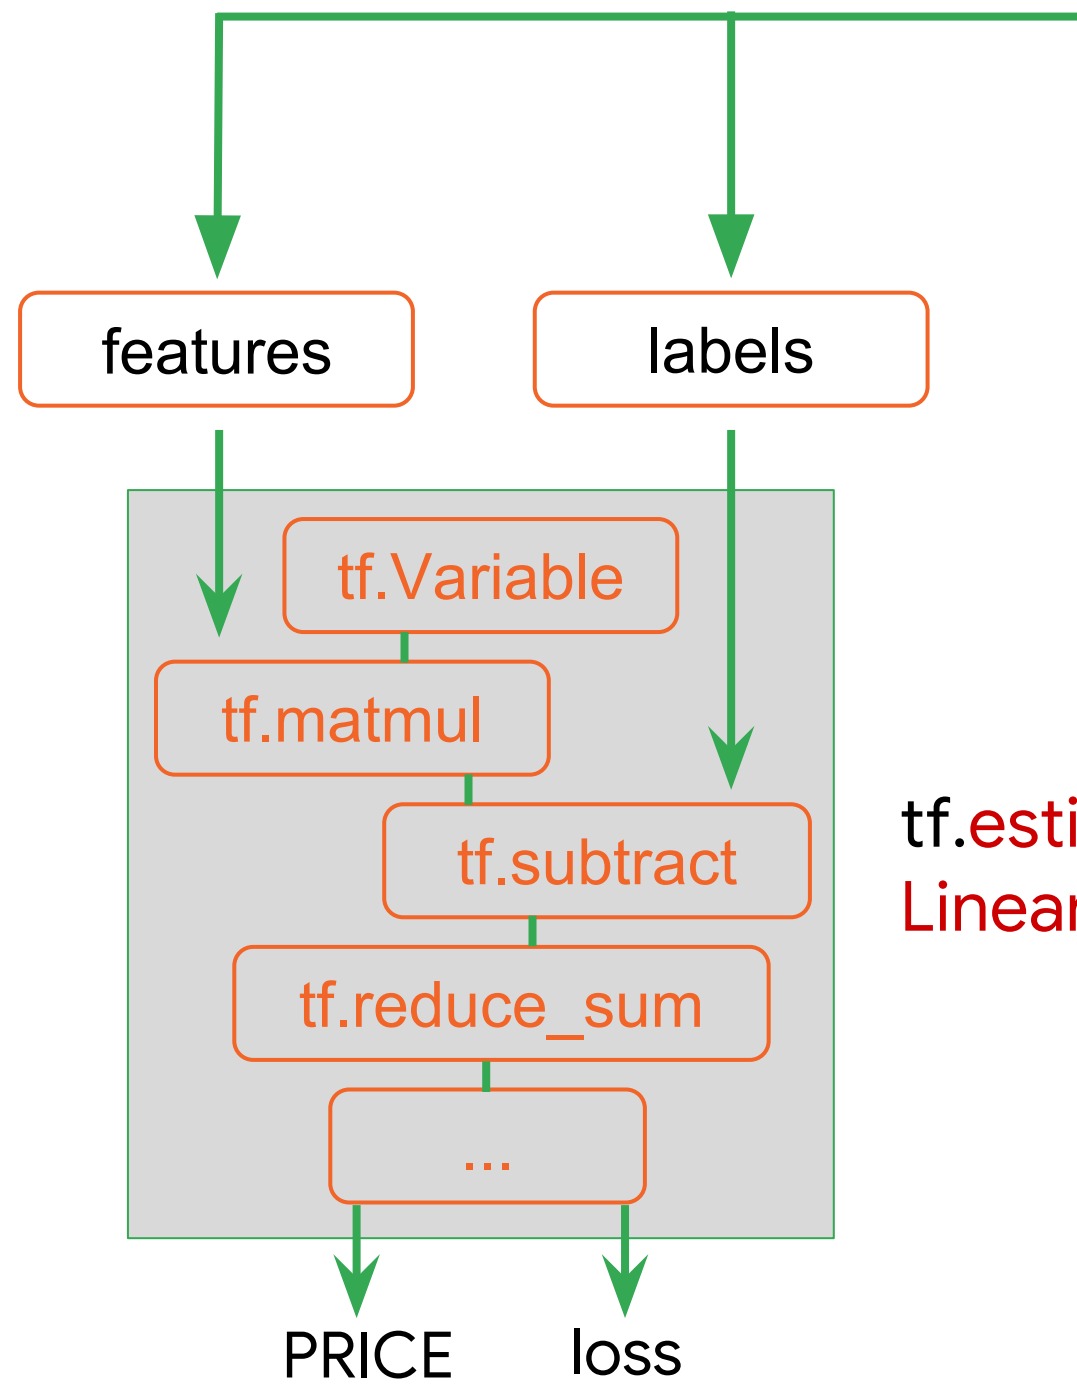
Under the hood: what does an input function do ?



Under the hood: what does an input function do ?



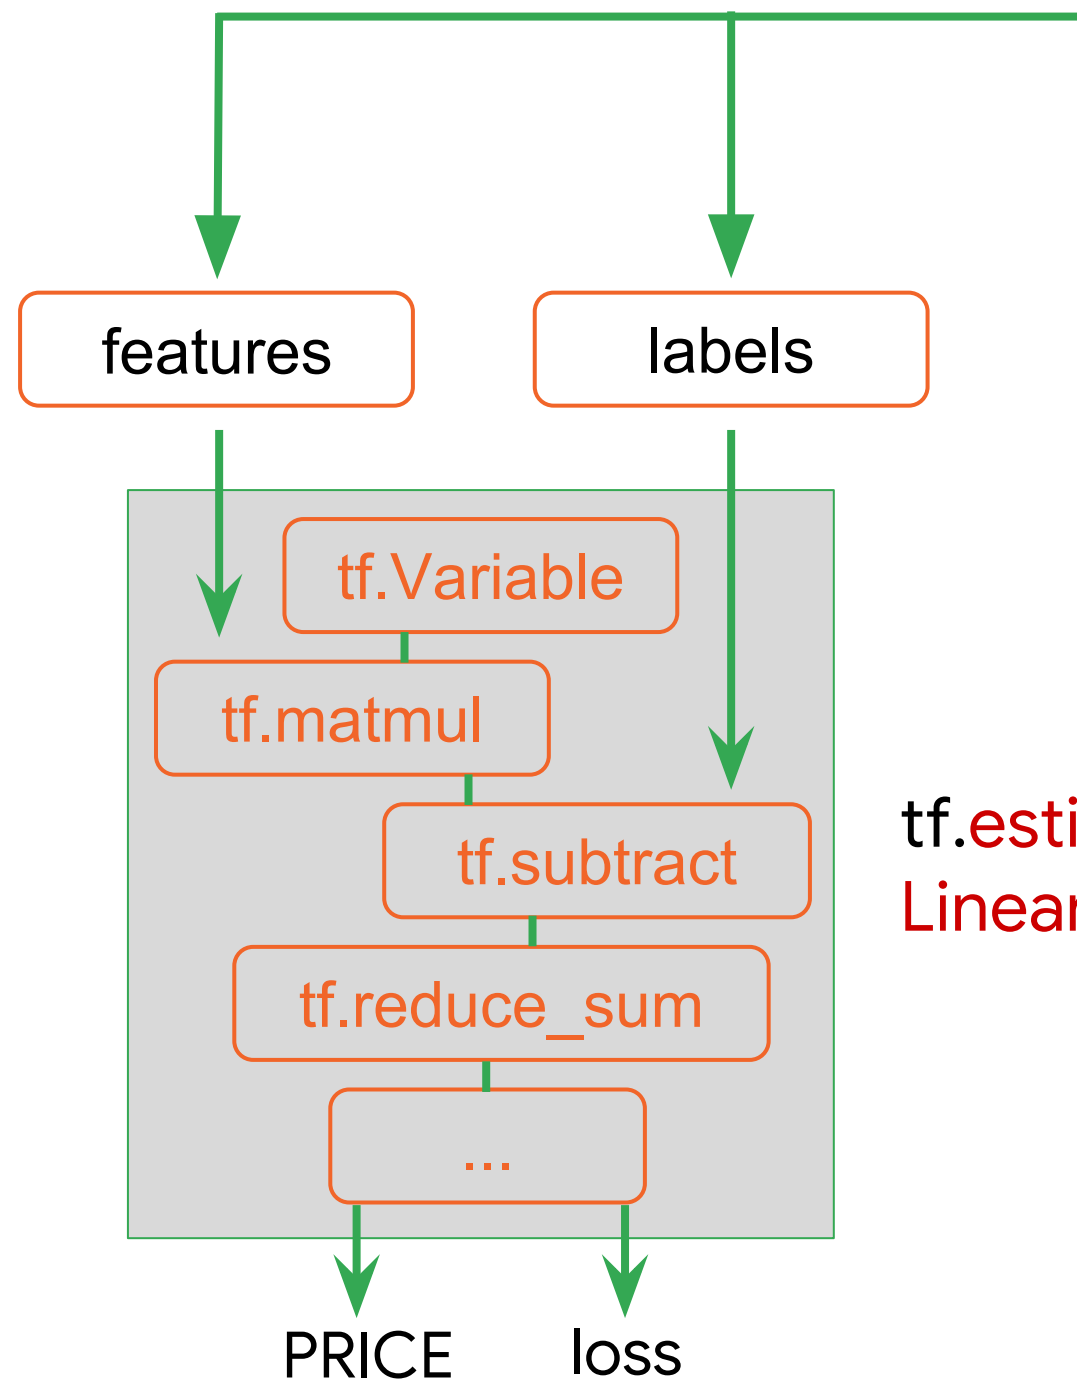
Under the hood: Dataset API



```
def input_fn():  
    features, label =  
        dataset.make_one_shot_iterator().get_next()  
    return features, labels
```

`tf.estimator.`
`LinearRegressor`

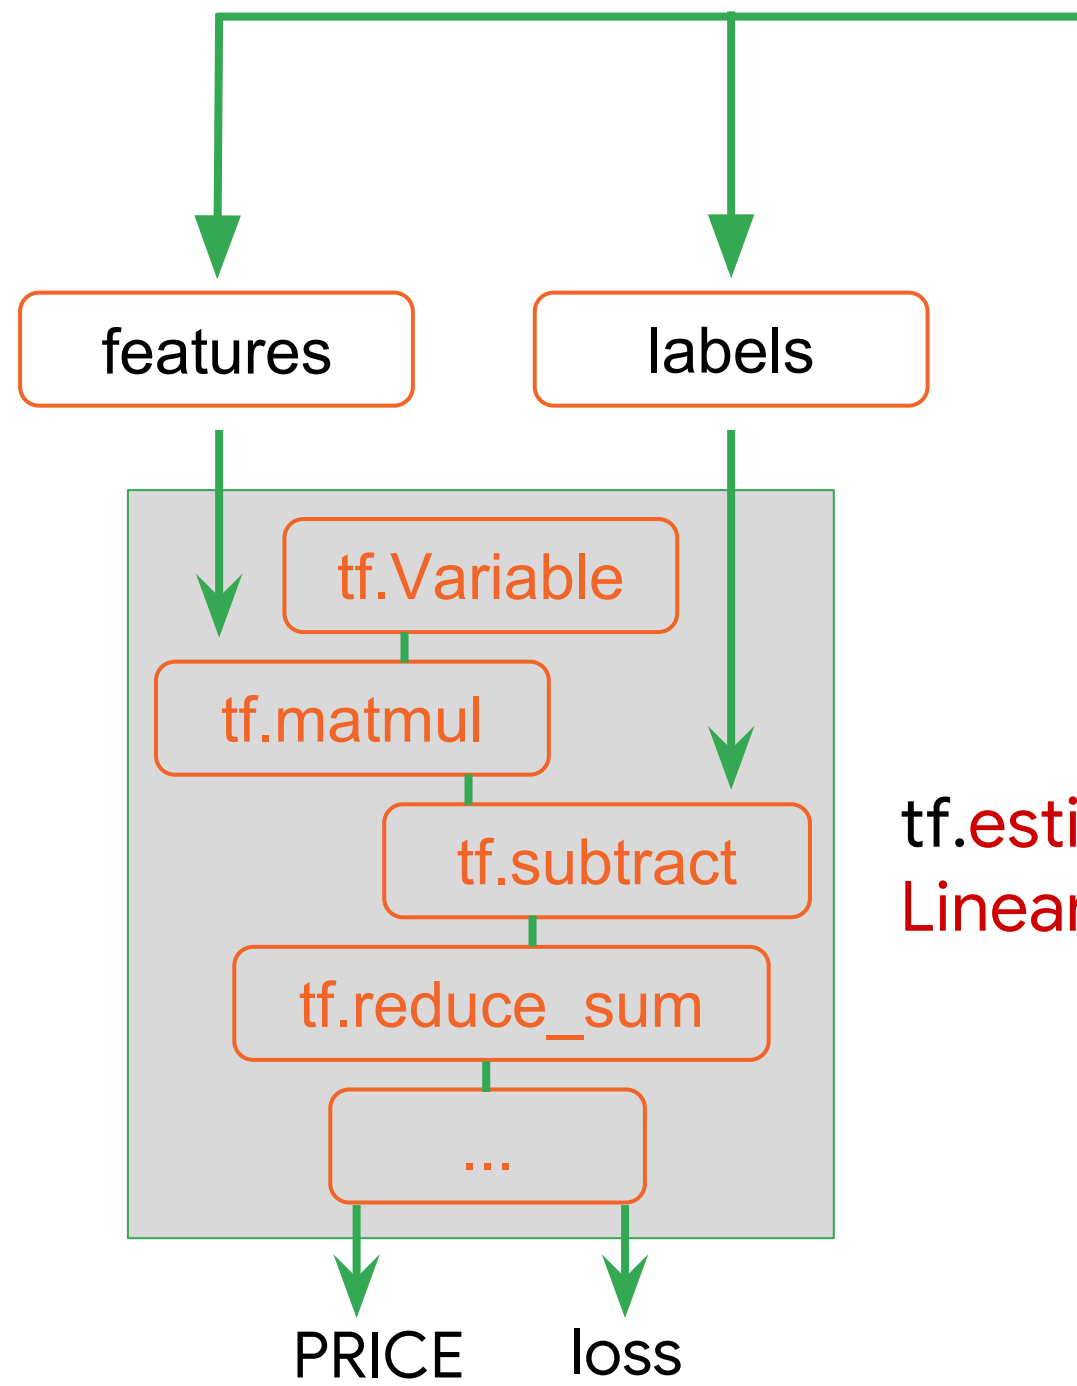
Under the hood: Dataset API



```
def input_fn():  
    features, label =  
        dataset.make_one_shot_iterator().get_next()  
    return features, labels
```

`tf.estimator.`
`LinearRegressor`

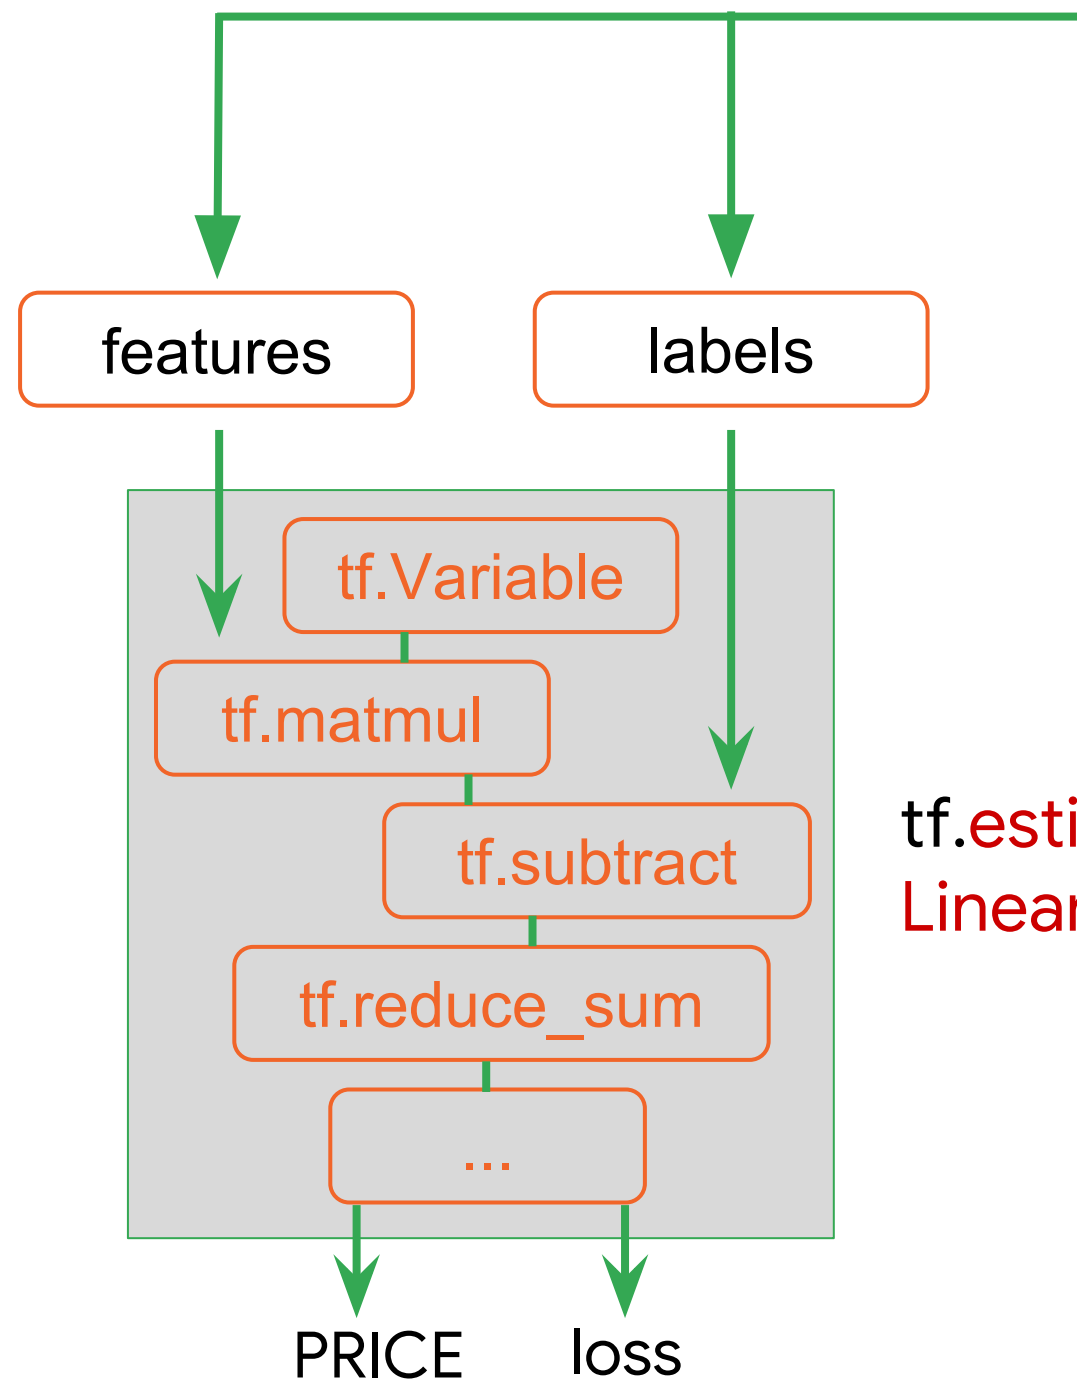
Common Misconceptions about input_fn



```
def input_fn():  
    features, label =  
        dataset.make_one_shot_iterator().get_next()  
    return features, labels
```

`tf.estimator.`
`LinearRegressor`

Common Misconceptions about input_fn

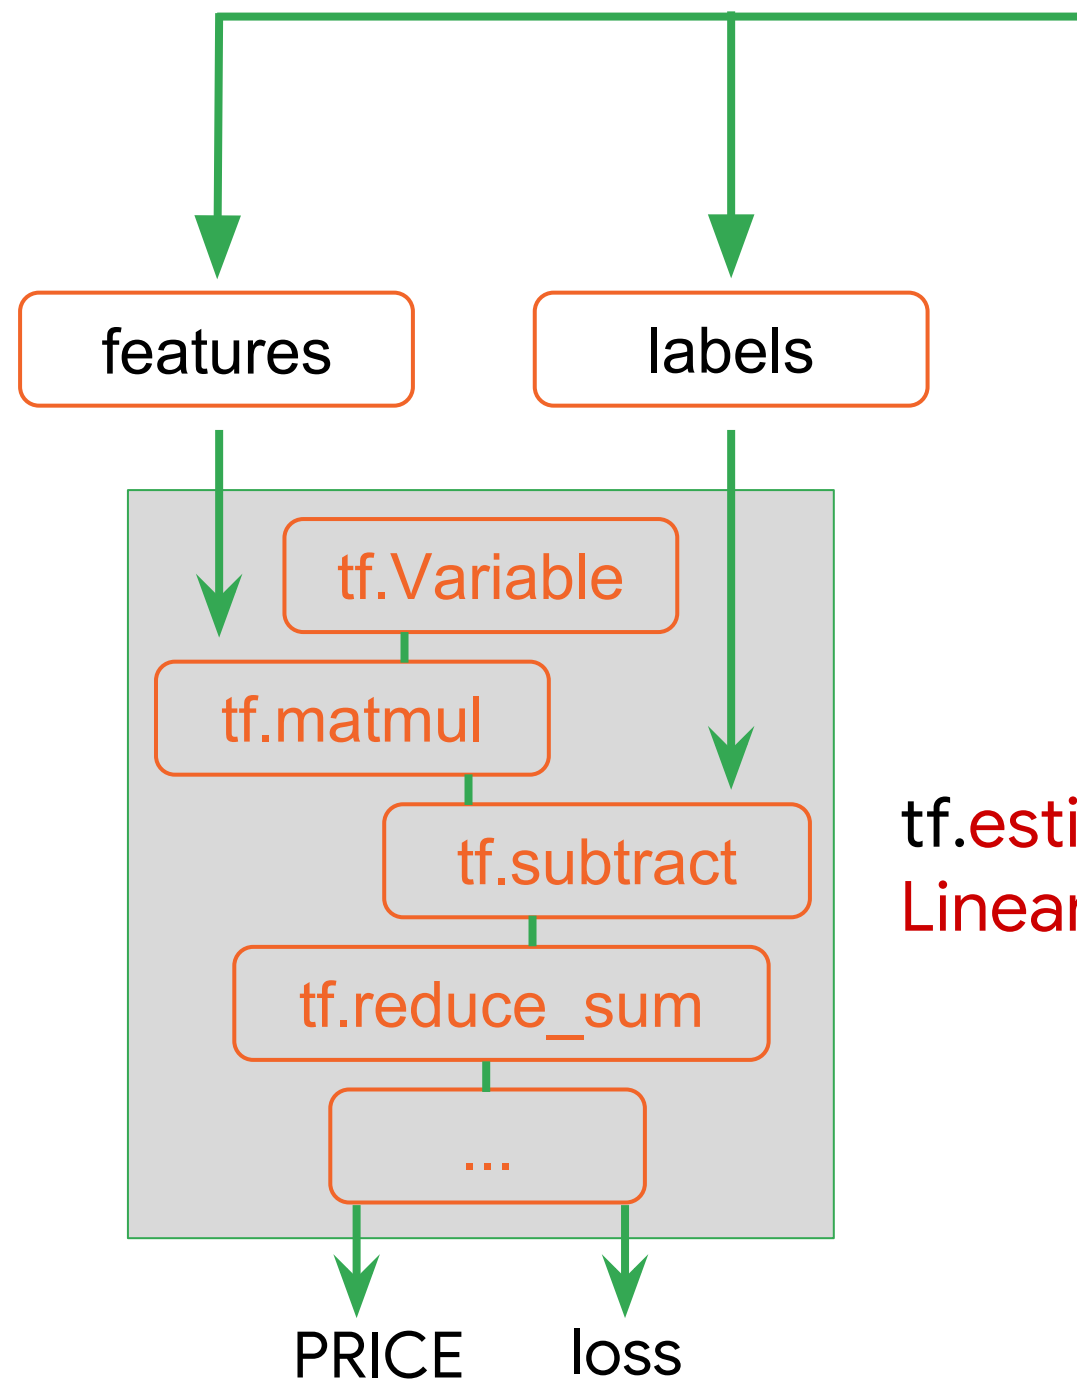


```
def input_fn():  
    features, label =  
        dataset.make_one_shot_iterator().get_next()  
    return features, labels
```

1. Input functions called only once

`tf.estimator.`
`LinearRegressor`

Common Misconceptions about input_fn



```
def input_fn():  
    features, label =  
        dataset.make_one_shot_iterator().get_next()  
    return features, labels
```

1. Input functions called only once
2. Input functions return Tf nodes (not data)

tf.estimator.
LinearRegressor

Read one CSV file using TextLineDataset

```
def decode_line(row):  
    cols = tf.decode_csv(row, record_defaults=[[0], ['house'], [0]])  
    features = {'sq_footage': cols[0], 'type': cols[1]}  
    label = cols[2] # price  
    return features, label
```

```
dataset = tf.data.TextLineDataset("train_1.csv") \  
    .map(decode_line)
```

```
dataset = dataset.shuffle(1000) \  
    .repeat(15) \  
    .batch(128)
```

```
def input_fn():  
    features, label = dataset.make_one_shot_iterator().get_next()  
    return features, label
```

```
model.train(input_fn)
```

property type
sq_footage PRICE in K\$

1001,	house,	501
2001,	house,	1001
3001,	house,	1501
1001,	apt,	701
2001,	apt,	1301
3001,	apt,	1901
1101,	house,	526
2101,	house,	1026

Read one CSV file using TextLineDataset

```
def decode_line(row):  
    cols = tf.decode_csv(row, record_defaults=[[0], ['house'], [0]])  
    features = {'sq_footage': cols[0], 'type': cols[1]}  
    label = cols[2] # price  
    return features, label
```

```
dataset = tf.data.TextLineDataset("train_1.csv") \  
    .map(decode_line)
```

```
dataset = dataset.shuffle(1000) \  
    .repeat(15) \  
    .batch(128)
```

```
def input_fn():  
    features, label = dataset.make_one_shot_iterator().get_next()  
    return features, label
```

```
model.train(input_fn)
```

property type
sq_footage PRICE in K\$

1001,	house,	501
2001,	house,	1001
3001,	house,	1501
1001,	apt,	701
2001,	apt,	1301
3001,	apt,	1901
1101,	house,	526
2101,	house,	1026

Read one CSV file using TextLineDataset

```
def decode_line(row):  
    cols = tf.decode_csv(row, record_defaults=[[0], ['house'], [0]])  
    features = {'sq_footage': cols[0], 'type': cols[1]}  
    label = cols[2] # price  
    return features, label
```

```
dataset = tf.data.TextLineDataset("train_1.csv") \  
    .map(decode_line)
```

```
dataset = dataset.shuffle(1000) \  
    .repeat(15) \  
    .batch(128)
```

```
def input_fn():  
    features, label = dataset.make_one_shot_iterator().get_next()  
    return features, label
```

```
model.train(input_fn)
```

property type
sq_footage PRICE in K\$

1001,	house,	501
2001,	house,	1001
3001,	house,	1501
1001,	apt,	701
2001,	apt,	1301
3001,	apt,	1901
1101,	house,	526
2101,	house,	1026

Read one CSV file using TextLineDataset

```
def decode_line(row):  
    cols = tf.decode_csv(row, record_defaults=[[0], ['house'], [0]])  
    features = {'sq_footage': cols[0], 'type': cols[1]}  
    label = cols[2] # price  
    return features, label
```

```
dataset = tf.data.TextLineDataset("train_1.csv") \  
    .map(decode_line)
```

```
dataset = dataset.shuffle(1000) \  
    .repeat(15) \  
    .batch(128)
```

```
def input_fn():  
    features, label = dataset.make_one_shot_iterator().get_next()  
    return features, label
```

```
model.train(input_fn)
```

property type
sq_footage PRICE in K\$

1001,	house,	501
2001,	house,	1001
3001,	house,	1501
1001,	apt,	701
2001,	apt,	1301
3001,	apt,	1901
1101,	house,	526
2101,	house,	1026

Read one CSV file using TextLineDataset

```
def decode_line(row):  
    cols = tf.decode_csv(row, record_defaults=[[0], ['house'], [0]])  
    features = {'sq_footage': cols[0], 'type': cols[1]}  
    label = cols[2] # price  
    return features, label
```

```
dataset = tf.data.TextLineDataset("train_1.csv") \  
    .map(decode_line)
```

```
dataset = dataset.shuffle(1000) \  
    .repeat(15) \  
    .batch(128)
```

```
def input_fn():  
    features, label = dataset.make_one_shot_iterator().get_next()  
    return features, label
```

```
model.train(input_fn)
```

property type
sq_footage PRICE in K\$

1001,	house,	501
2001,	house,	1001
3001,	house,	1501
1001,	apt,	701
2001,	apt,	1301
3001,	apt,	1901
1101,	house,	526
2101,	house,	1026

Read a set of sharded CSV files using TextLineDataset

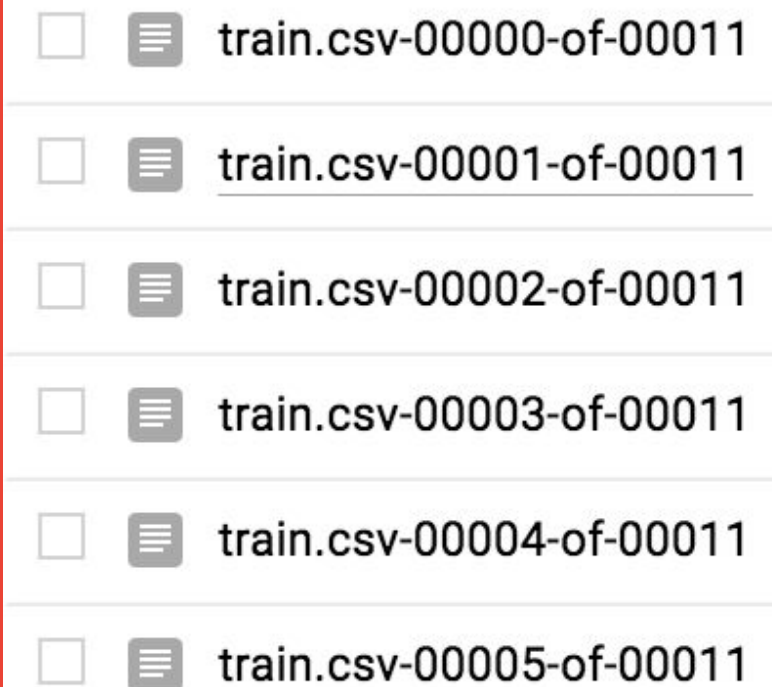
```
def decode_line(row):  
    cols = tf.decode_csv(row, record_defaults=[[0], ['house'], [0]])  
    features = {'sq_footage': cols[0], 'type': cols[1]}  
    label = cols[2] # price  
    return features, label
```







```
dataset = tf.data.Dataset.list_files("train.csv-*") \  
    .flat_map(tf.data.TextLineDataset) \  
    .map(decode_line)
```

```
dataset = dataset.shuffle(1000) \  
    .repeat(15) \  
    .batch(128)
```

```
def input_fn():  
    return dataset.make_one_shot_iterator().get_next()
```

```
model.train(input_fn)
```



<input type="checkbox"/>		train.csv-00000-of-00011
<input type="checkbox"/>		train.csv-00001-of-00011
<input type="checkbox"/>		train.csv-00002-of-00011
<input type="checkbox"/>		train.csv-00003-of-00011
<input type="checkbox"/>		train.csv-00004-of-00011
<input type="checkbox"/>		train.csv-00005-of-00011

Read a set of sharded CSV files using TextLineDataset







```
def decode_line(row):  
    cols = tf.decode_csv(row, record_defaults=[[0], ['house'], [0]])  
    features = {'sq_footage': cols[0], 'type': cols[1]}  
    label = cols[2] # price  
    return features, label
```

```
dataset = tf.data.Dataset.list_files("train.csv-*") \  
    .flat_map(tf.data.TextLineDataset) \  
    .map(decode_line)
```

```
dataset = dataset.shuffle(1000) \  
    .repeat(15) \  
    .batch(128)
```

```
def input_fn():  
    return dataset.make_one_shot_iterator().get_next()
```

```
model.train(input_fn)
```

<input type="checkbox"/>		train.csv-00000-of-00011
<input type="checkbox"/>		train.csv-00001-of-00011
<input type="checkbox"/>		train.csv-00002-of-00011
<input type="checkbox"/>		train.csv-00003-of-00011
<input type="checkbox"/>		train.csv-00004-of-00011
<input type="checkbox"/>		train.csv-00005-of-00011

Read a set of sharded CSV files using TextLineDataset







```
def decode_line(row):  
    cols = tf.decode_csv(row, record_defaults=[[0], ['house'], [0]])  
    features = {'sq_footage': cols[0], 'type': cols[1]}  
    label = cols[2] # price  
    return features, label
```

```
dataset = tf.data.Dataset.list_files("train.csv-*") \  
    .flat_map(tf.data.TextLineDataset) \  
    .map(decode_line)
```

```
dataset = dataset.shuffle(1000) \  
    .repeat(15) \  
    .batch(128)
```

```
def input_fn():  
    return dataset.make_one_shot_iterator().get_next()
```

```
model.train(input_fn)
```

<input type="checkbox"/>		train.csv-00000-of-00011
<input type="checkbox"/>		train.csv-00001-of-00011
<input type="checkbox"/>		train.csv-00002-of-00011
<input type="checkbox"/>		train.csv-00003-of-00011
<input type="checkbox"/>		train.csv-00004-of-00011
<input type="checkbox"/>		train.csv-00005-of-00011

Read a set of sharded CSV files using TextLineDataset







```
def decode_line(row):  
    cols = tf.decode_csv(row, record_defaults=[[0], ['house'], [0]])  
    features = {'sq_footage': cols[0], 'type': cols[1]}  
    label = cols[2] # price  
    return features, label
```

```
dataset = tf.data.Dataset.list_files("train.csv-*") \  
    .flat_map(tf.data.TextLineDataset) \  
    .map(decode_line)
```

```
dataset = dataset.shuffle(1000) \  
    .repeat(15) \  
    .batch(128)
```

```
def input_fn():  
    return dataset.make_one_shot_iterator().get_next()
```

```
model.train(input_fn)
```

<input type="checkbox"/>		train.csv-00000-of-00011
<input type="checkbox"/>		train.csv-00001-of-00011
<input type="checkbox"/>		train.csv-00002-of-00011
<input type="checkbox"/>		train.csv-00003-of-00011
<input type="checkbox"/>		train.csv-00004-of-00011
<input type="checkbox"/>		train.csv-00005-of-00011

Read a set of sharded CSV files using TextLineDataset







```
def decode_line(row):  
    cols = tf.decode_csv(row, record_defaults=[[0], ['house'], [0]])  
    features = {'sq_footage': cols[0], 'type': cols[1]}  
    label = cols[2] # price  
    return features, label
```

```
dataset = tf.data.Dataset.list_files("train.csv-*") \  
    .flat_map(tf.data.TextLineDataset) \  
    .map(decode_line)
```

```
dataset = dataset.shuffle(1000) \  
    .repeat(15) \  
    .batch(128)
```

```
def input_fn():  
    return dataset.make_one_shot_iterator().get_next()
```

```
model.train(input_fn)
```

<input type="checkbox"/>		train.csv-00000-of-00011
<input type="checkbox"/>		train.csv-00001-of-00011
<input type="checkbox"/>		train.csv-00002-of-00011
<input type="checkbox"/>		train.csv-00003-of-00011
<input type="checkbox"/>		train.csv-00004-of-00011
<input type="checkbox"/>		train.csv-00005-of-00011

The real benefit of Dataset is
that you can do more than just
ingest data

```
dataset =  
tf.data.TextLineDataset(filename)\  
    .skip(num_header_lines)\  
    .map(add_key)\  
    .map(decode_csv)\  
    .map(lambda feats,  
          labels: preproc(feats), labels)\  
    .filter(is_valid)\  
    .cache()
```

LAB


Scaling up TensorFlow ingest
using batching



Big Jobs, Distributed training

Martin Gerner

Real World ML Models

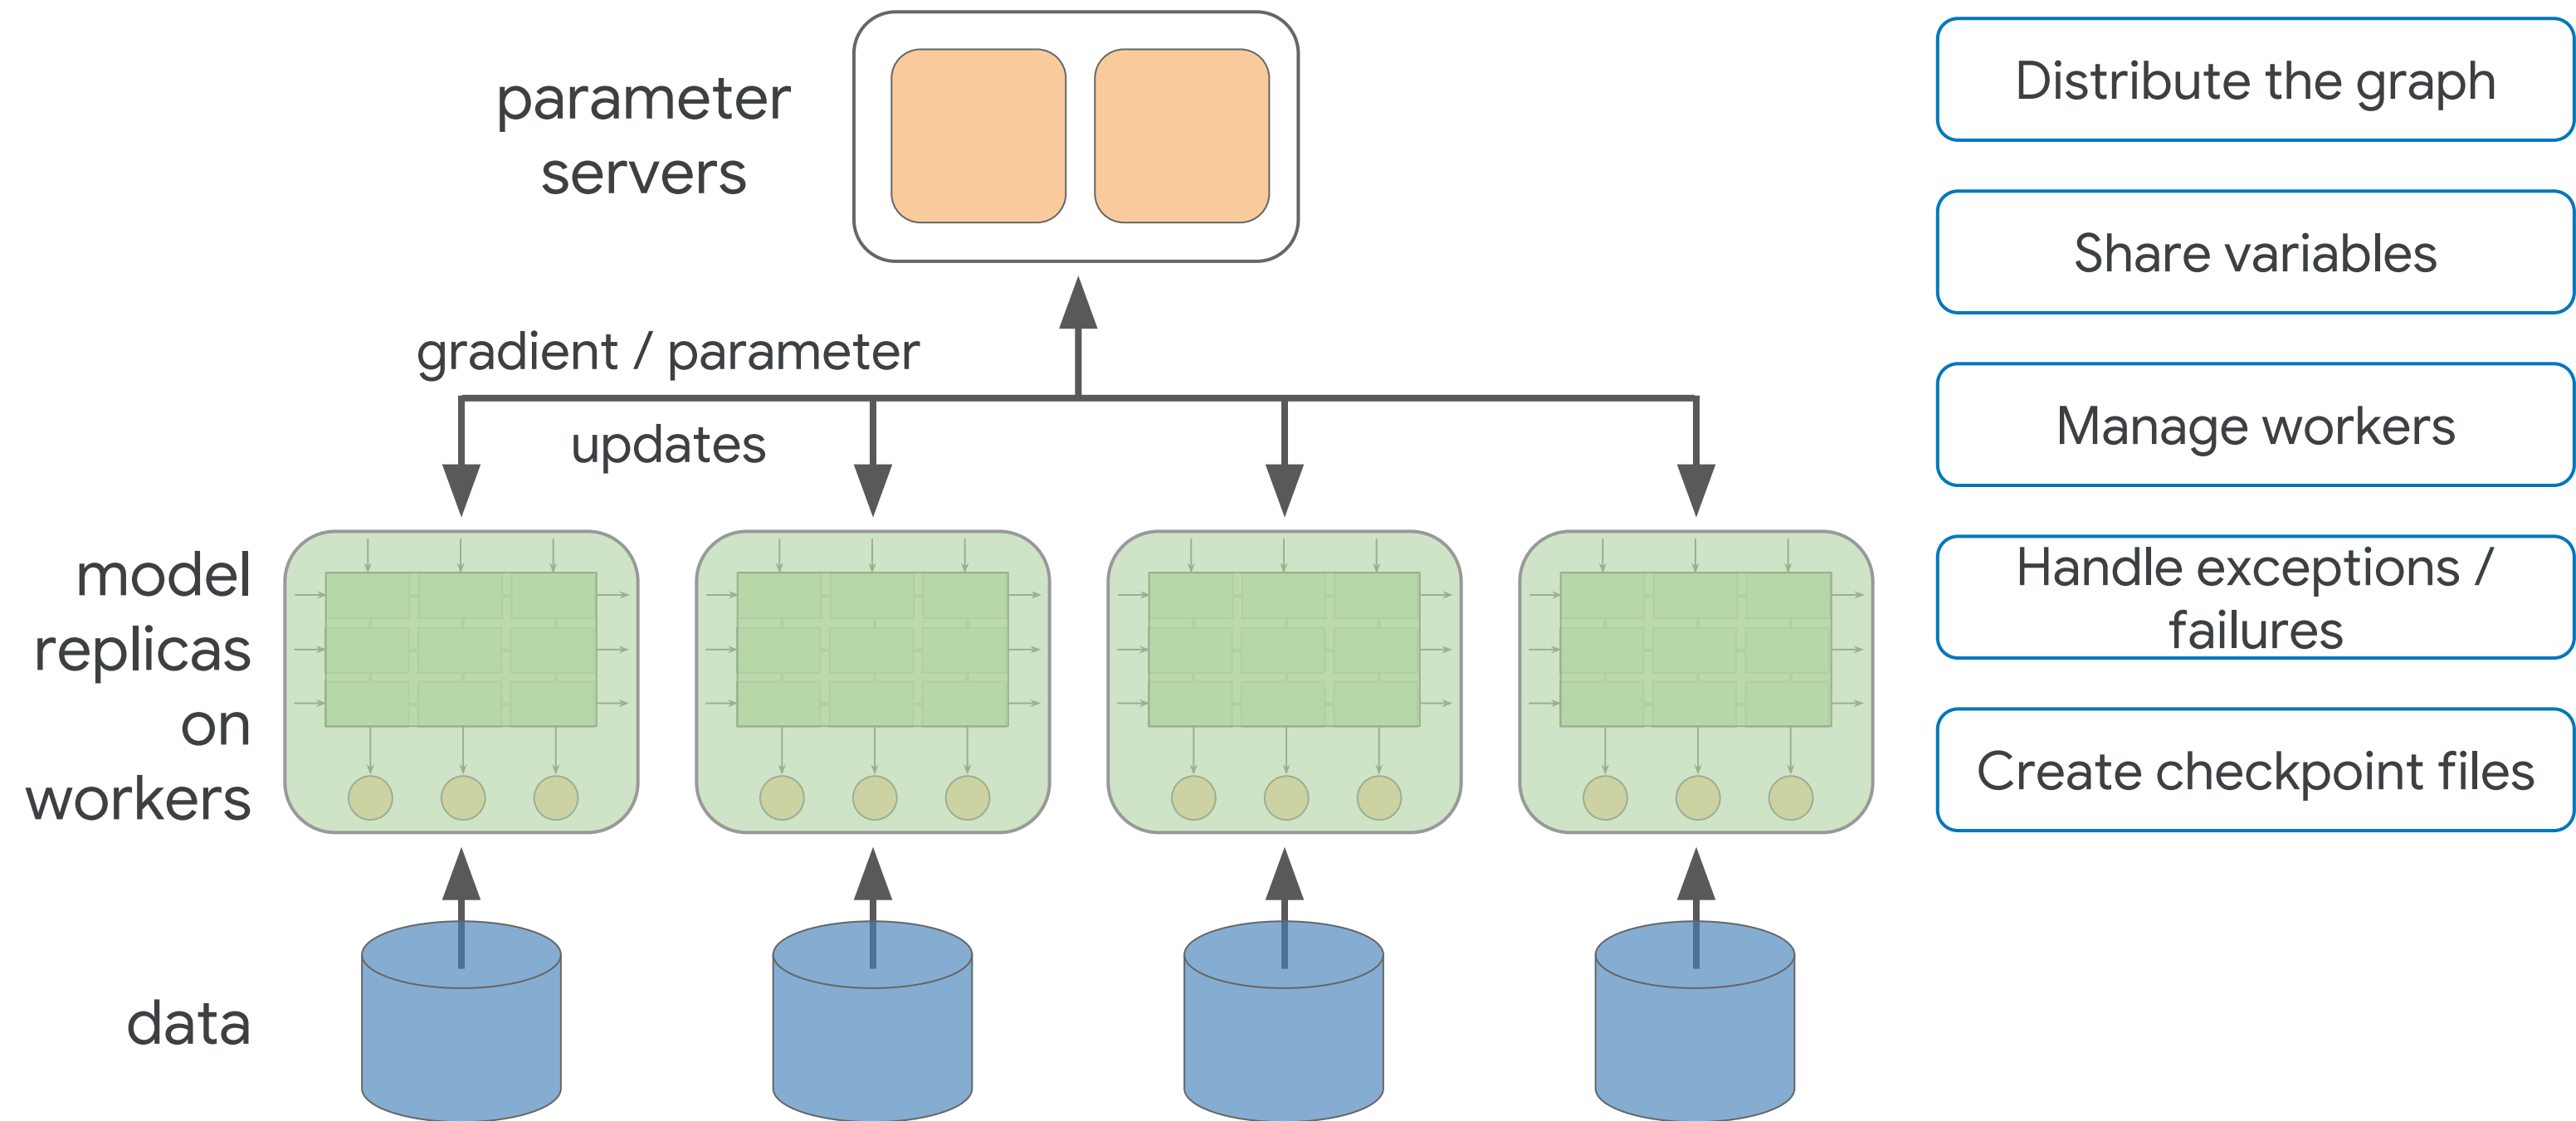
Problem	Solution
 Out of memory data	Use the Dataset API
Distribution	Use train_and_evaluate
Need to evaluate during training	?
Deployments that scale	?

`estimator.train_and_evaluate`
is the preferred method for training
real-world models

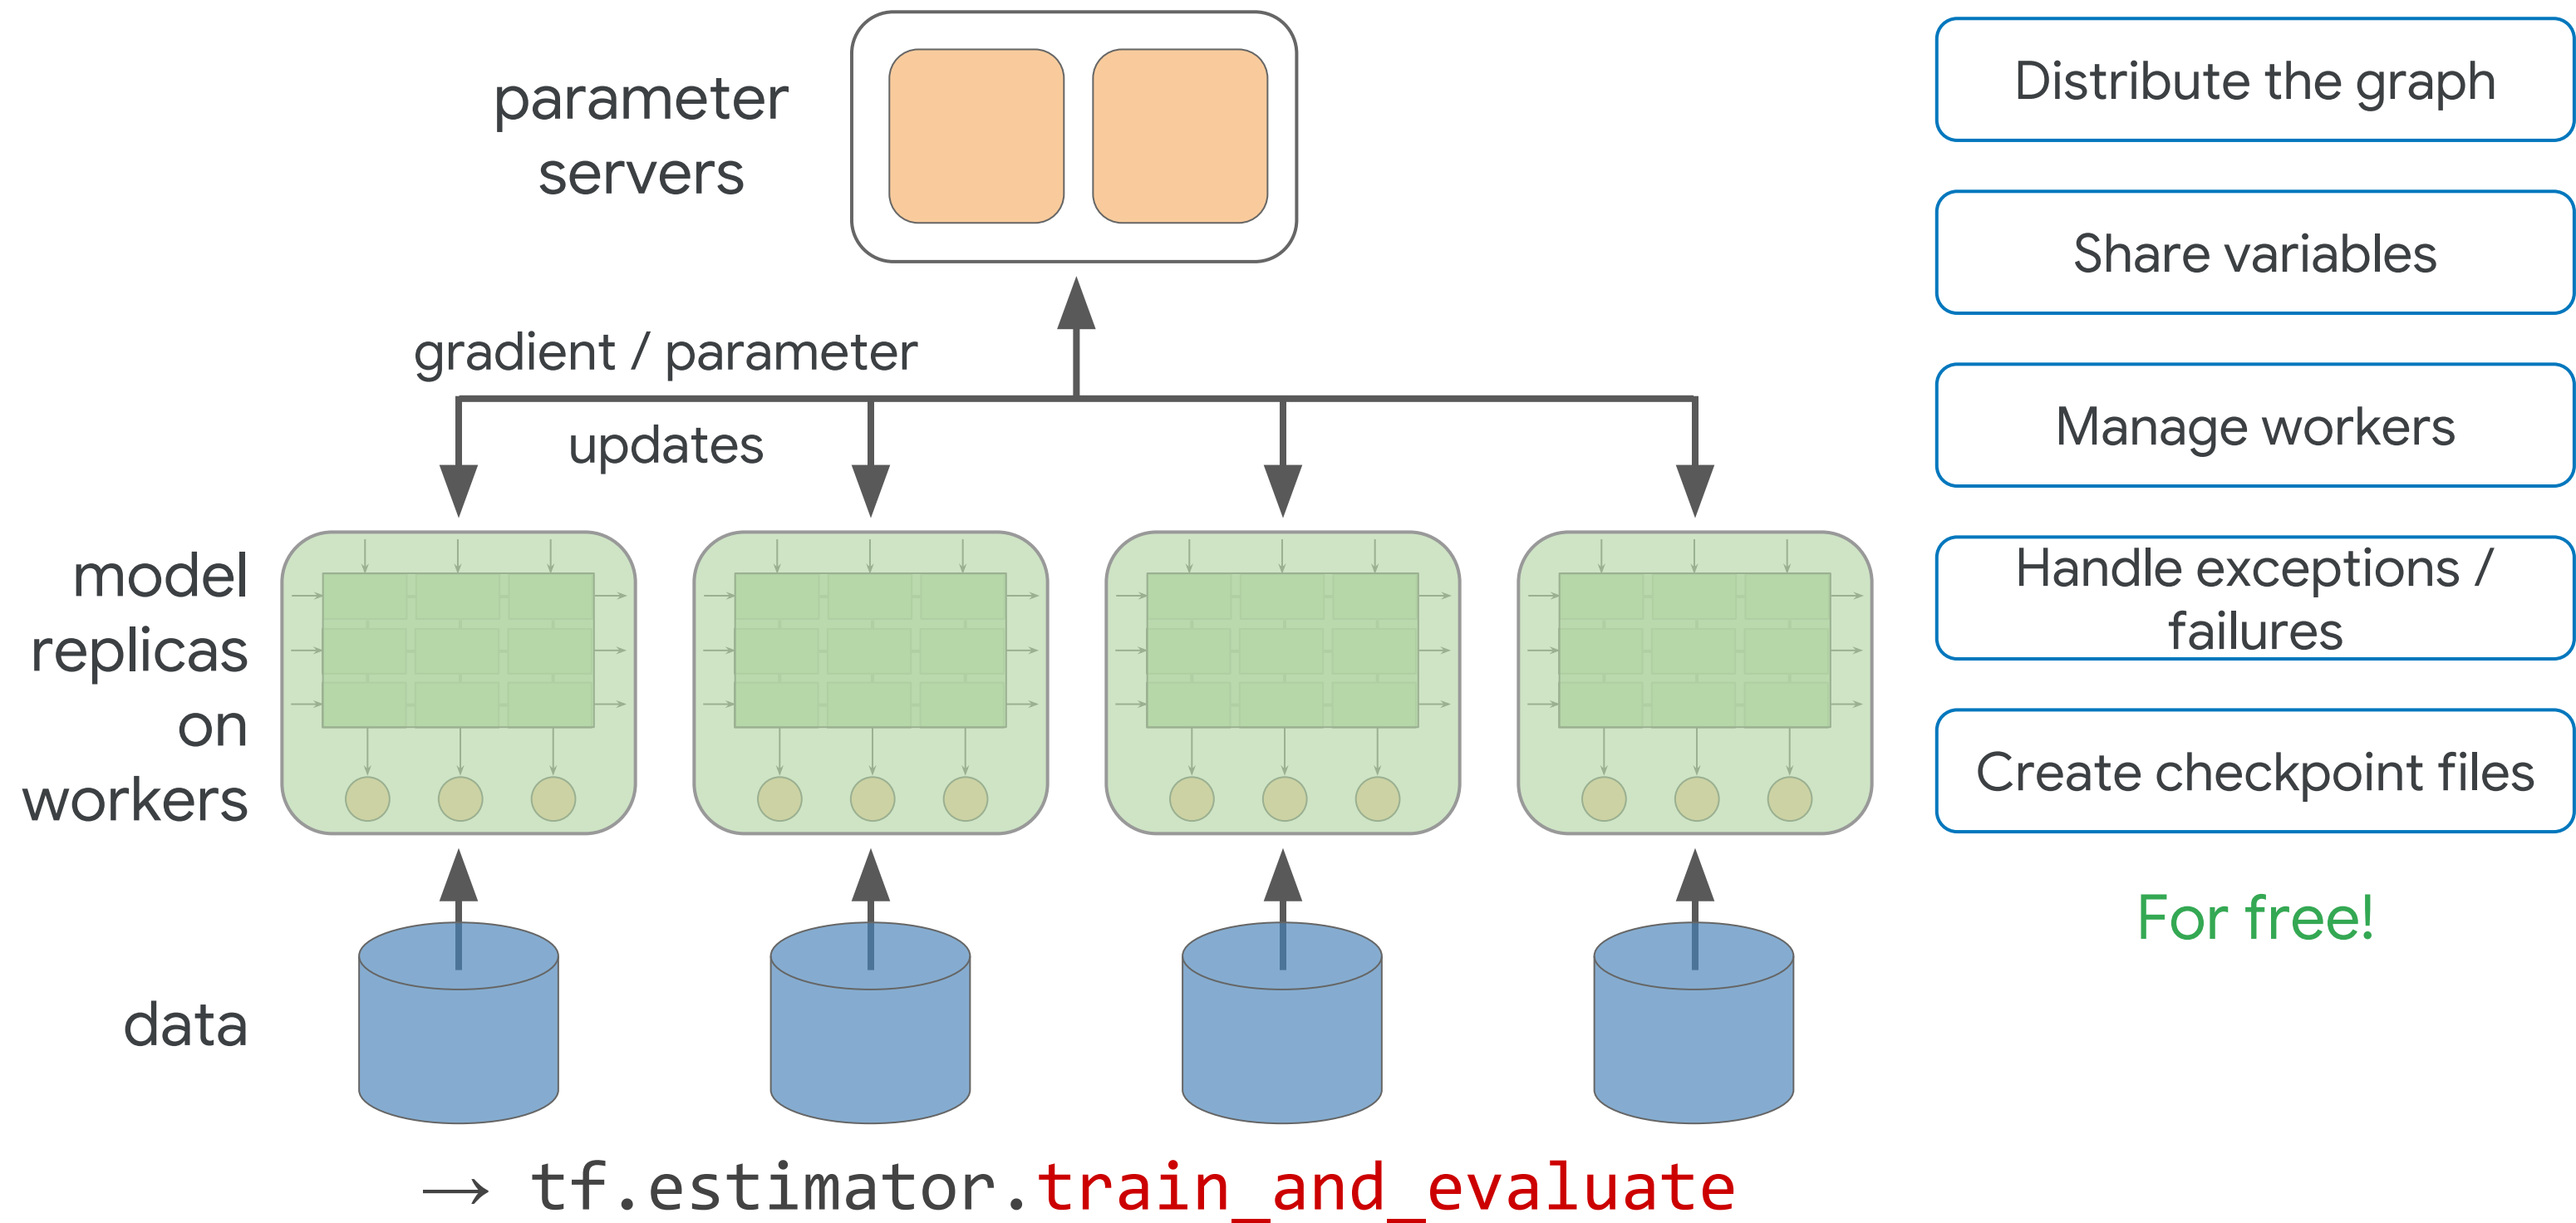
```
estimator = tf.estimator.LinearRegressor(...)  
tf.estimator.train_and_evaluate(estimator, ...)
```

data parallelism =
replicate your model on
multiple workers

Distributed training using “data parallelism”



Distributed training using “data parallelism”



`estimator.train_and_evaluate` is the preferred method for training real-world models

```
estimator = tf.estimator.LinearRegressor(  
    feature_columns=featcols,  
    config=run_config)  
  
...  
  
tf.estimator.train_and_evaluate(estimator,  
    train_spec,  
    eval_spec)
```

You need:

1. Estimator

2. Run Config

3. Train Spec

4. Eval Spec

run_config tells the estimator where and how often
to write Checkpoints and Tensorboard logs
("summaries")

```
run_config = tf.estimator.RunConfig(  
    model_dir=output_dir,  
    save_summary_steps=100,  
    save_checkpoints_steps=2000)  
  
estimator = tf.estimator.LinearRegressor(config=run_config, ...)
```

The TrainSpec tells the estimator how to get training data

Use Datasets




```
train_spec =  
tf.estimator.TrainSpec(input_fn=train_input_fn, max_steps=50000)  
  
...  
  
tf.estimator.train_and_evaluate(estimator, train_spec, eval_spec)
```


The EvalSpec controls the evaluation and the checkpointing of the model since they happen at the same time

```
eval_spec =  
tf.estimator.EvalSpec(  
    input_fn=eval_input_fn,  
    steps=100, # evals on 100 batches  
    throttle_secs=600 # eval no more than every 10 min  
    exporters=...)
  
tf.estimator.train_and_evaluate(estimator, train_spec, eval_spec)
```

The EvalSpec controls the evaluation and the checkpointing of the model since they happen at the same time

Use Datasets



```
eval_spec =  
tf.estimator.EvalSpec(  
    input_fn=eval_input_fn,  
    steps=100, # evals on 100 batches  
    throttle_secs=600 # eval no more than every 10 min  
    exporters=...)   
  
tf.estimator.train_and_evaluate(estimator, train_spec, eval_spec)
```

The EvalSpec controls the evaluation and the checkpointing of the model since they happen at the same time

```
eval_spec =  
tf.estimator.EvalSpec(  
    input_fn=eval_input_fn,  
    steps=100, # evals on 100 batches  
    throttle_secs=600 # eval no more than every 10 min  
    exporters=...)   
  
tf.estimator.train_and_evaluate(estimator, train_spec, eval_spec)
```

Recap with all the code

```
run_config =  
tf.estimator.RunConfig(model_dir=output_dir, ...)  
  
estimator =  
tf.estimator.LinearRegressor(featcols, config=run_config)  
  
train_spec =  
tf.estimator.TrainSpec(input_fn=train_input_fn, max_steps=1000)  
  
export_latest =  
tf.estimator.LatestExporter(serving_input_receiver_fn=serving_input_fn)  
  
eval_spec =  
tf.estimator.EvalSpec(input_fn=eval_input_fn, exporters=export_latest)  
  
tf.estimator.train_and_evaluate(estimator, train_spec, eval_spec)
```

Recap with all the code

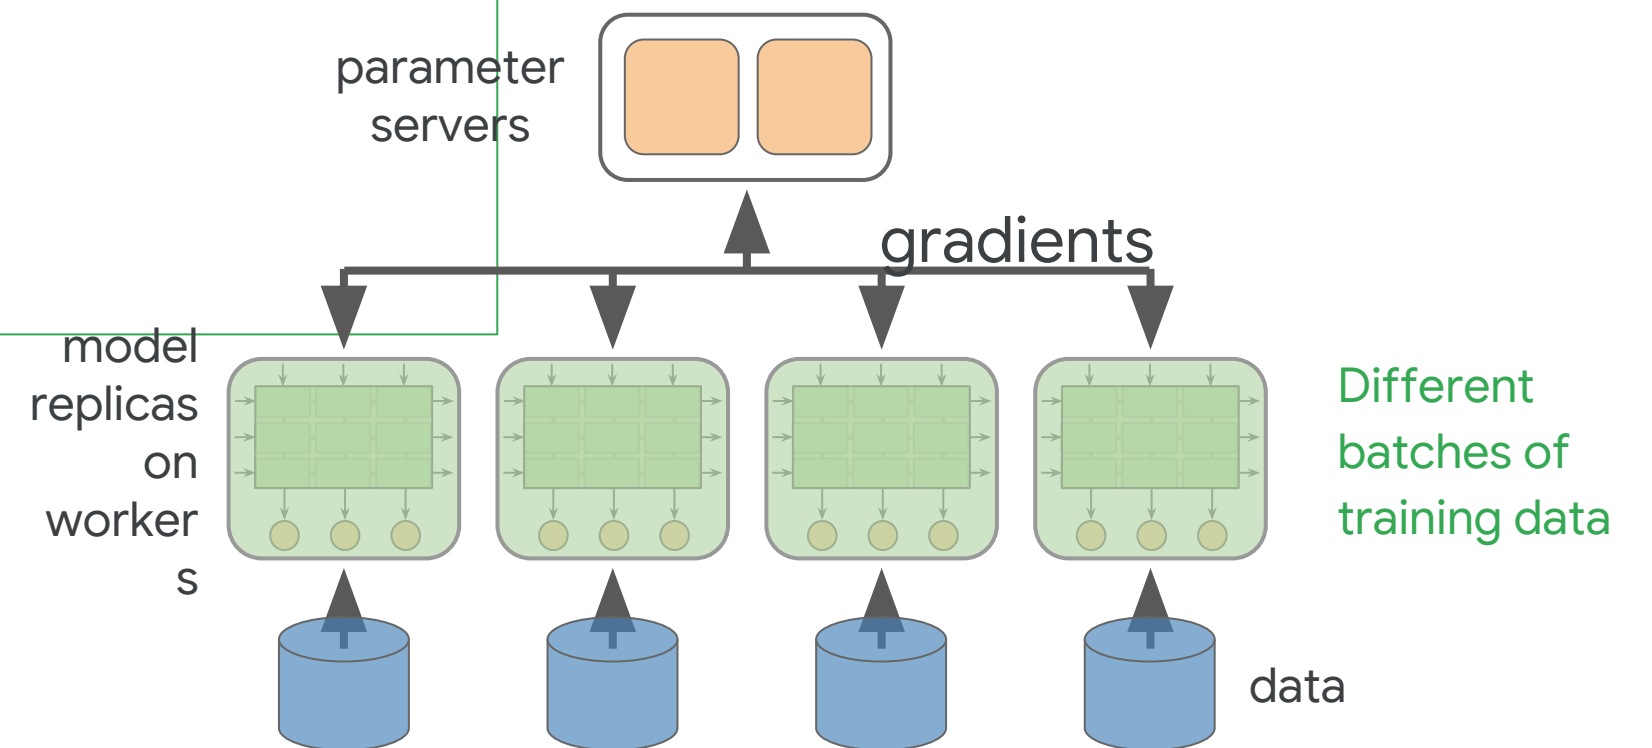
```
run_config =  
tf.estimator.RunConfig(model_dir=output_dir, ...)  
  
estimator =  
tf.estimator.LinearRegressor(featcols, config=run_config)  
  
train_spec =  
tf.estimator.TrainSpec(input_fn=train_input_fn, max_steps=1000)  
  
export_latest =  
tf.estimator.LatestExporter(serving_input_receiver_fn=serving_input_fn)  
  
eval_spec =  
tf.estimator.EvalSpec(input_fn=eval_input_fn, exporters=export_latest)  
  
tf.estimator.train_and_evaluate(estimator, train_spec, eval_spec)
```

Shuffling is even more important in distributed training

```
dataset =  
tf.data.Dataset.list_files("train.csv-*") \  
    .shuffle(100) \  
    .flat_map(tf.data.TextLineDataset) \  
    .map(decode_csv)  
  
dataset = dataset.shuffle(1000) \  
    .repeat(15) \  
    .batch(128)
```

Shuffling is even more important in distributed training.

```
dataset =  
tf.data.Dataset.list_files("train.csv-*") \  
    .shuffle(100) \  
    .flat_map(tf.data.TextLineDataset) \  
    .map(decode_csv)  
  
dataset = dataset.shuffle(1000) \  
    .repeat(15) \  
    .batch(128)
```





Monitoring with TensorBoard


Martin Gerner

Real World ML Models

Problem	Solution
✓ Out of memory data	Use the Dataset API
✓ Distribution	Use <code>train_and_evaluate</code>
Need to evaluate during training	Use <code>train_and_evaluate</code> + TensorBoard
Deployments that scale	?

☐ Show data download links☒ Ignore outliers in chart scalingTooltip sorting method: default

Smoothing

 0.463

Horizontal Axis

STEP

RELATIVE

WALL

Runs

Write a regex to filter runs

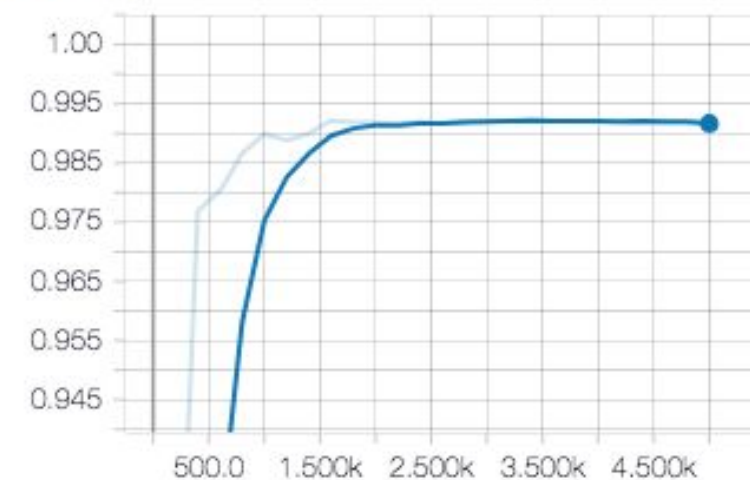
☒ ☐ .
☒ ☐ eval

Q acc*|lear*|loss

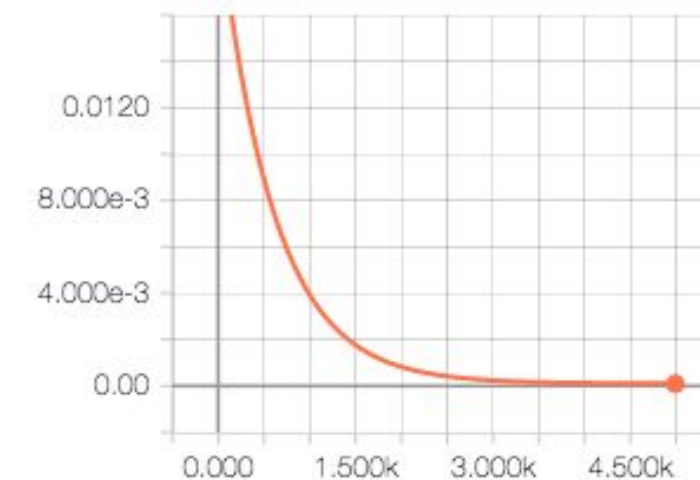
Tags matching /acc*|lear*|loss/

3

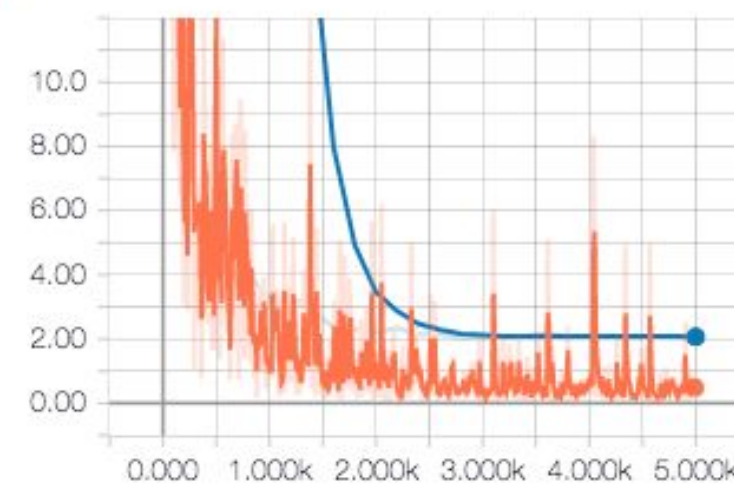
accuracy

learn_rate

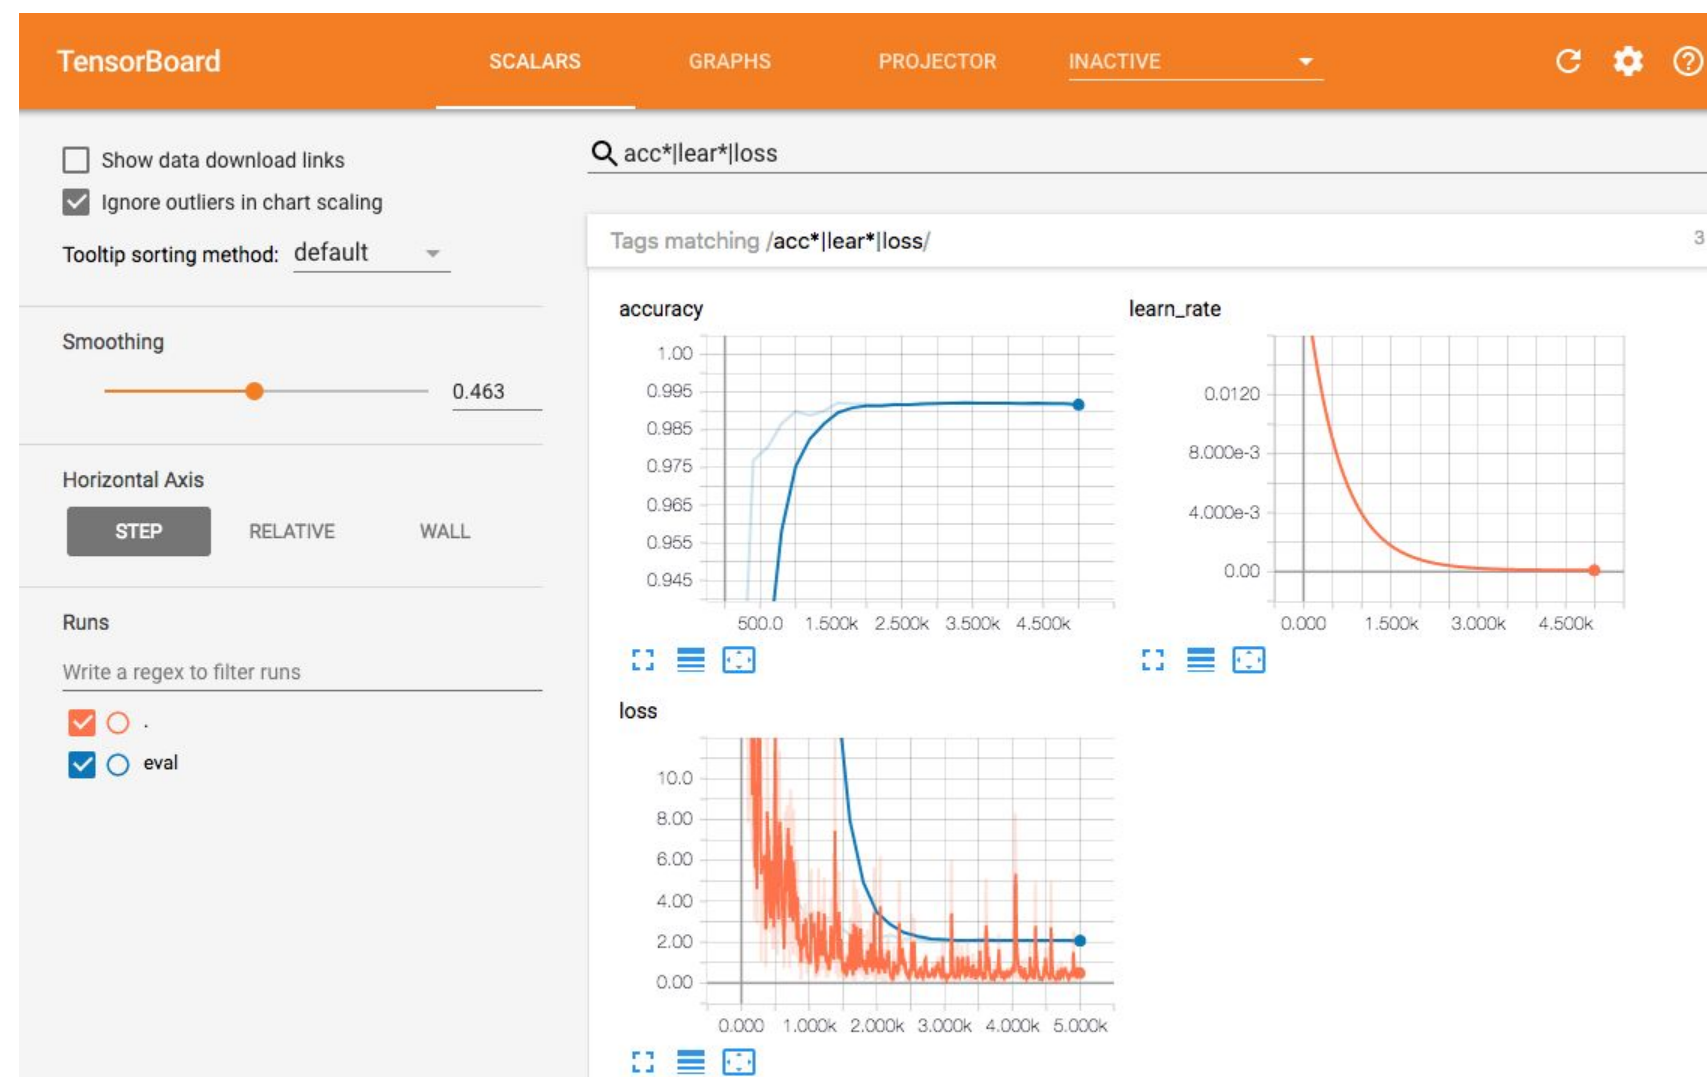
  

loss

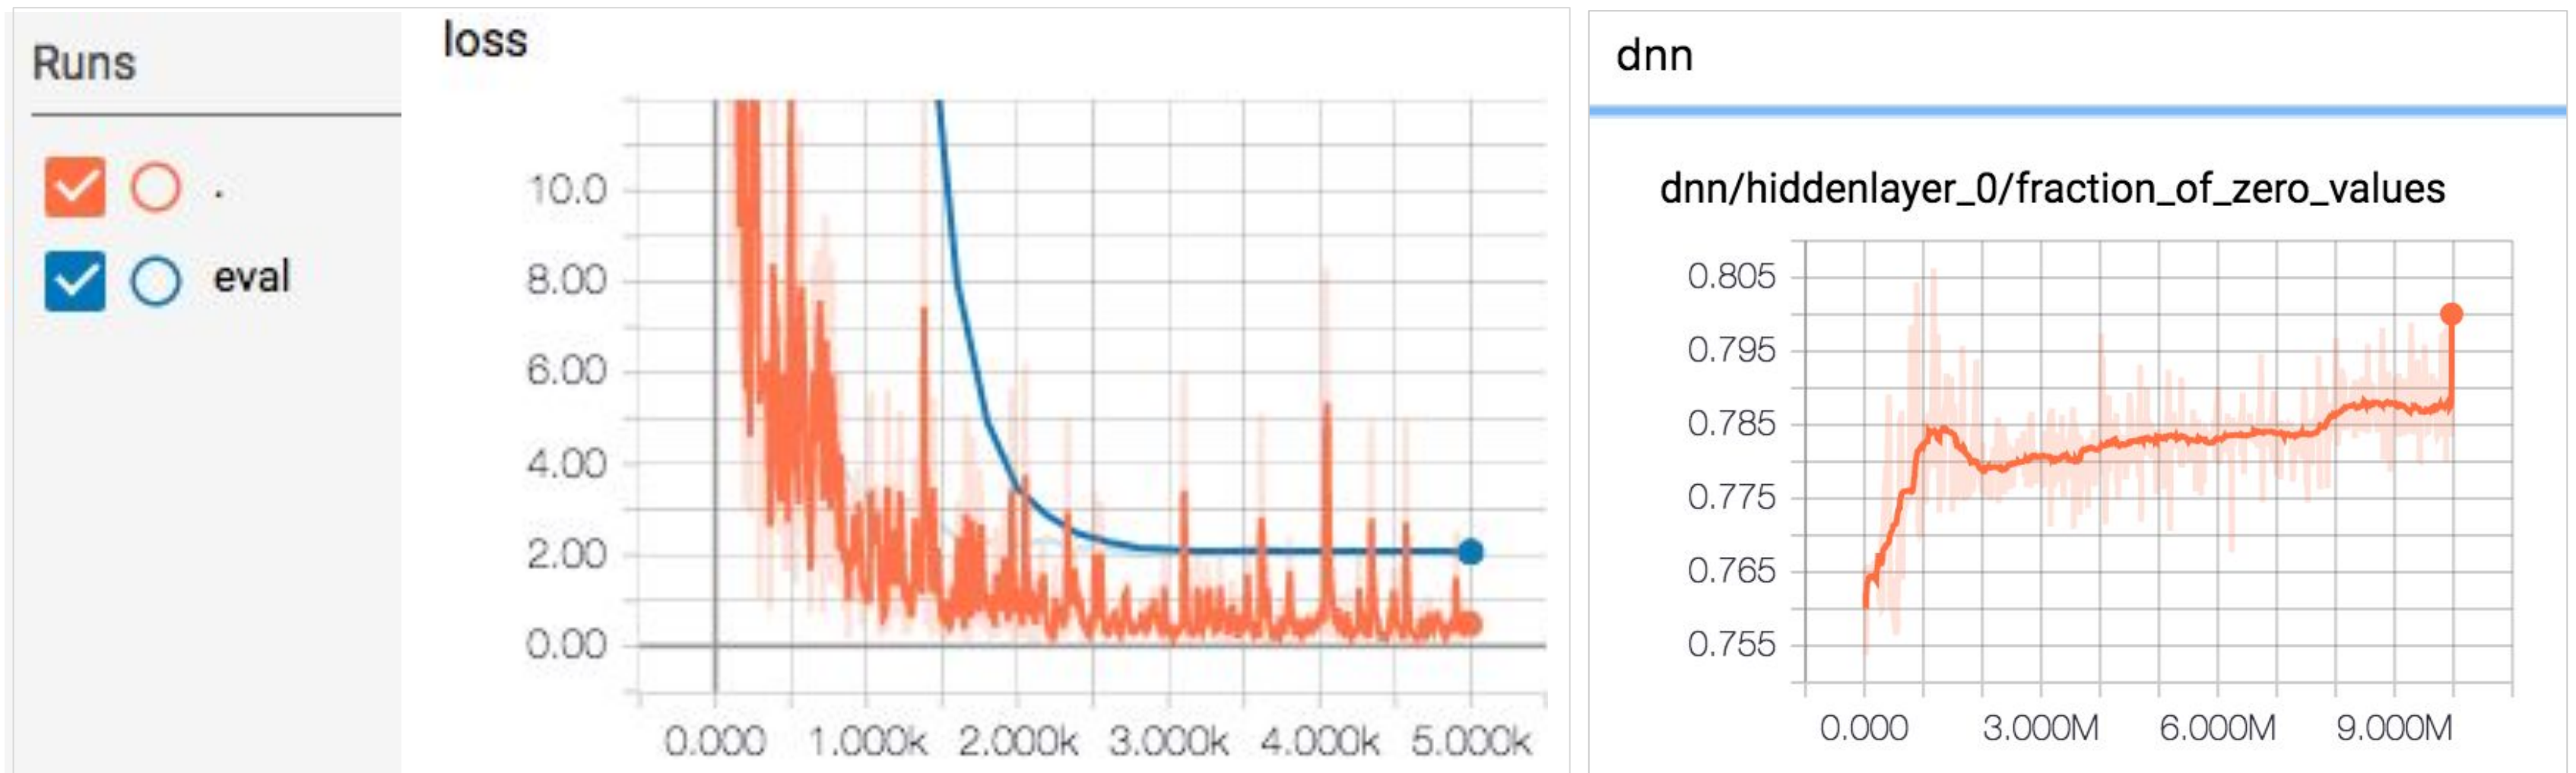
  

Point Tensorboard to your output directory and the dashboards appear in your browser at localhost:6006

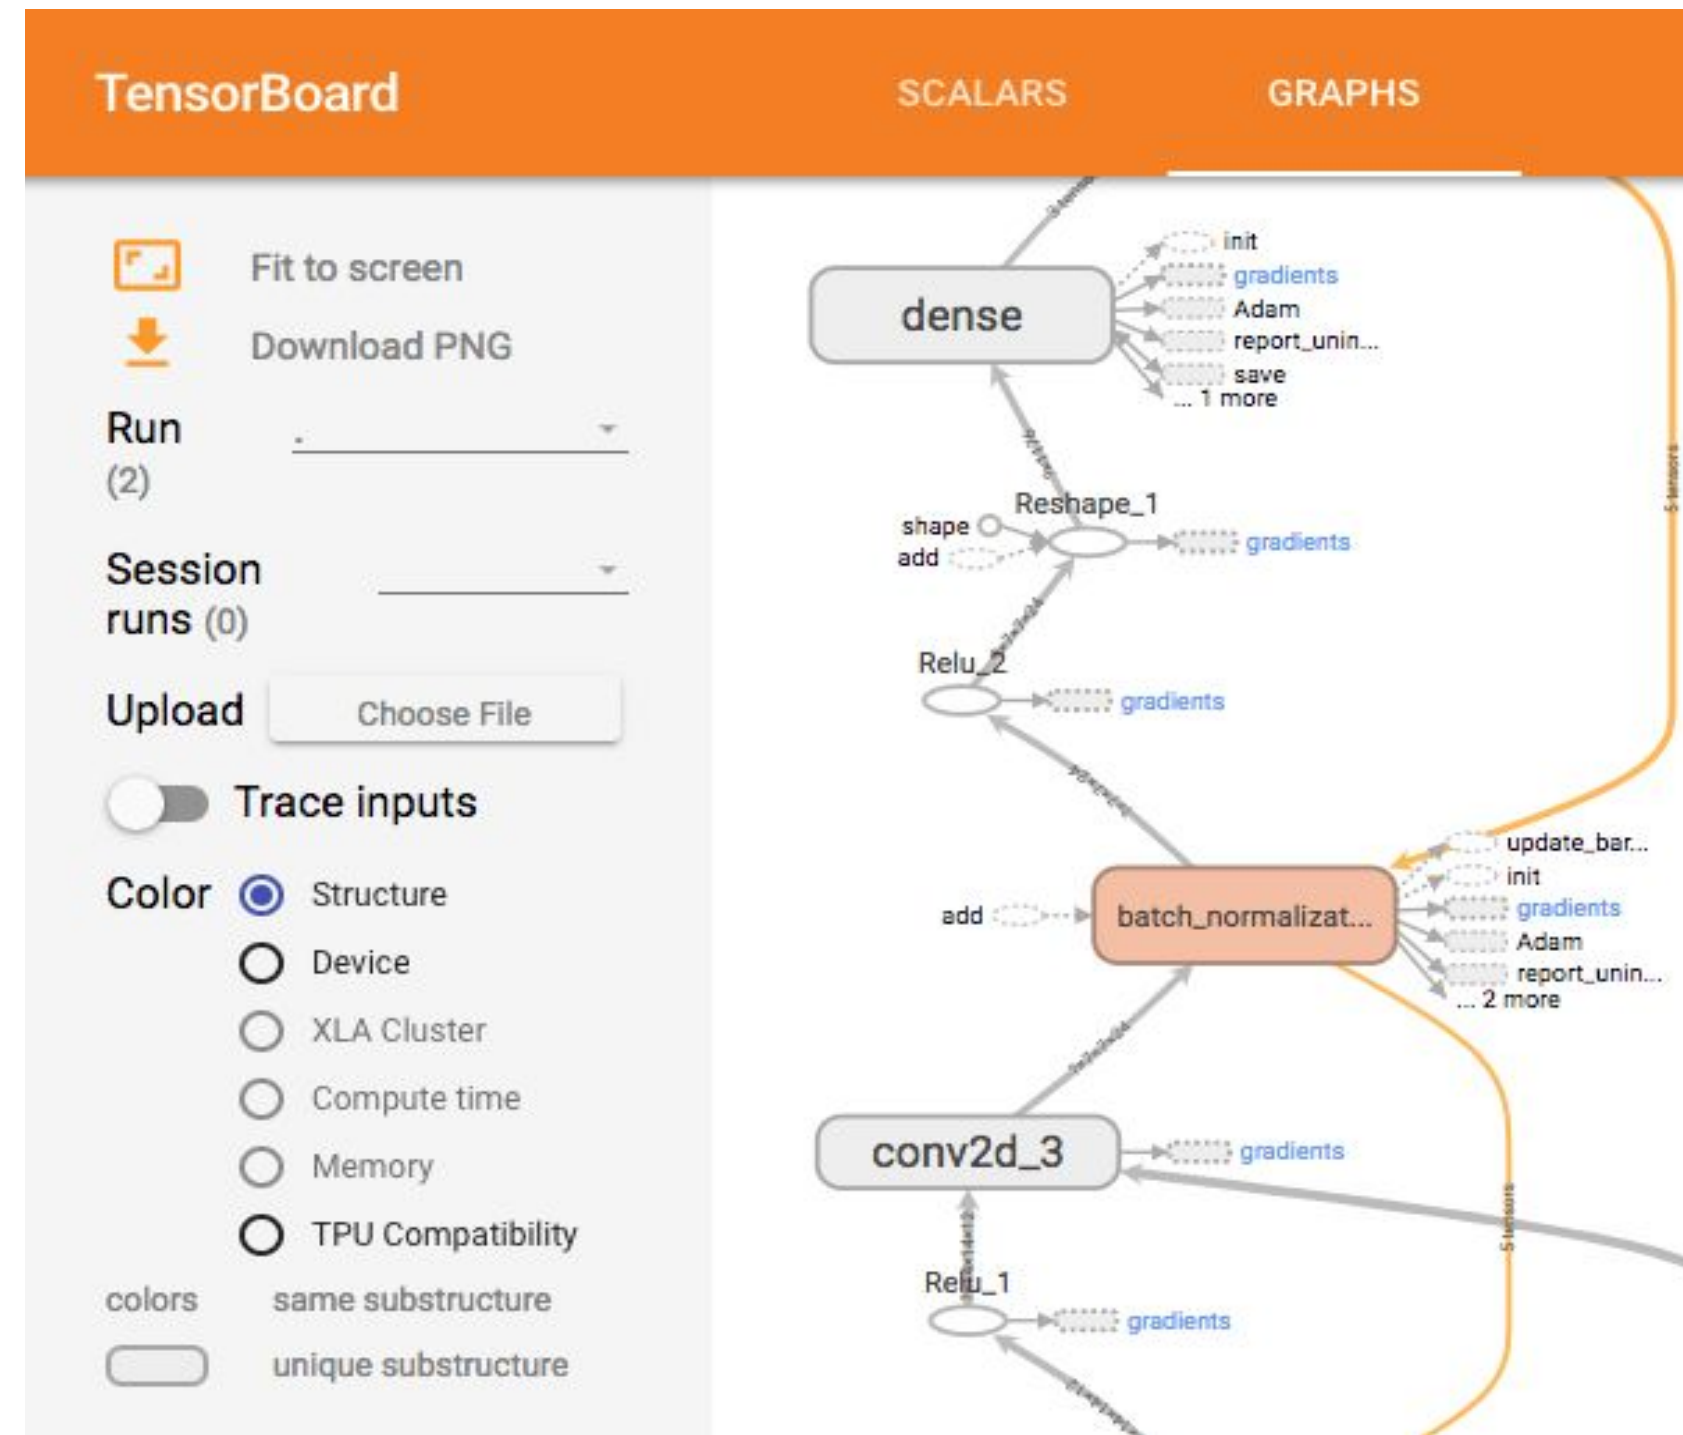
```
tf.estimator.RunConfig(model_dir='output_dir')  
> tensorboard --logdir output_dir
```



Pre-made Estimators export relevant metrics, embeddings, histograms, etc. for TensorBoard, so there is nothing more to do



The dashboard for the graph



If you are writing a custom Estimator model, you can add summaries for Tensorboard with a single line

```
tf.summary.scalar('meanVar1', tf.reduce_mean(var1))  
...  
tf.summary.text('outClass', stringvar))
```

these names will show
up in TensorBoard

Sprinkle appropriate summary ops throughout your code:

`tf.summary.scalar`

`tf.summary.image`

`tf.summary.audio`

`tf.summary.text`

`tf.summary.histogram`

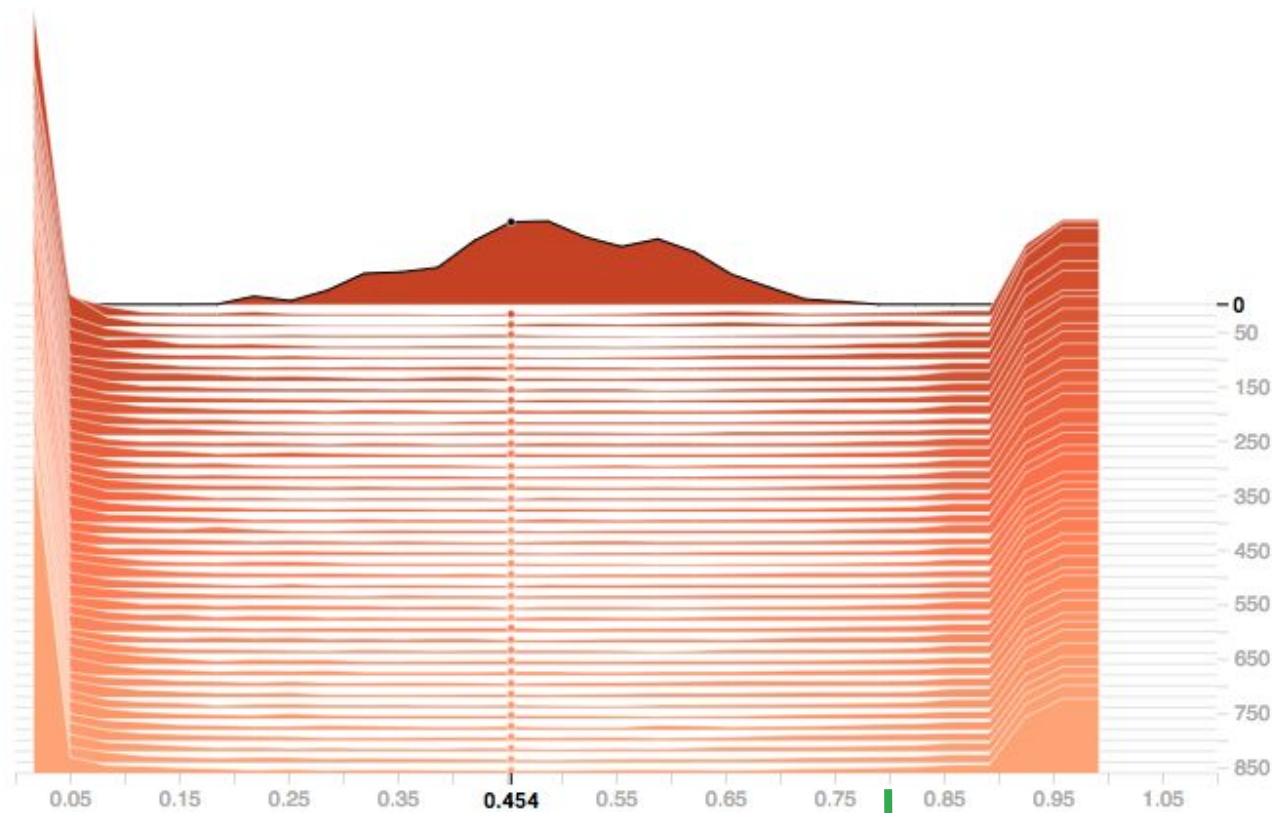
https://www.tensorflow.org/get_started/summaries_and_tensorboard

The dashboard for histograms

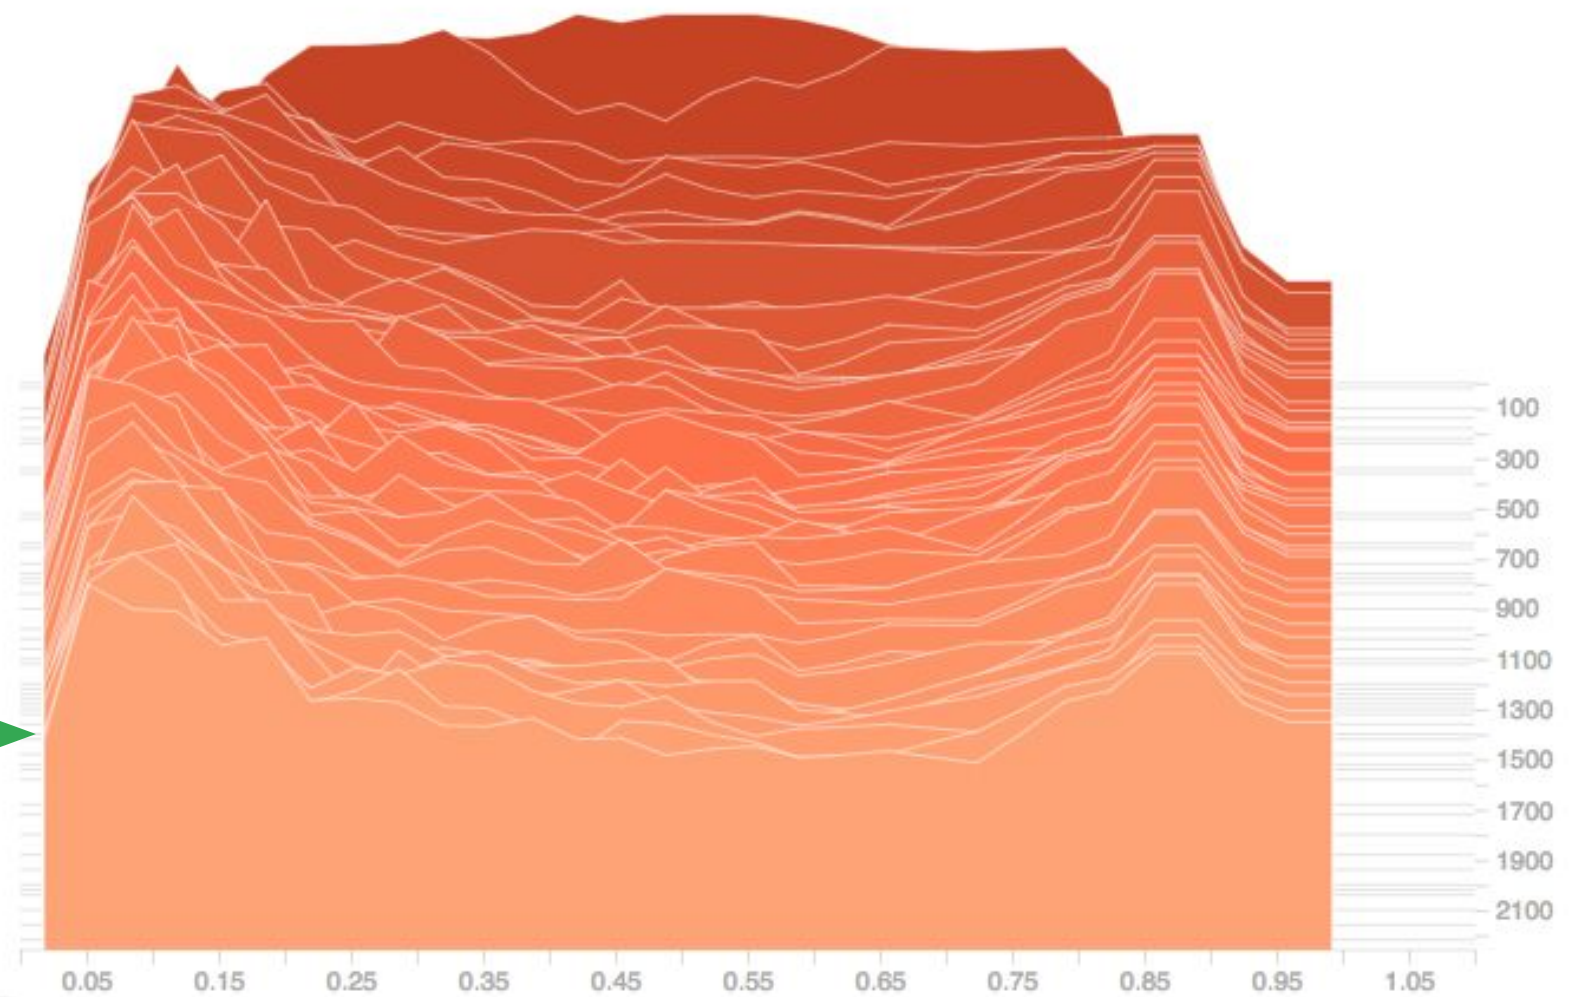


after-activation

756

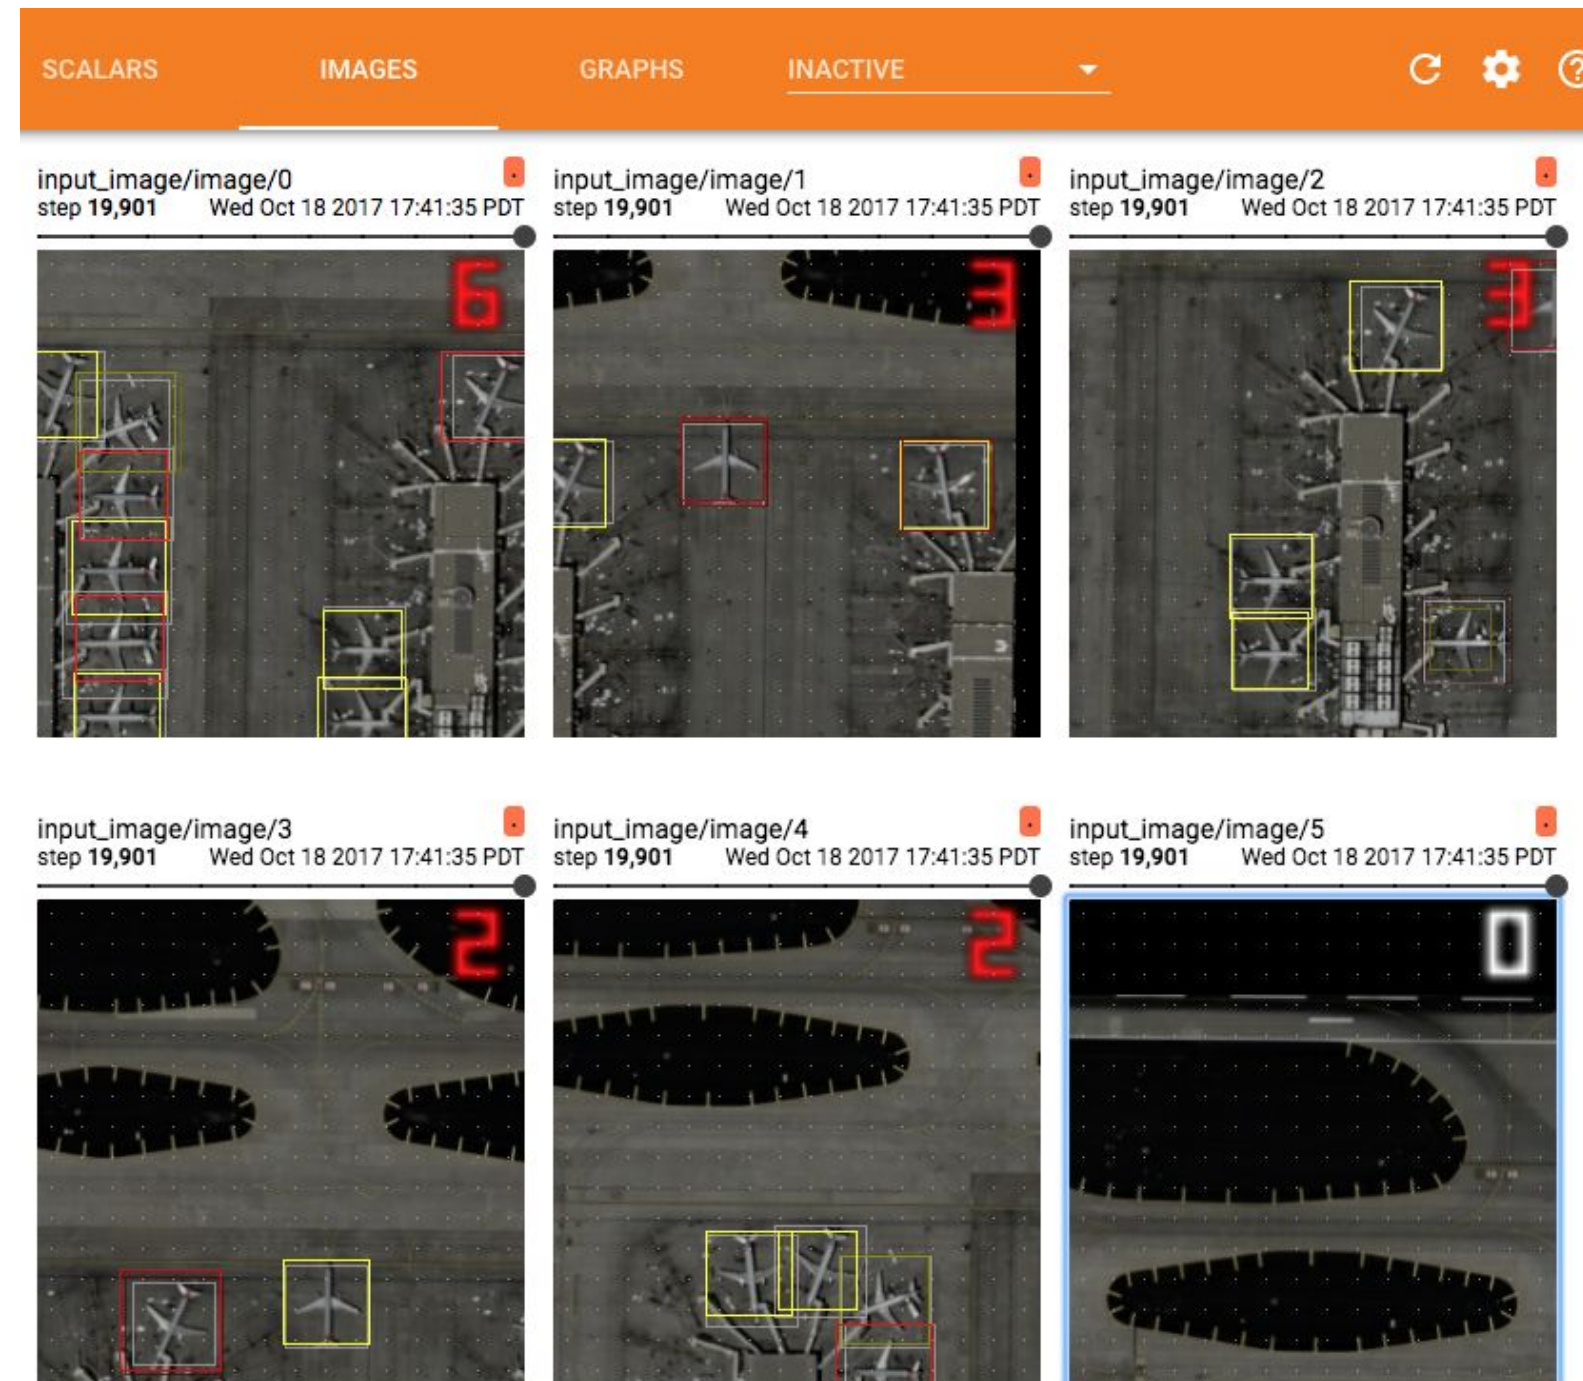


after-activation



After Batch
Normalization

Images and Audio have their own (non-scalar) summary ops

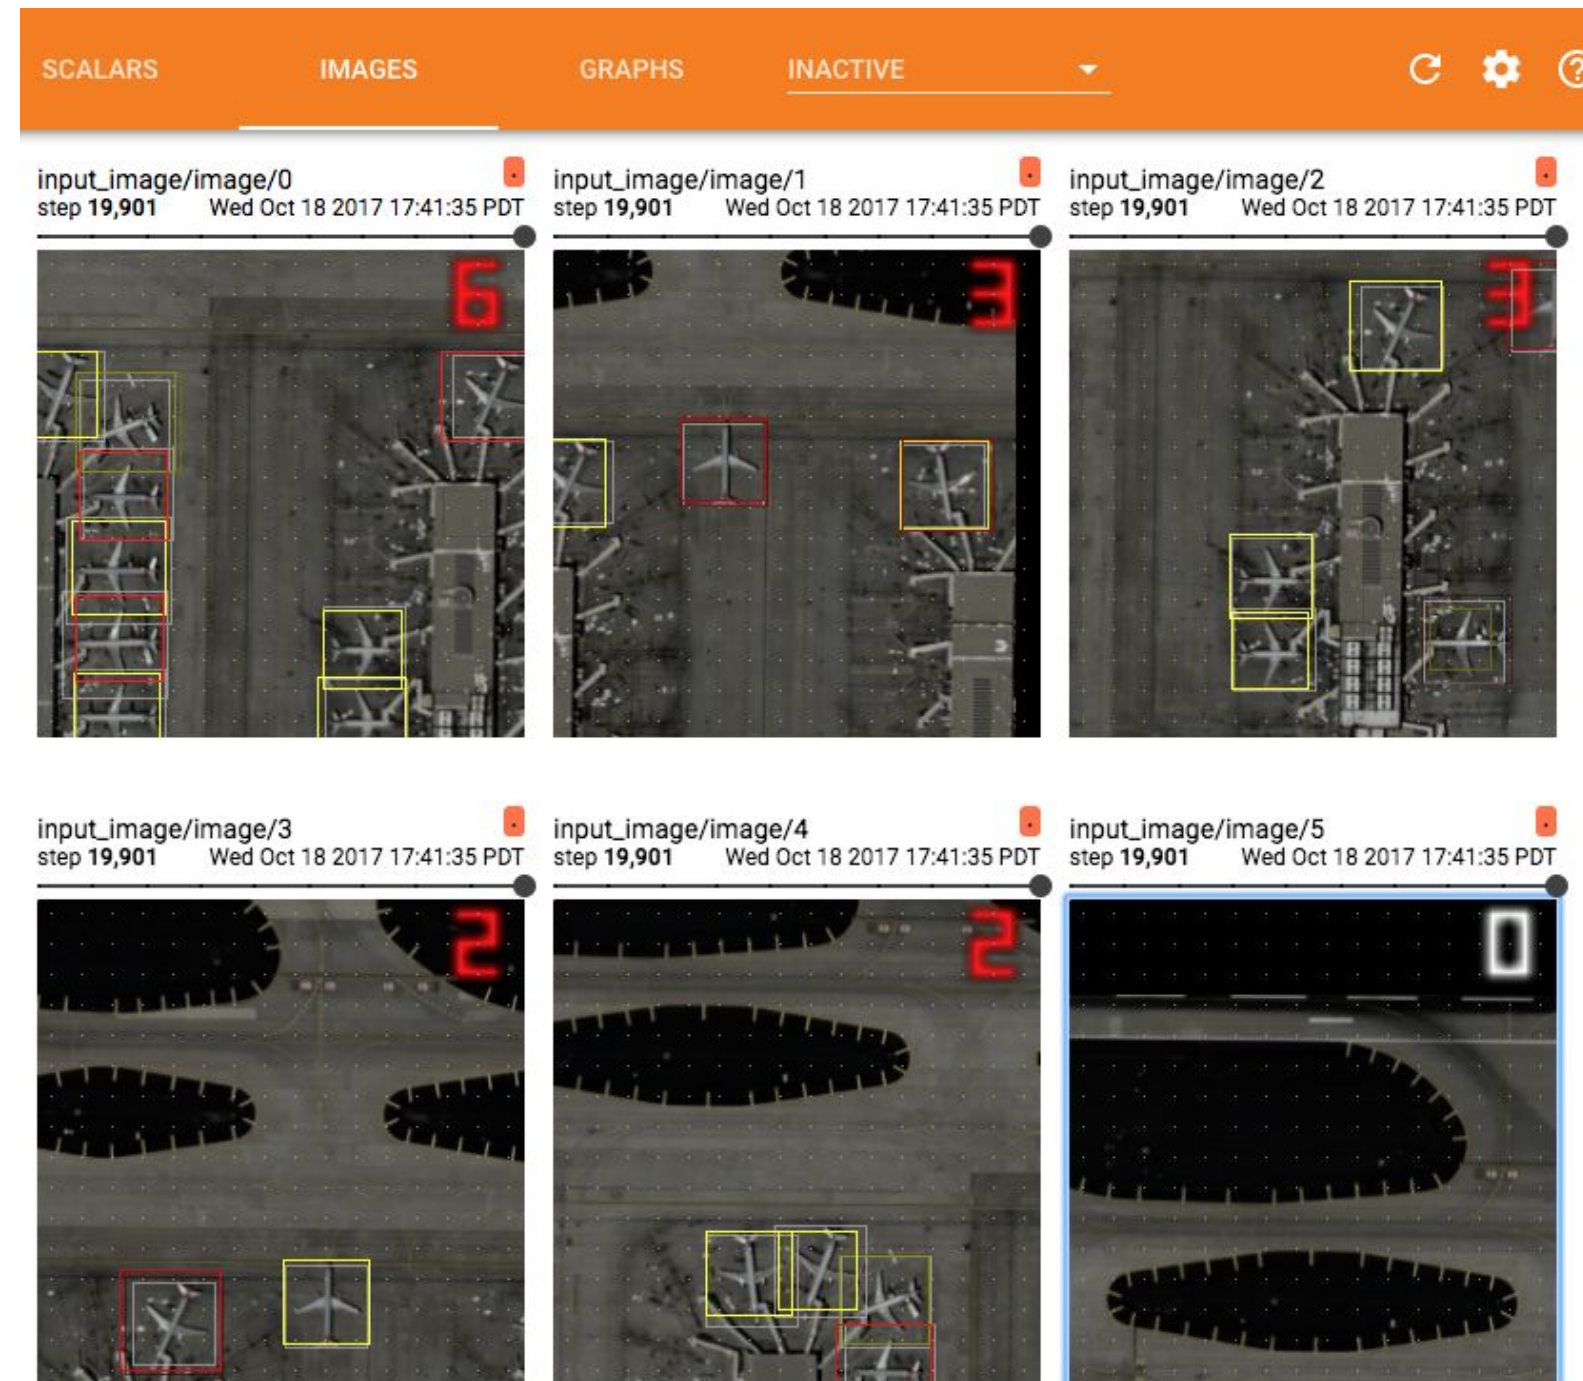


Images and Audio have their own (non-scalar) summary ops

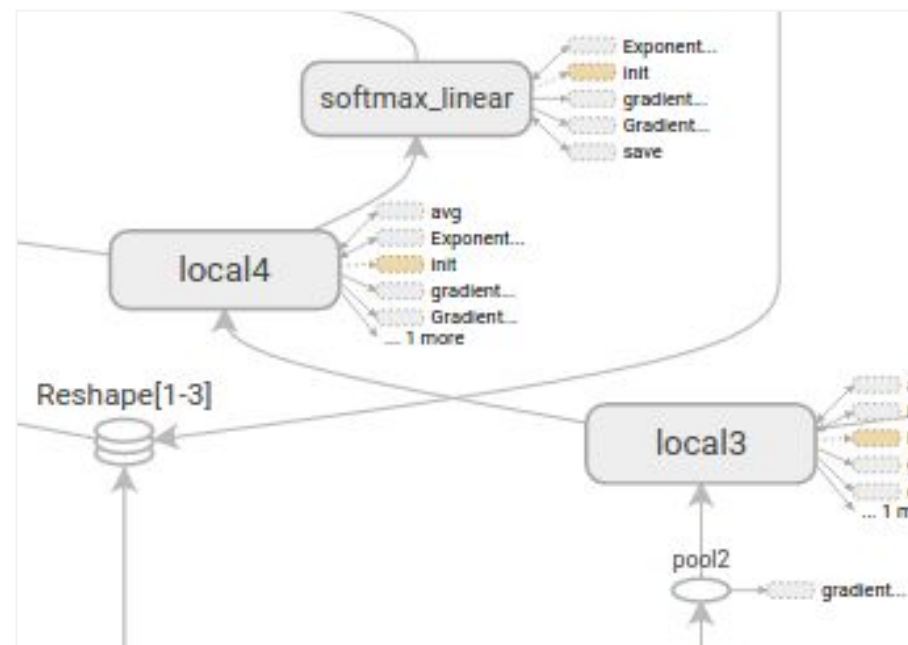
```
p = tf.placeholder("uint8", (None, ht, wd, num_channels))  
s = tf.summary.image("im1", p)
```

```
p = tf.placeholder("float32",  
                  (None, duration_frames, num_channels))  
s = tf.summary.audio(name='au1', tensor=p, sample_rate=4000)
```

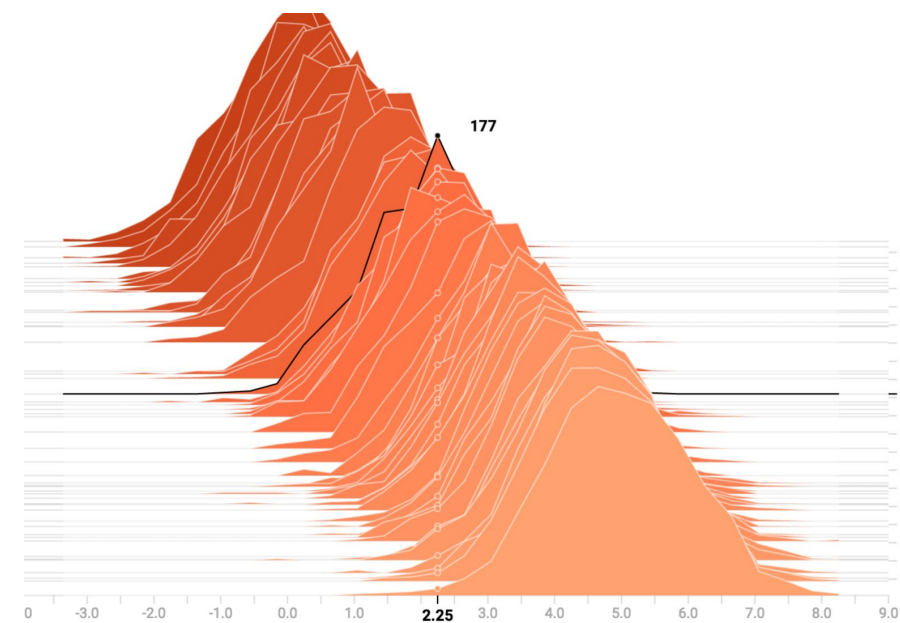
Images and Audio have their own (non-scalar) summary ops



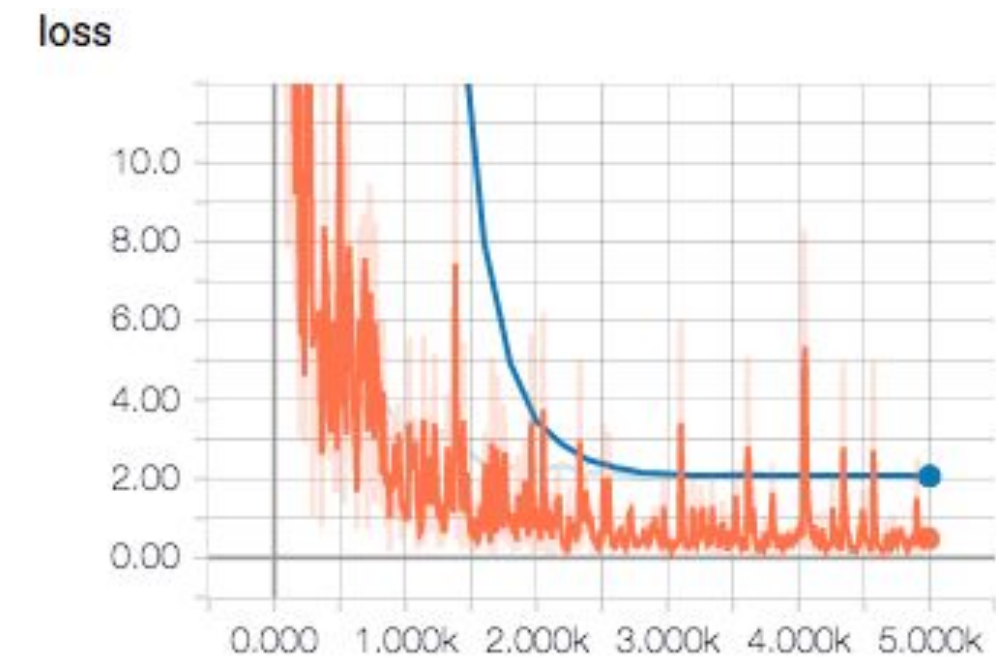
TensorBoard has a suite of visualization tools to explore and explain your model and results



Graph Explorer



Histogram Dashboard



Scalar Dashboard

Many more including Audio,
Image and Text Dashboards ...

Real World ML Models

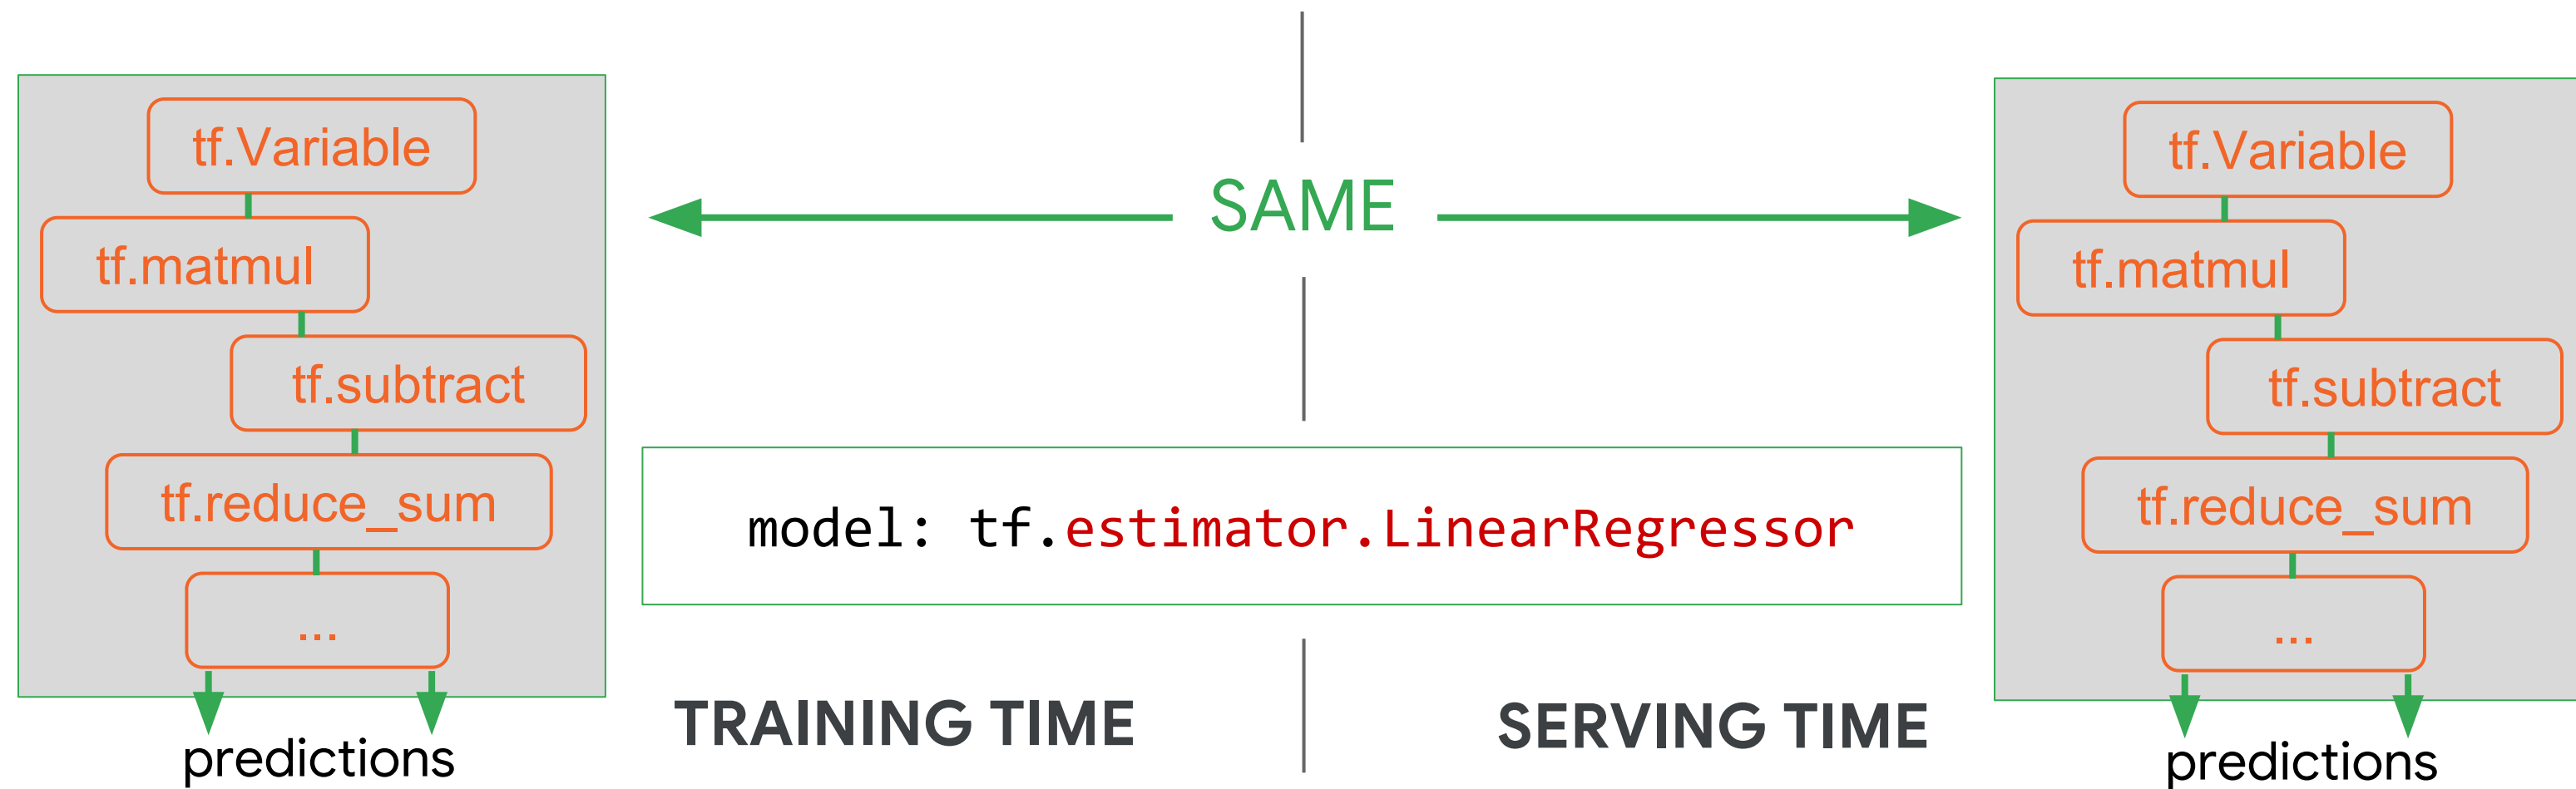
Problem	Solution
✓ Out of memory data	Use the Dataset API
✓ Distribution	Use <code>train_and_evaluate</code>
✓ Need to evaluate during training	Use <code>train_and_evaluate</code> + TensorBoard
Deployments that scale	Use serving input function

Recap with all the code

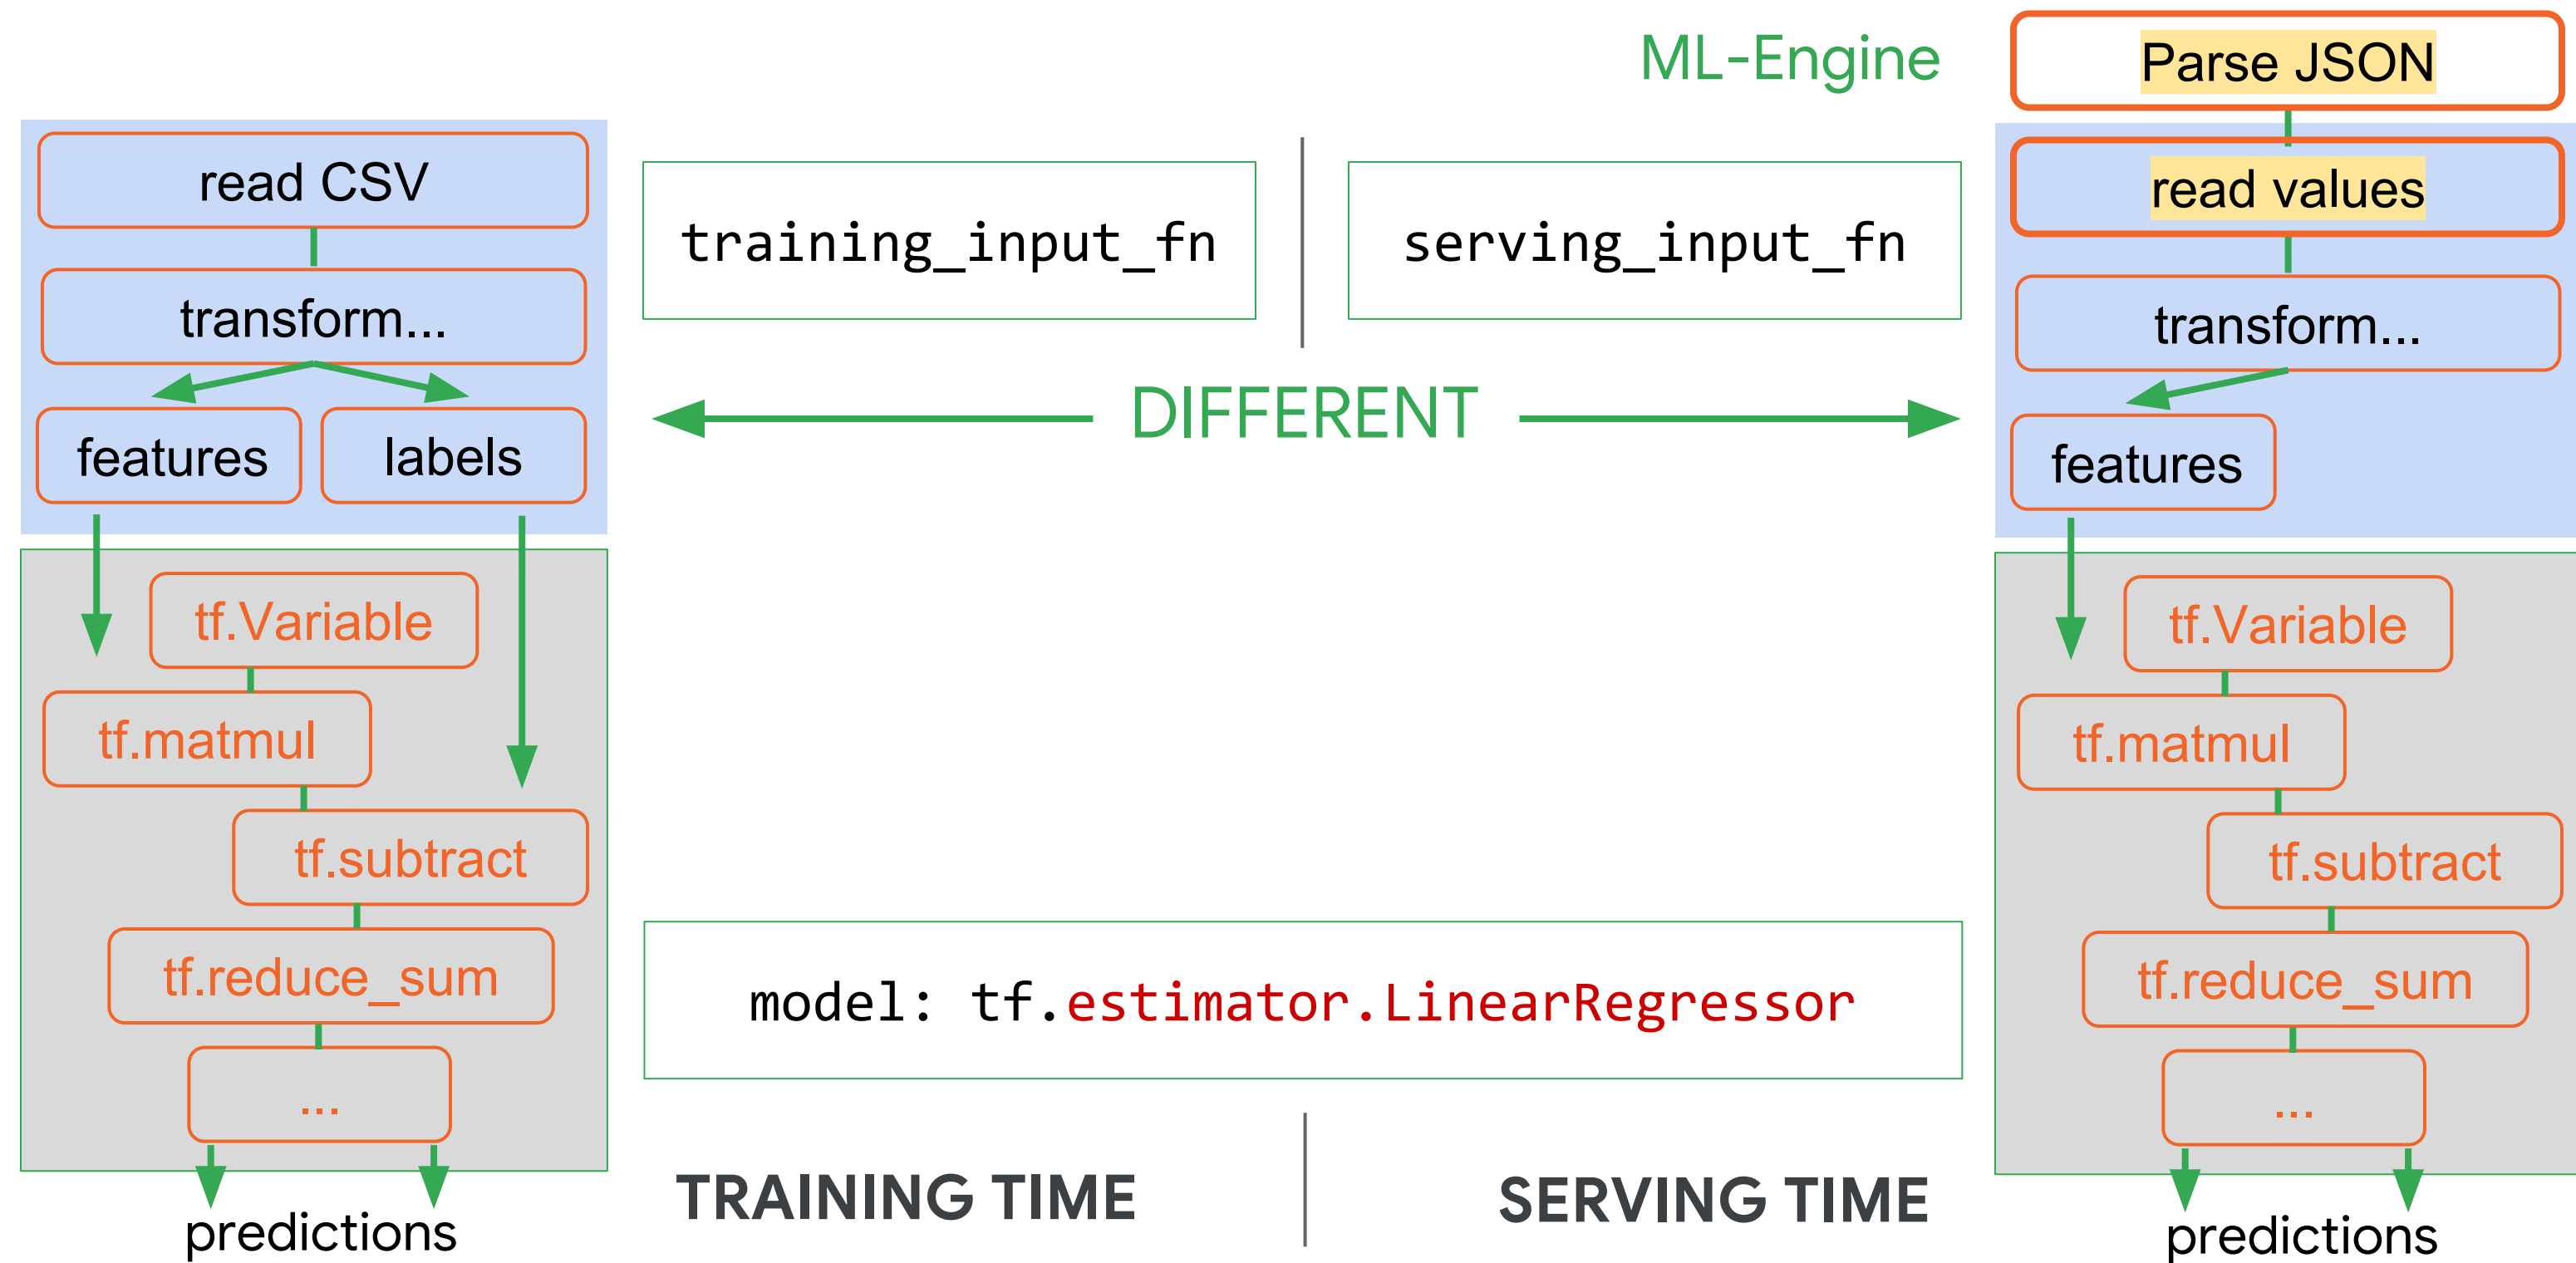
```
run_config =  
tf.estimator.RunConfig(model_dir=output_dir, ...)  
  
estimator =  
tf.estimator.LinearRegressor(featcols, config=run_config)  
  
train_spec =  
tf.estimator.TrainSpec(input_fn=train_input_fn, max_steps=1000)  
  
export_latest =  
tf.estimator.LatestExporter(serving_input_receiver_fn=serving_input_fn)  
  
eval_spec =  
tf.estimator.EvalSpec(input_fn=eval_input_fn, exporters=export_latest)  
  
tf.estimator.train_and_evaluate(estimator, train_spec, eval_spec)
```


Serving and training-time
inputs are often very different

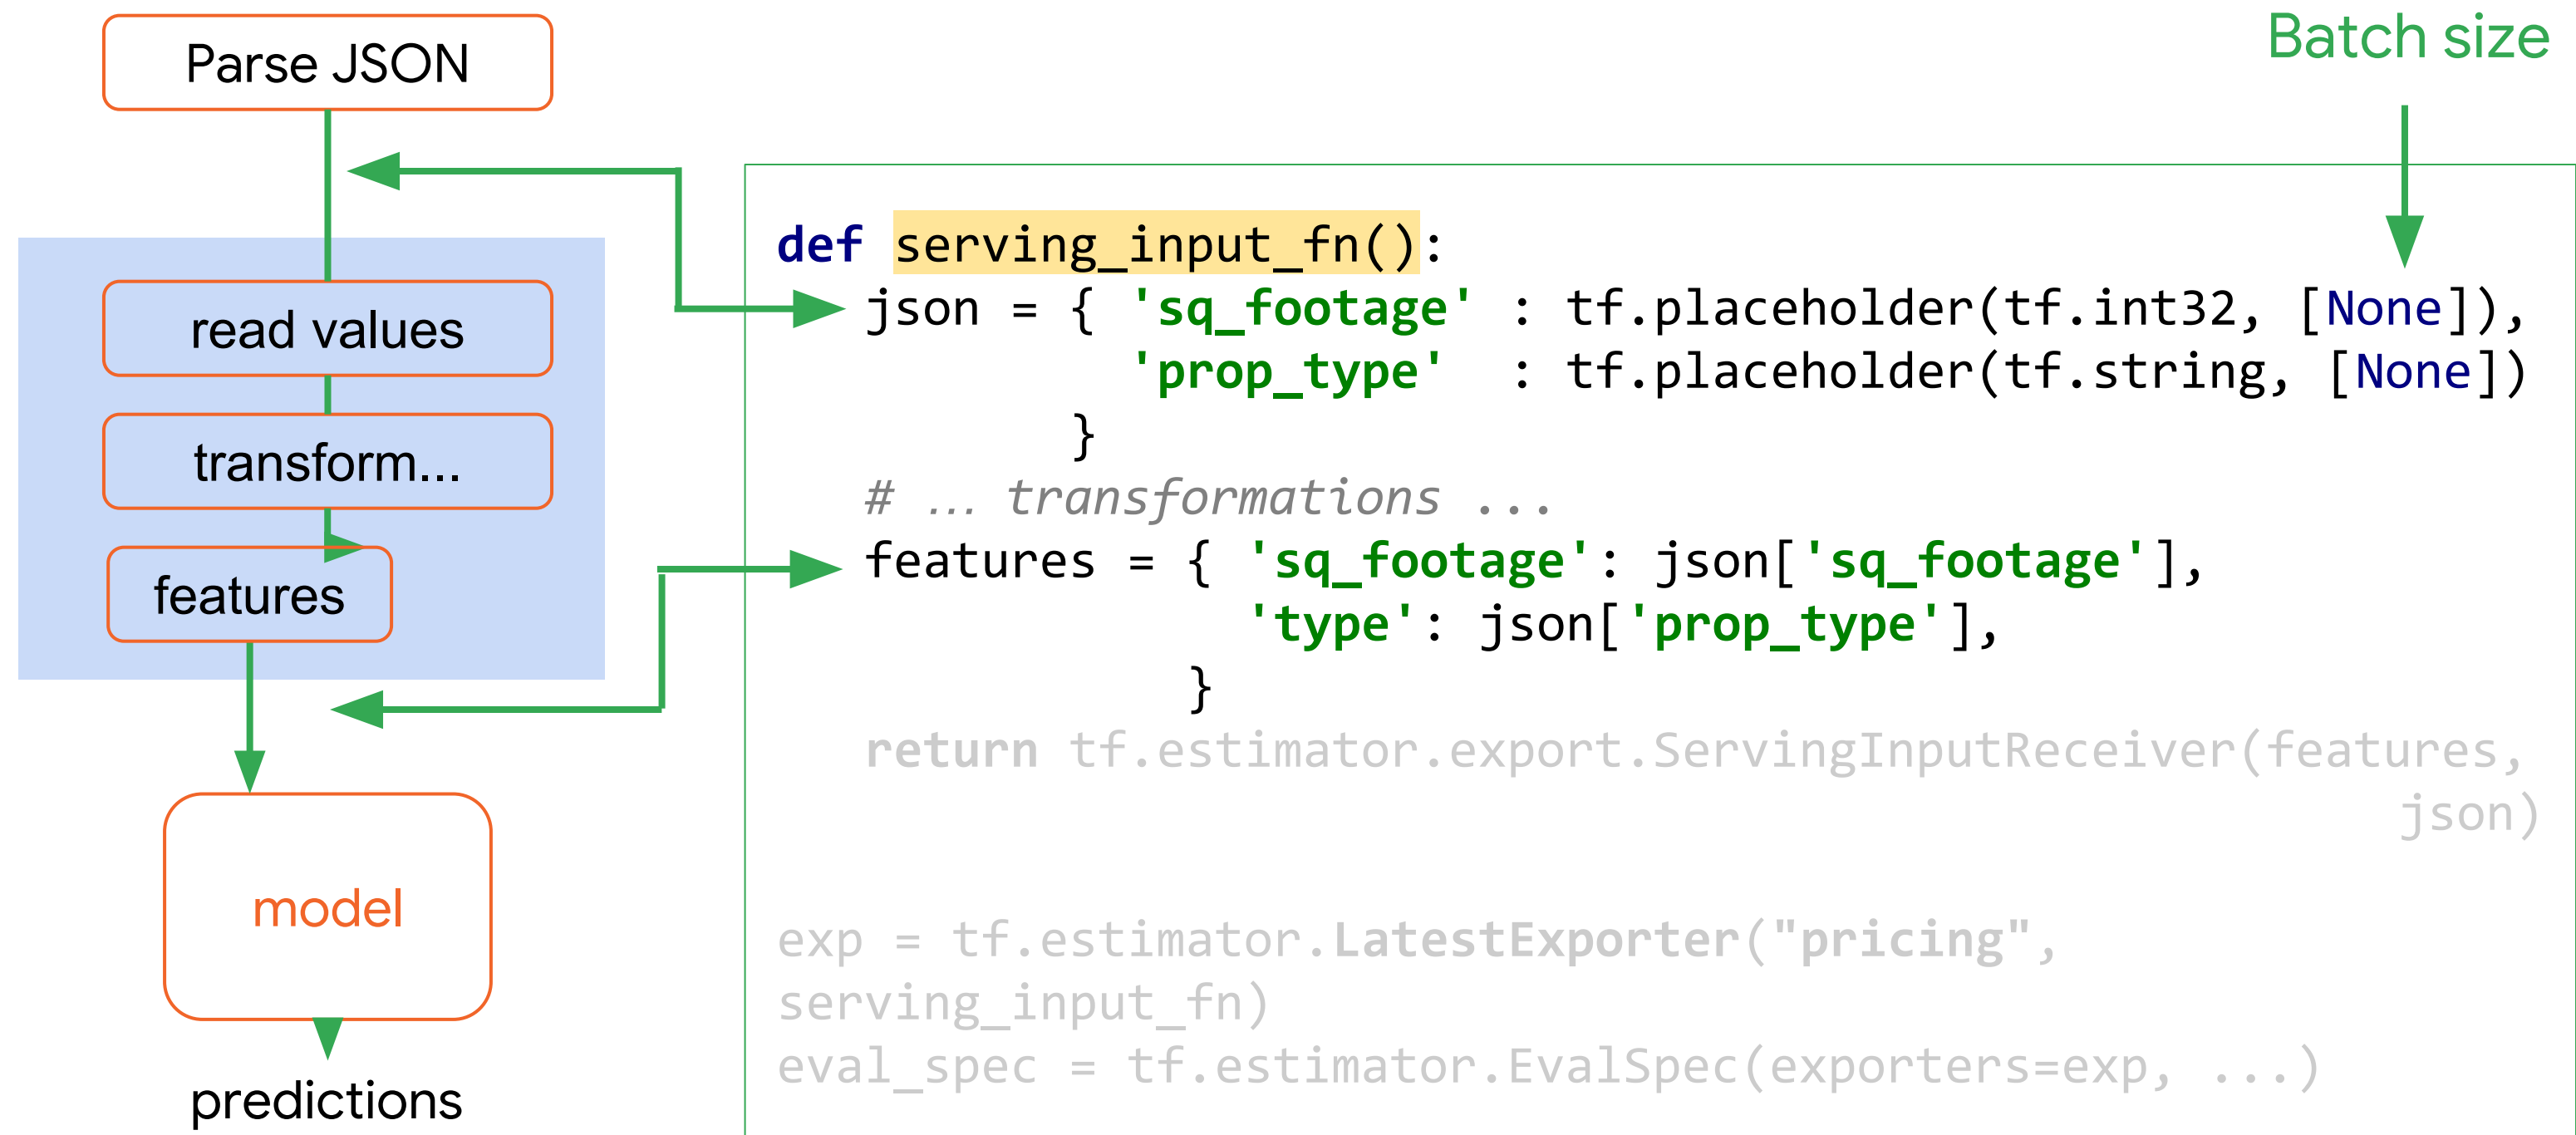
Serving and training-time inputs are often very different



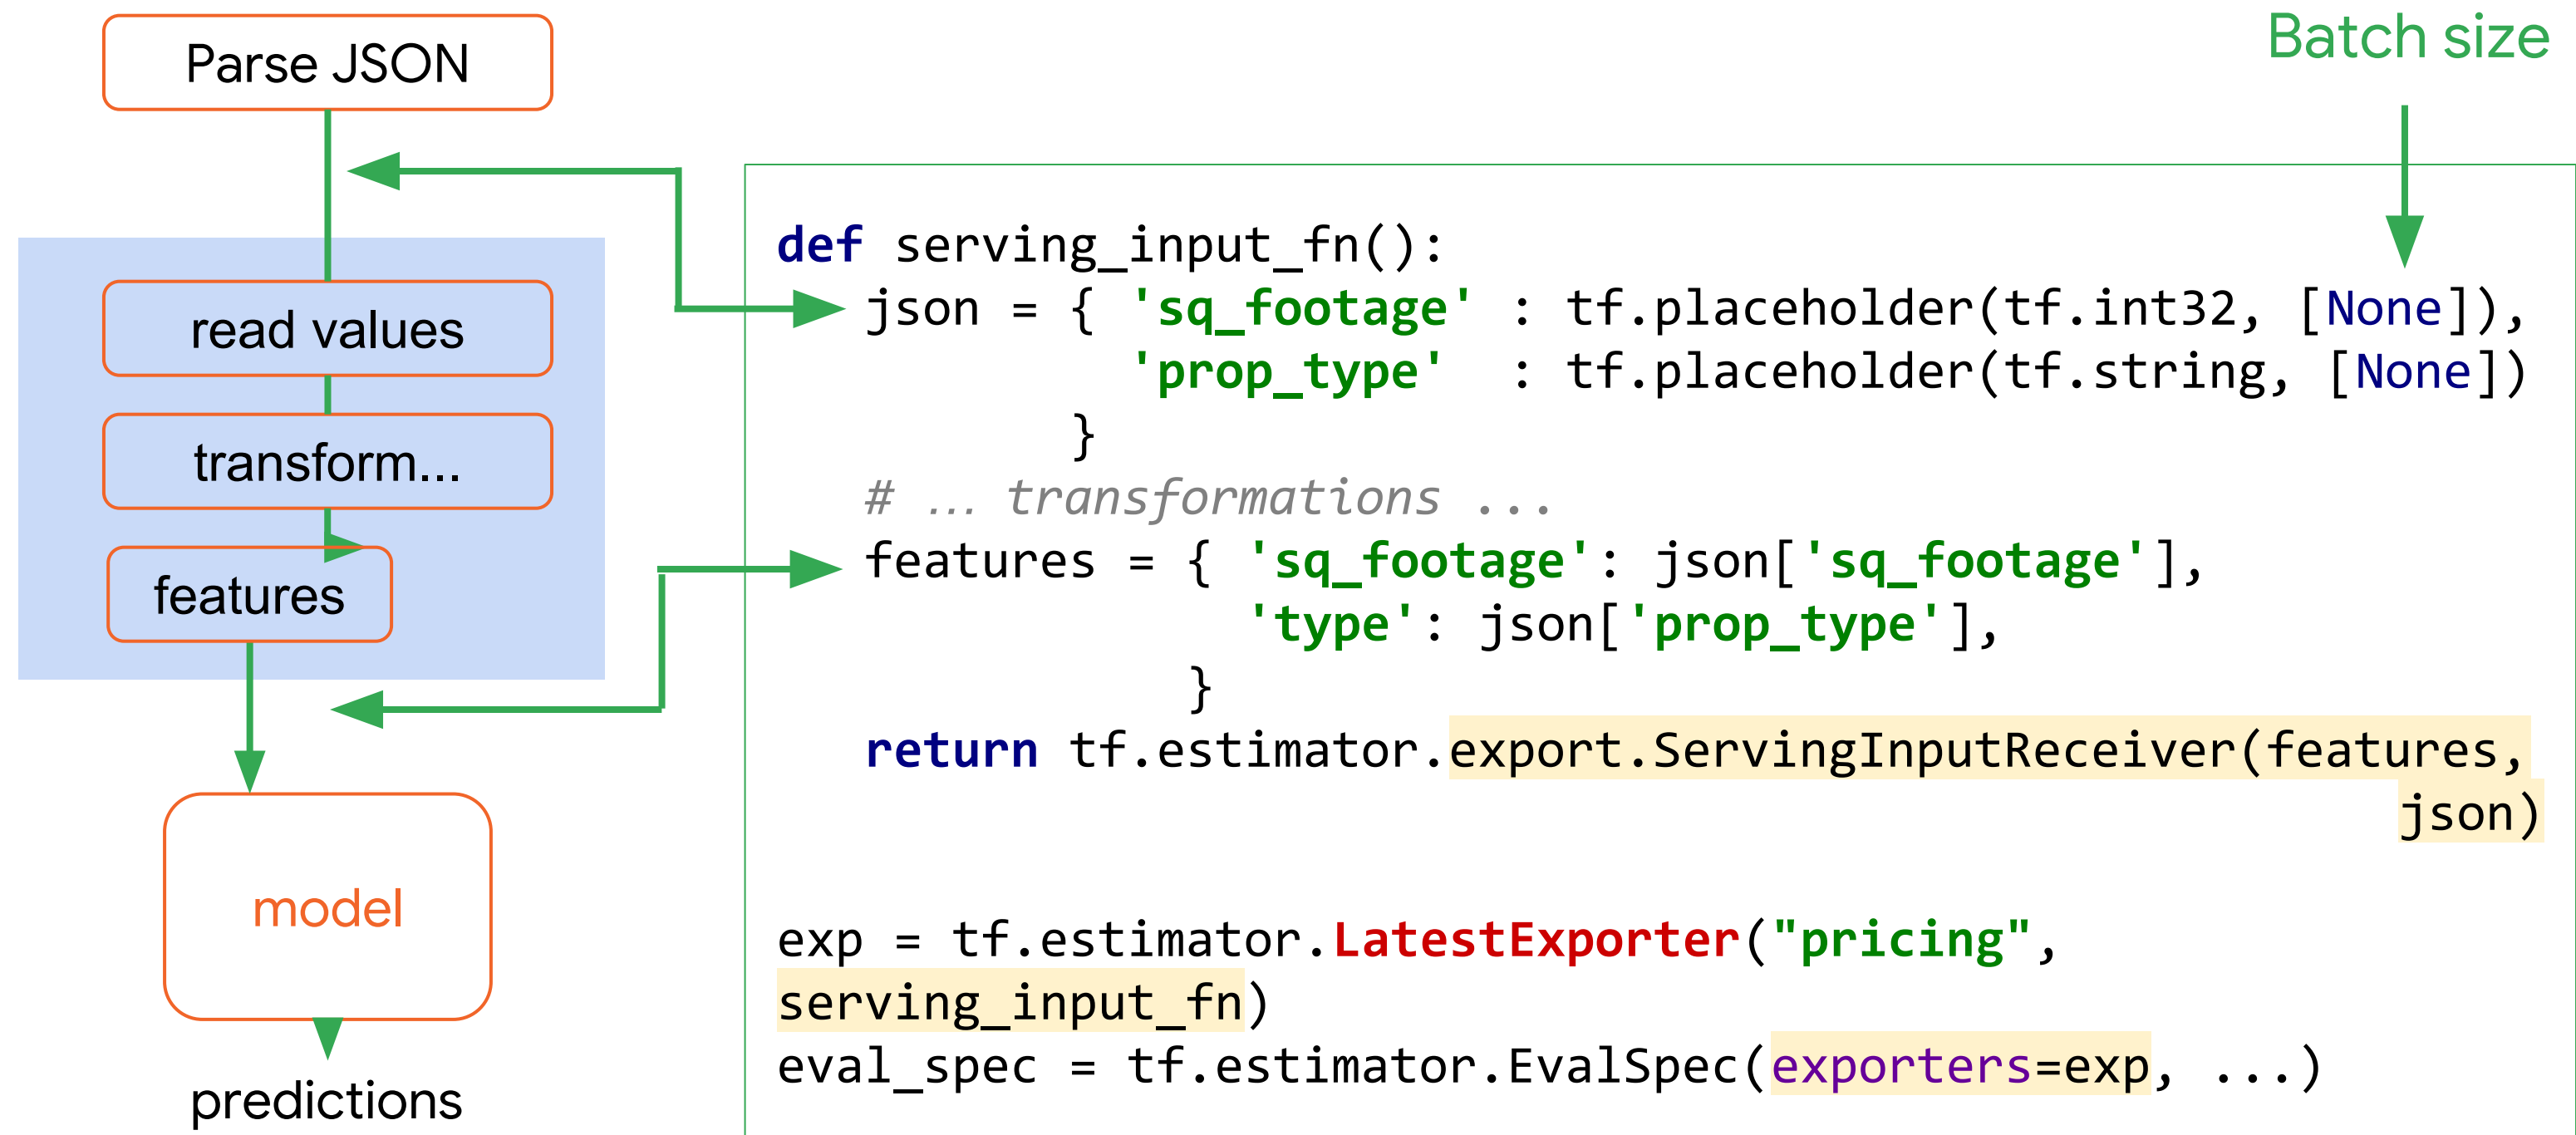
Serving and training-time inputs are often very different



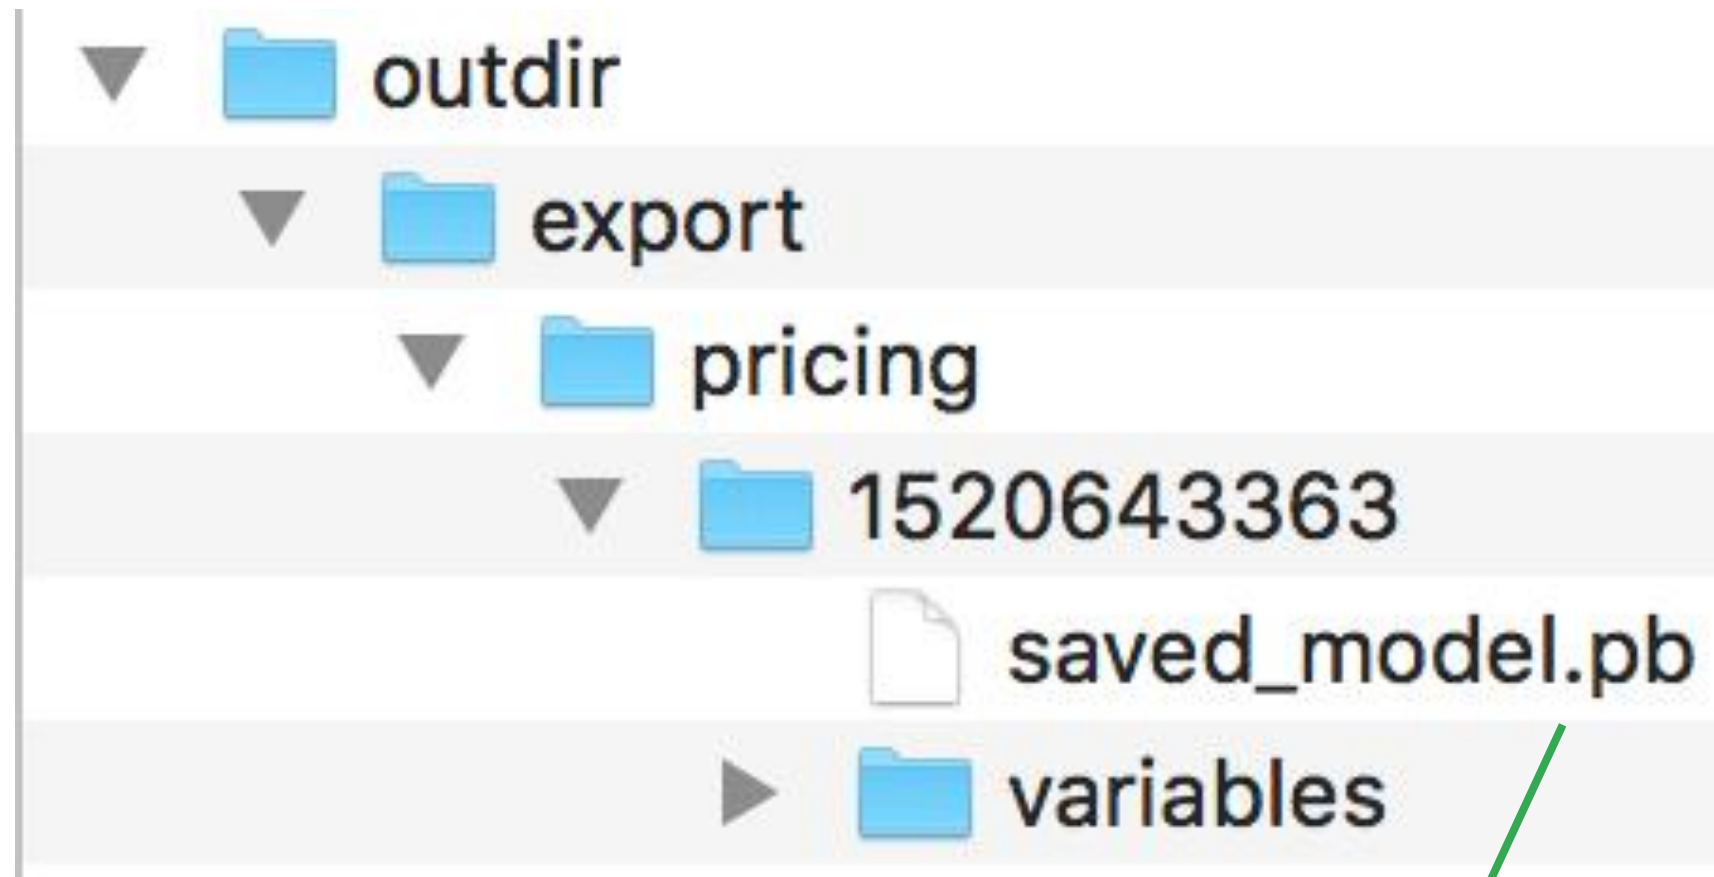
Serving input function transforms from parsed JSON data to the data your model expects



Serving input function transforms from parsed JSON data to the data your model expects

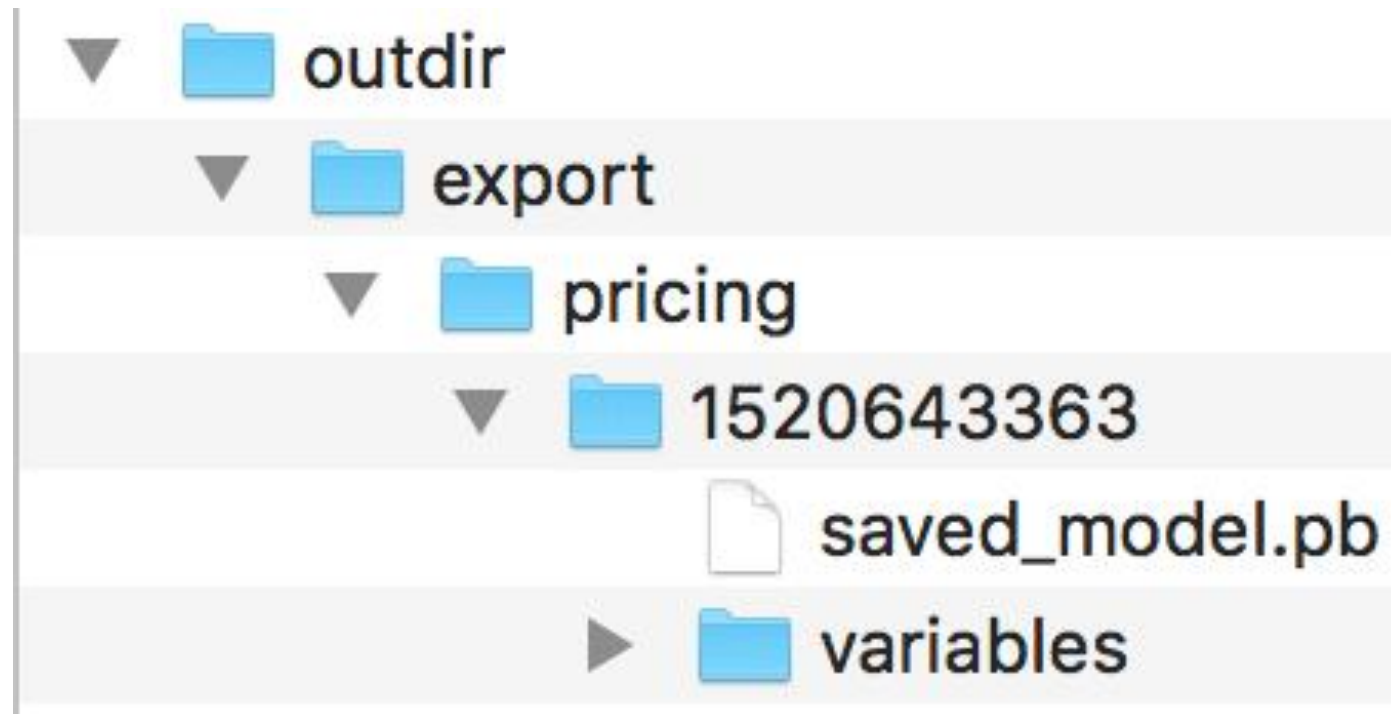


The exported model is
ready to deploy



One click deploy to ML Engine

The exported model is ready to deploy



Online predictions :

```
{"instances": [
  {"sq_footage": 1520, "prop_type": "apt"},
  {"sq_footage": 1520, "prop_type": "house"}
]}
```

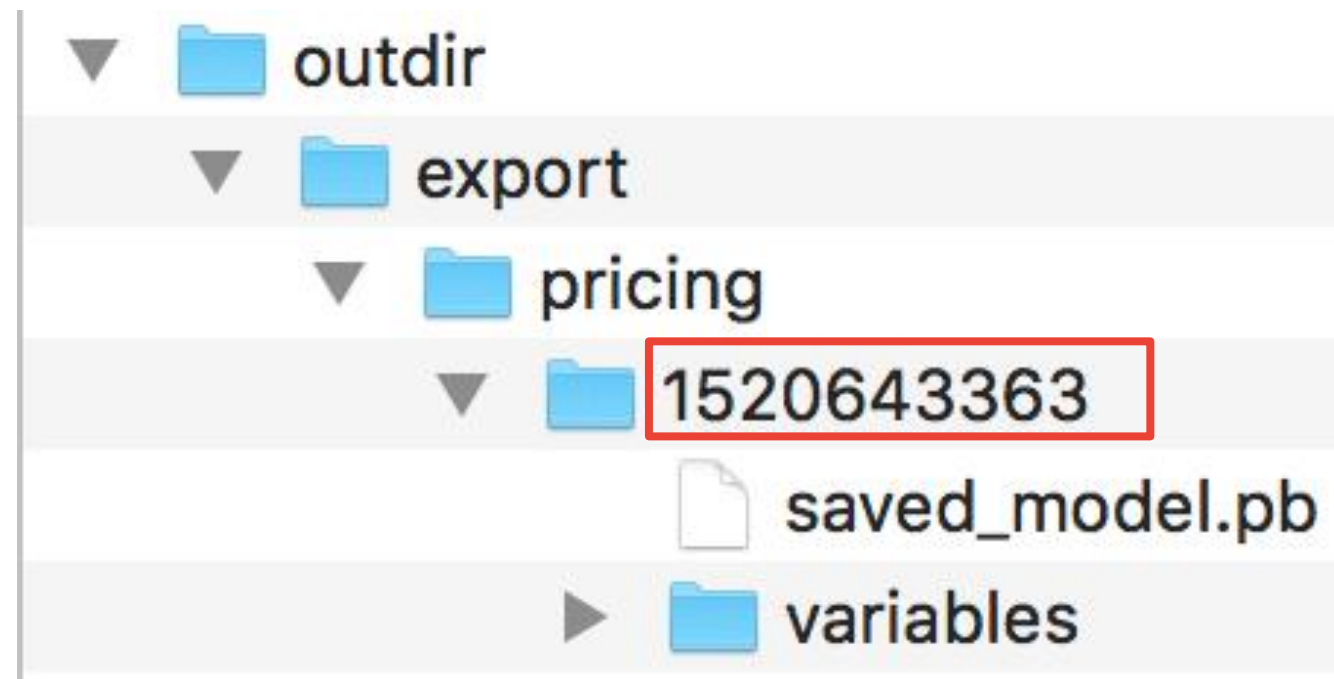
REST
JSON



predictions

```
gcloud ml-engine predict
--model <model_name>
--json-instances data.json
```

The exported model is ready to deploy



To test locally :

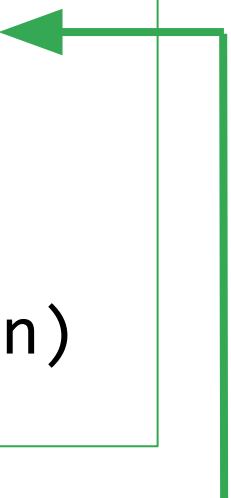
```
{"sq_footage": 1520, "prop_type": "apt"},  
{"sq_footage": 1520, "prop_type": "house"}
```

A diagram with two green arrows. One arrow starts from the JSON instances and points down to the command line. The other arrow starts from the word 'model' and points down to the folder '1520643363' in the command line.

```
gcloud ml-engine local predict \  
  --model-dir outdir/export/pricing/1520643363 \  
  --json-instances test_jsons.txt
```

Example serving input function that decodes JPEGs

```
def serving_input_fn():  
    json = {'jpeg_bytes': tf.placeholder(tf.string, [None])}  
  
    def decode(jpeg):  
        pixels = tf.image.decode_jpeg(jpeg, channels=3)  
        return pixels  
  
    pics = tf.map_fn(decode, json['jpeg_bytes'], dtype=tf.uint8)  
  
    features = {'pics': pics}  
    return tf.estimator.export.ServingInputReceiver(features, json)
```



Output shape:
[batch, width, height, 3]

Example serving input function that decodes JPEGs

```
def serving_input_fn():  
    json = {'jpeg_bytes': tf.placeholder(tf.string, [None])}  
  
    def decode(jpeg):  
        pixels = tf.image.decode_jpeg(jpeg, channels=3)  
        return pixels  
  
    pics = tf.map_fn(decode, json['jpeg_bytes'], dtype=tf.uint8)  
  
    features = {'pics': pics}  
    return tf.estimator.export.ServingInputReceiver(features, json)
```

Output shape:
[batch, width, height, 3]

Example serving input function that decodes JPEGs

JSON

```
{ "jpeg_bytes": { "b64": "/9j/4DAwQDBAgEBAgQCwkLEBAQEBAQEQEBA..." } }
```

Base 64: special syntax `_bytes` and `b64`

Lab: Implementing a distributed TensorFlow model



- ✓ Use train_and_evaluate
- ✓ Monitor training using TensorBoard

	Problem	Solution
✓	Out of memory data	Use the Dataset API
✓	Distribution	Use <code>train_and_evaluate</code>
✓	Need to evaluate during training	Use <code>train_and_evaluate</code> + TensorBoard
✓	Deployments that scale	Use serving input function

cloud.google.com